



CS 423

Operating System Design: Concurrency

Tianyin Xu

* Thanks for Prof. Adam Bates for the slides.

Concurrency vs Parallelism



Two tasks

1. Get a visa
2. Prepare slides

1. Sequential execution
2. Concurrent execution
3. Parallel execution
4. Concurrent but not parallel
5. Parallel but not concurrent
6. Parallel and concurrent

Why Concurrency?



- Servers
 - Multiple connections handled simultaneously
- Parallel programs
 - To achieve better performance
- Programs with user interfaces
 - To achieve user responsiveness while doing computation
- Network and disk bound programs
 - To hide network/disk latency

Definitions

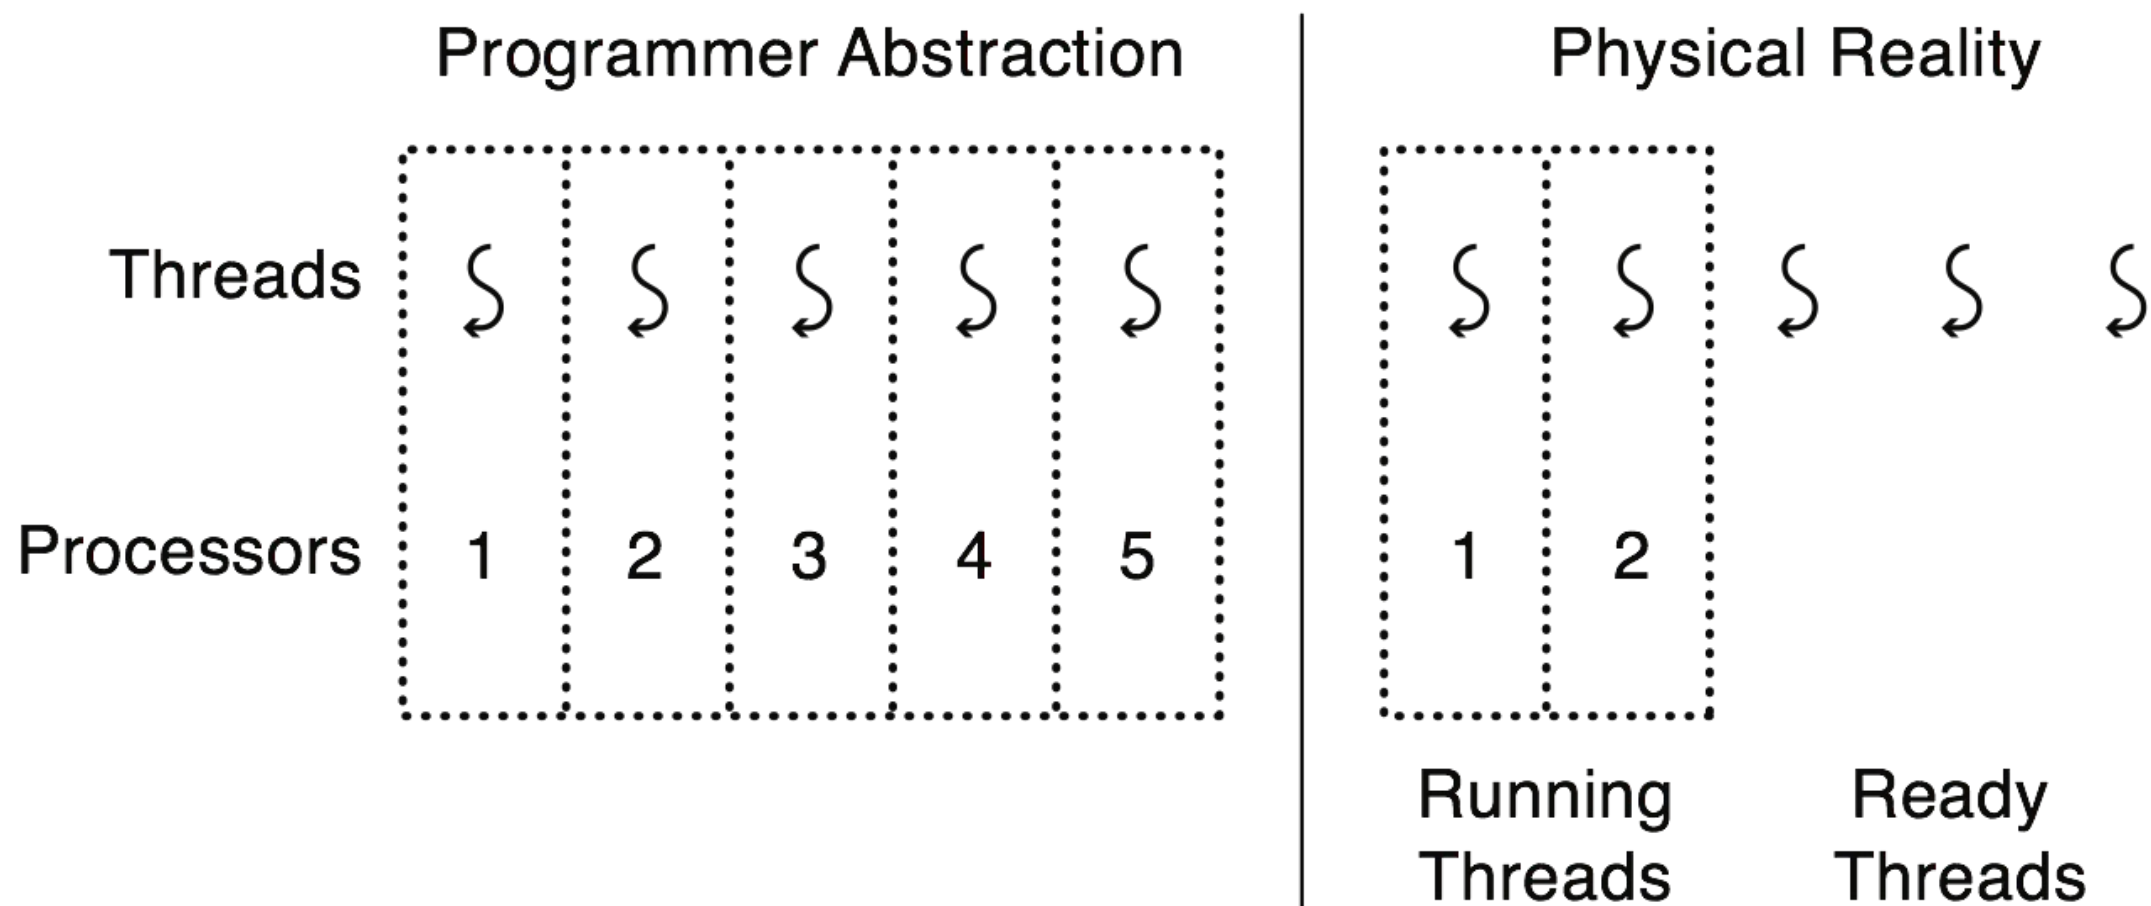


- Thread: A single execution sequence that represents a separately schedulable task.
 - *Single execution sequence*: intuitive and familiar programming model
 - *separately schedulable*: OS can run or suspend a thread at any time.
 - Schedulers operate over threads/tasks, both kernel and user threads.
- *Does the OS protect all threads from one another?*

The Thread Abstraction



- Infinite number of processors
- Threads execute with variable speed



Programmer vs. Processor View



Programmer View

Programmer's View

.
.
.
x = x + 1;
y = y + x;
z = x + 5y;
.
.
.

Possible Execution #1

.
.
.
x = x + 1;
y = y + x;
z = x + 5y;
.
.
.

Possible Execution #2

.
.
.
x = x + 1;
.....
Thread is suspended.
Other thread(s) run.
Thread is resumed.
.....
y = y + x;
z = x + 5y;

Possible Execution #3

.
.
.
x = x + 1;
y = y + x;
.....
Thread is suspended.
Other thread(s) run.
Thread is resumed.
.....
z = x + 5y;


Variable Speed: Program must anticipate all of these possible executions

Possible Executions




Processor View


One Execution



Another Execution



Another Execution



Something to look forward to when we discuss scheduling!

Thread Ops



- `thread_create(thread, func, args)`
Create a new thread to run `func(args)`
- `thread_yield()`
Relinquish processor voluntarily
- `thread_join(thread)`
In parent, wait for forked thread to exit, then return
- `thread_exit`
Quit thread and clean up, wake up joiner if any

Ex: threadHello



```
#define NTHREADS 10
thread_t threads[NTHREADS];
main() {
    for (i = 0; i < NTHREADS; i++) thread_create(&threads[i], &go, i);
    for (i = 0; i < NTHREADS; i++) {
        exitValue = thread_join(threads[i]);
        printf("Thread %d returned with %ld\n", i, exitValue);
    }
    printf("Main thread done.\n");
}
void go (int n) {
    printf("Hello from thread %d\n", n);
    thread_exit(100 + n);
    // REACHED?
}
```

Ex: threadHello output



```
bash-3.2$ ./threadHello
Hello from thread 0
Hello from thread 1
Thread 0 returned 100
Hello from thread 3
Hello from thread 4
Thread 1 returned 101
Hello from thread 5
Hello from thread 2
Hello from thread 6
Hello from thread 8
Hello from thread 7
Hello from thread 9
Thread 2 returned 102
Thread 3 returned 103
Thread 4 returned 104
Thread 5 returned 105
Thread 6 returned 106
Thread 7 returned 107
Thread 8 returned 108
Thread 9 returned 109
Main thread done.
```

- Must “thread returned” print in order?
- What is maximum # of threads that exist when thread 5 prints hello?
- Minimum?
- Why aren’t any messages interrupted mid-string?

Create/Join Concurrency



- Threads can create children, and wait for their completion
- Data only shared before fork/after join
- Examples:
 - Web server: fork a new thread for every new connection
 - As long as the threads are completely independent
 - Merge sort
 - Parallel memory copy


Ex: bzero




```
void blockzero (unsigned char *p, int length) {
    int i, j;
    thread_t threads[NTHREADS];
    struct bzeroparams params[NTHREADS];

    // For simplicity, assumes length is divisible by NTHREADS.
    for (i = 0, j = 0; i < NTHREADS; i++, j += length/NTHREADS) {
        params[i].buffer = p + i * length/NTHREADS;
        params[i].length = length/NTHREADS;
        thread_create_p(&(threads[i]), &go, &params[i]);
    }
    for (i = 0; i < NTHREADS; i++) {
        thread_join(threads[i]);
    }
}
```

Thread Data Structures



Thread Lifecycle




Thread Implementations



- Kernel threads
 - Thread abstraction only available to kernel
 - To the kernel, a kernel thread and a single threaded user process look quite similar
- Multithreaded processes using kernel threads
 - Kernel thread operations available via syscall
- User-level threads
 - Thread operations without system calls

Multithreaded OS Kernel



Implementing Threads



- Thread_fork(func, args)
 - Allocate thread control block
 - Allocate stack
 - Build stack frame for base of stack (stub)
 - Put func, args on stack
 - Put thread on ready list
 - Will run sometime later (maybe right away!)
- stub(func, args):
 - Call (*func)(args)
 - If return, call thread_exit()

Implementing Threads



- Thread_Exit
 - Remove thread from the ready list so that it will never run again
 - Free the per-thread state allocated for the thread

Ex: Two Threads call Yield



Thread 1's instructions

“return” from thread_switch
into stub
call go
call thread_yield
choose another thread
call thread_switch
save thread 1 state to TCB
load thread 2 state

return from thread_switch
return from thread_yield
call thread_yield
choose another thread
call thread_switch

Thread 2's instructions

“return” from thread_switch
into stub
call go
call thread_yield
choose another thread
call thread_switch
save thread 2 state to TCB
load thread 1 state

Processor's instructions

“return” from thread_switch
into stub
call go
call thread_yield
choose another thread
call thread_switch
save thread 1 state to TCB
load thread 2 state
“return” from thread_switch
into stub
call go
call thread_yield
choose another thread
call thread_switch
save thread 2 state to TCB
load thread 1 state
return from thread_switch
return from thread_yield
call thread_yield
choose another thread
call thread_switch



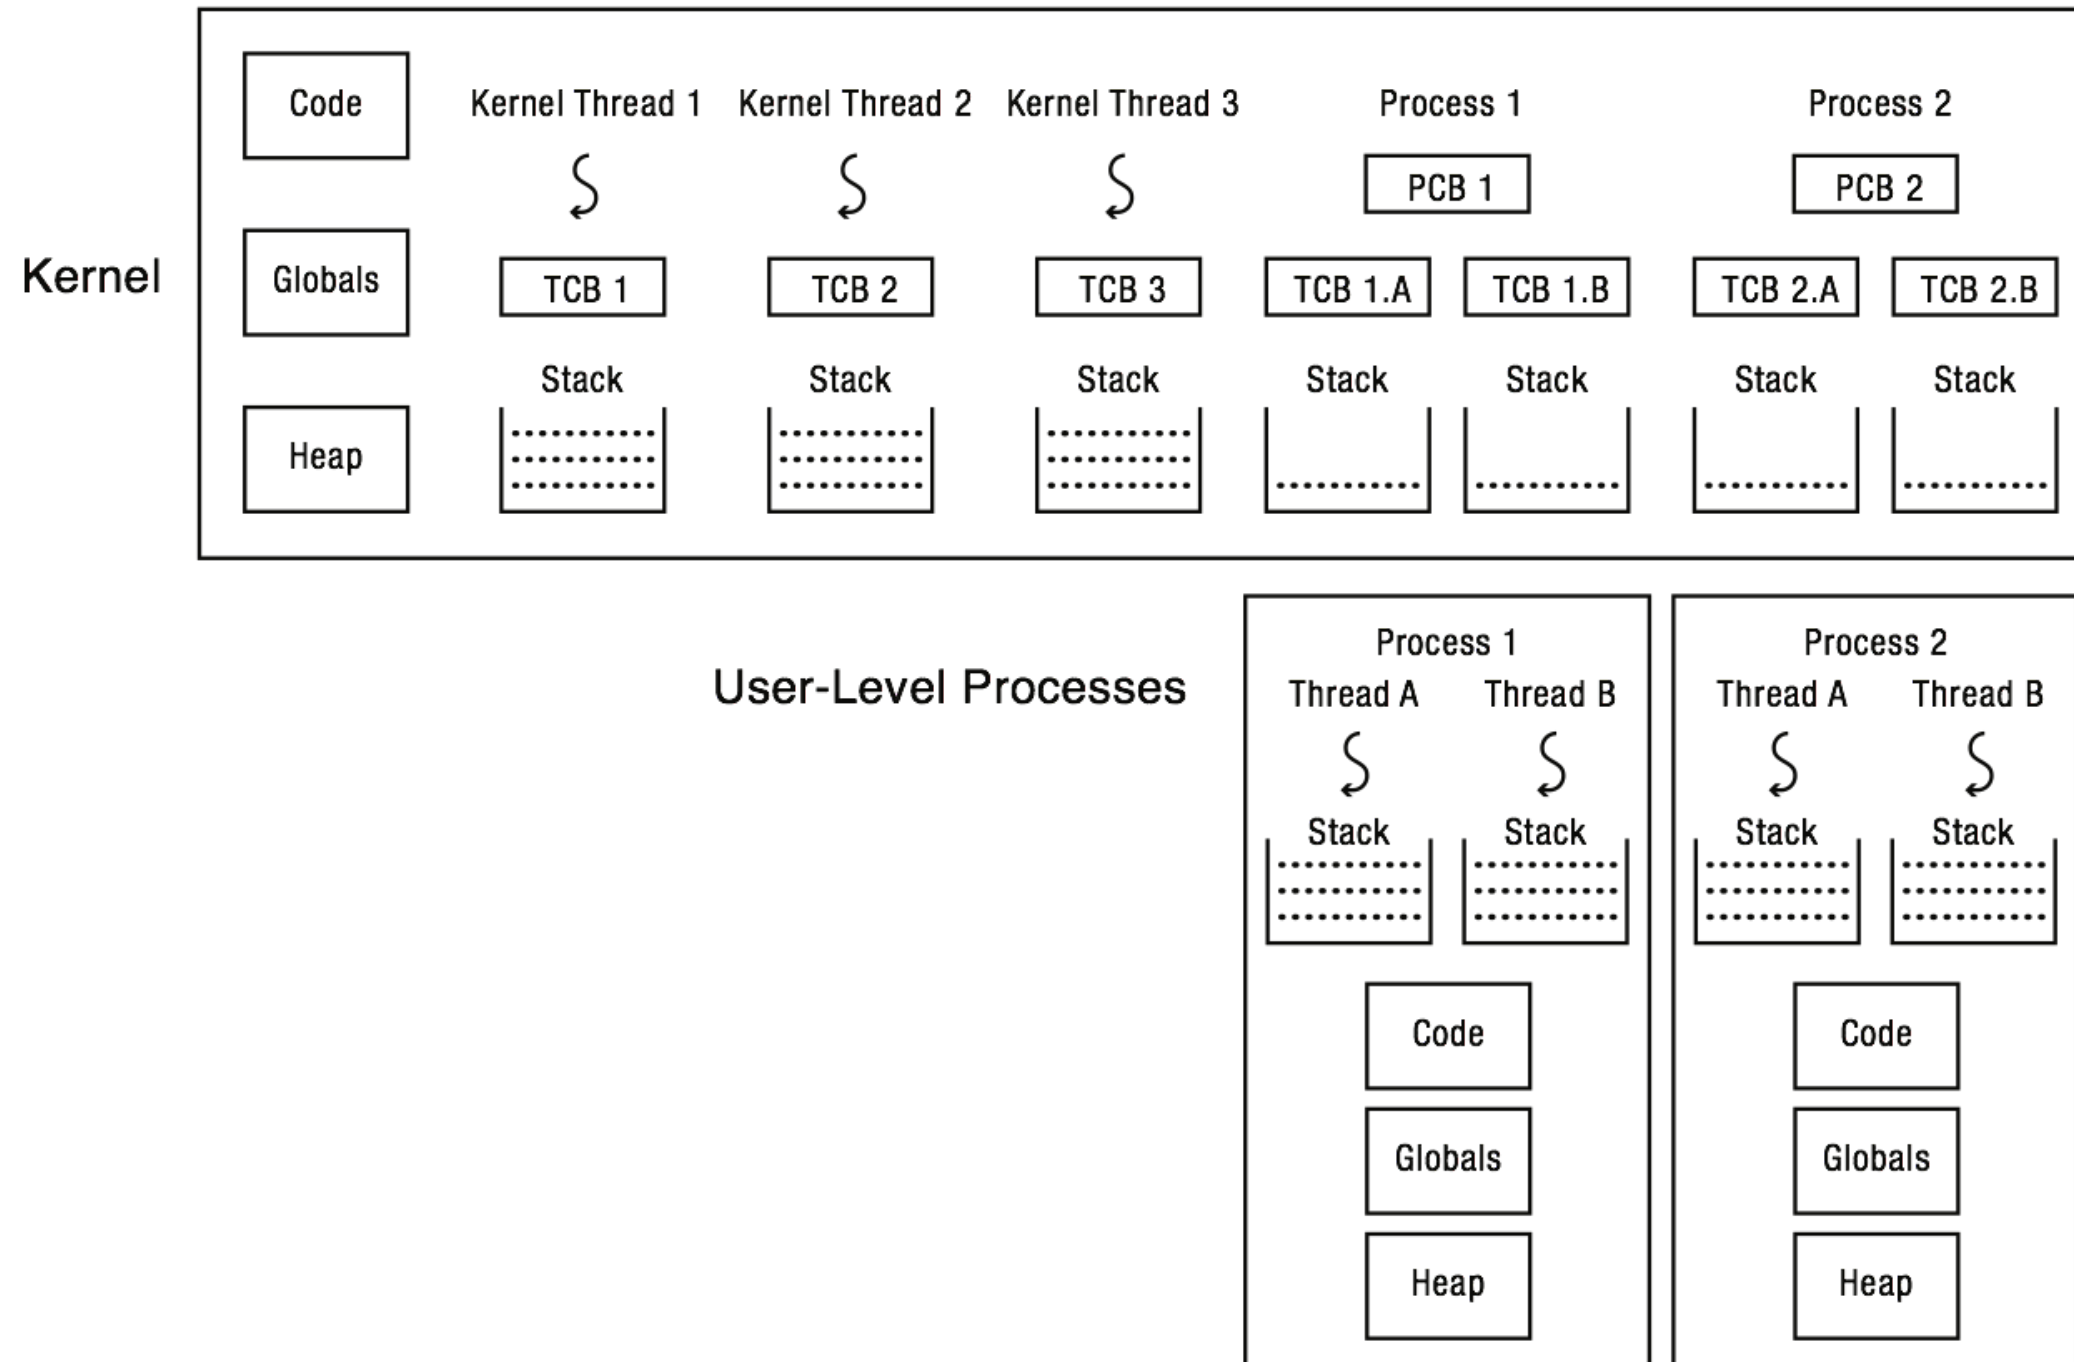
Take 1:

- User thread = kernel thread (Linux, MacOS)
 - System calls for thread fork, join, exit (and lock, unlock,...)
 - Kernel does context switch
 - Simple, but a lot of transitions between user and kernel mode

Multi-threaded User Processes



Take 1:





Take 2:

- Green threads (early Java)
 - User-level library, within a single-threaded process
 - Library does thread context switch
 - Preemption via upcall/UNIX signal on timer interrupt
 - Use multiple processes for parallelism
 - Shared memory region mapped into each process




Take 3:

- Scheduler activations (Windows 8):
 - Kernel allocates processors to user-level library
 - Thread library implements context switch
 - Thread library decides what thread to run next
- Upcall whenever kernel needs a user-level scheduling decision:
 - Process assigned a new processor
 - Processor removed from process
 - System call blocks in kernel

Multi-threaded User Processes



Take 3: (What's old is new again)



M:N model multiplexes N user-level threads onto M kernel-level threads

Good idea? Bad Idea?

Question



Compare event-driven programming with multithreaded concurrency. Which is better in which circumstances, and why?