

# MP1 Walkthrough

Cathy Cai  
9/16

Credits: Siyuan Chai, Jinghao Jia

# Get Starter Code

- [https://classroom.github.com/a/HXioa\\_mx](https://classroom.github.com/a/HXioa_mx)
- Find your name and click (don't click on other's name!)
- Due Oct. 2nd at 11:59 PM CT

Join the classroom:

**cs423-uiuc-fall25**

To join the GitHub Classroom for this course, please select yourself from the list below to associate your GitHub account with your school's identifier (i.e., your name, ID, or email).

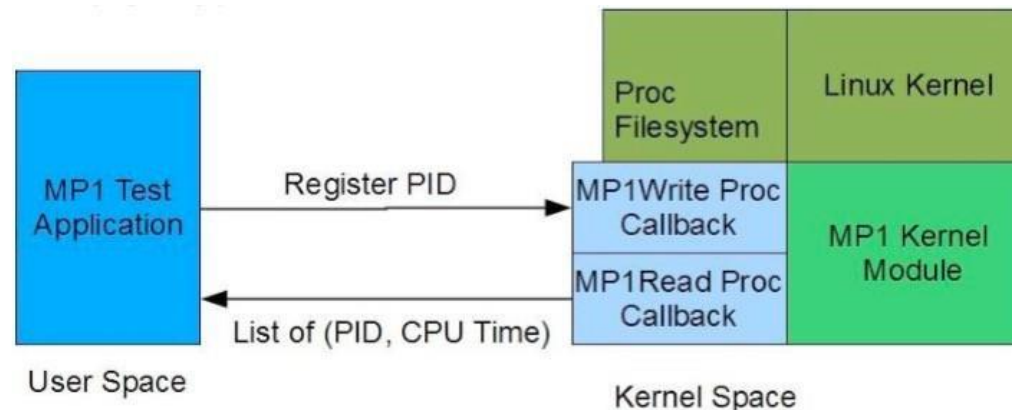
Can't find your name? [Skip to the next step →](#)

- P.S. Don't forget to submit your MP0!

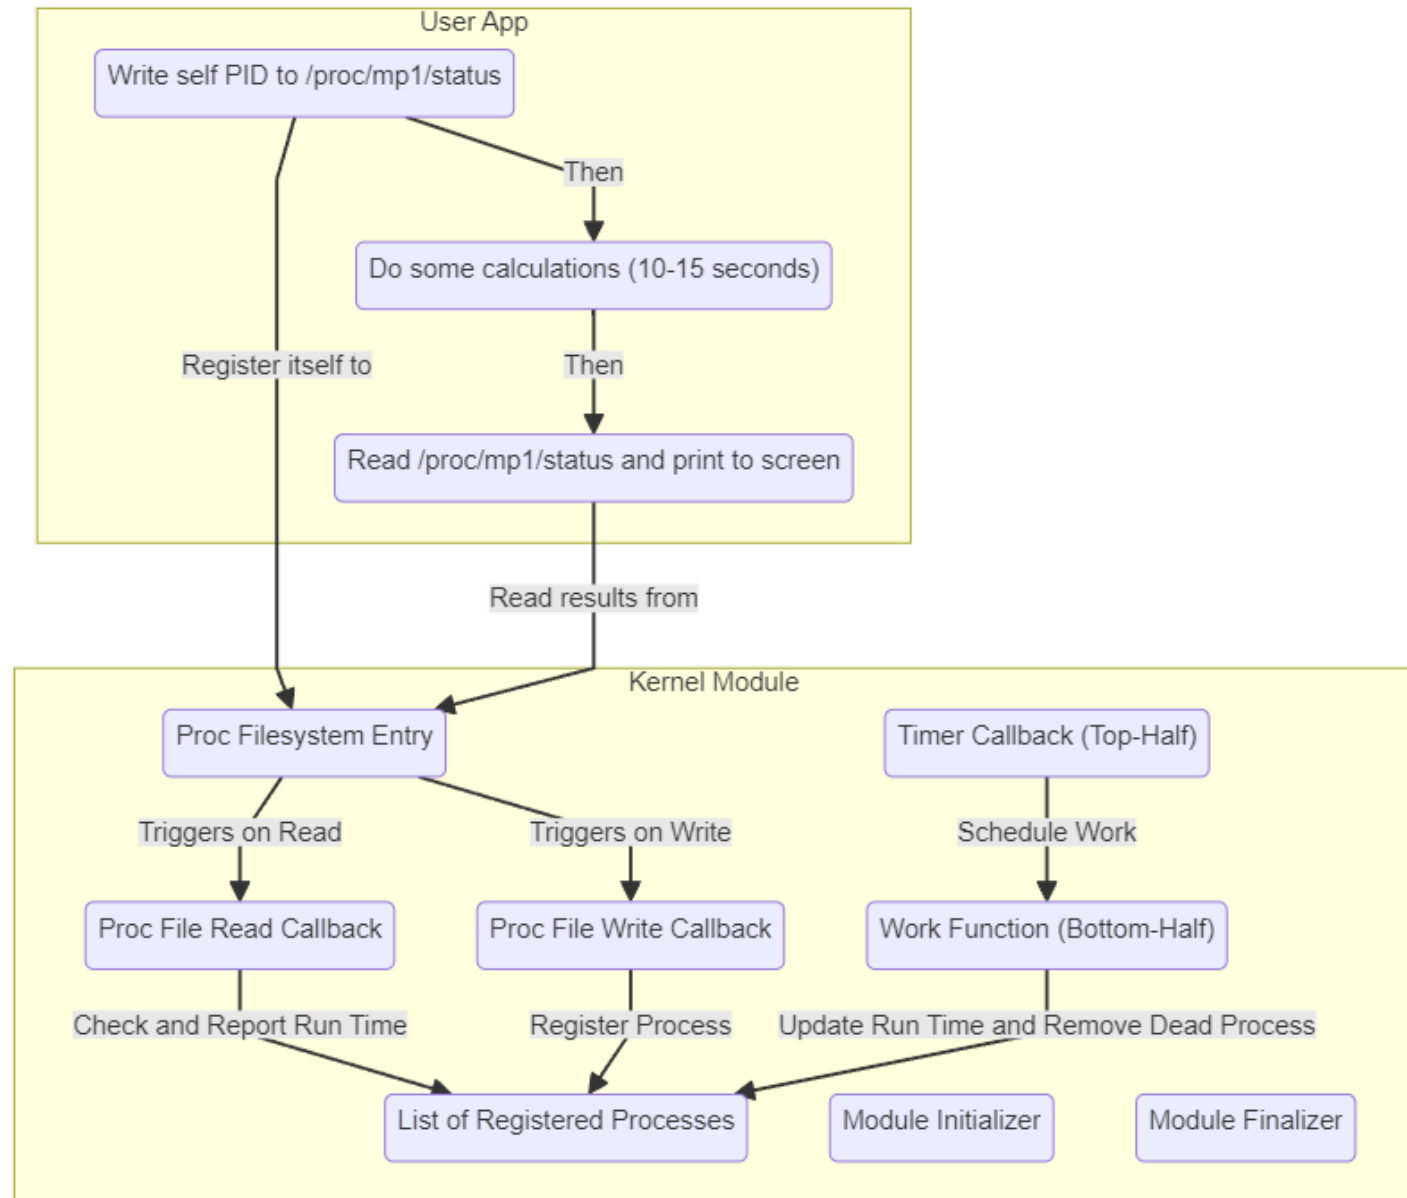
Identifiers

# Problem Description

- Write a kernel module that measures the userspace CPU
- Time of processes registered within the kernel module
- Register processes using PID through the Proc Filesystem
- Kernel module updates the userspace CPU time of each registered process every 5s
- Print the userspace CPU time of each registered process

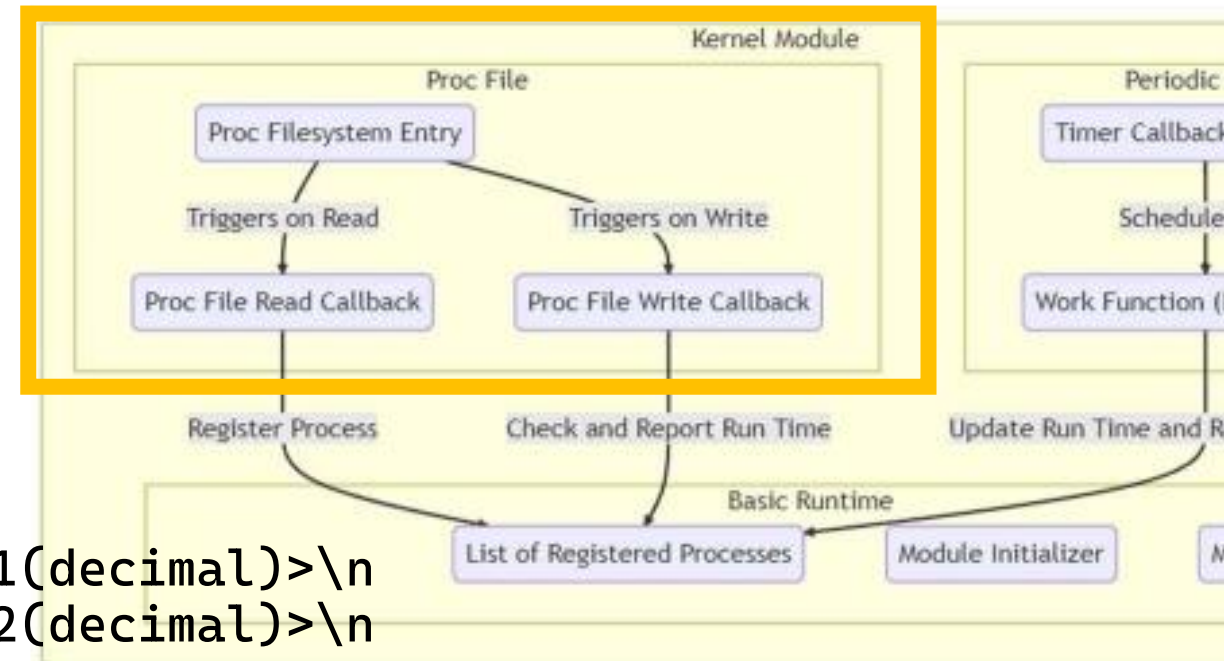


# Overview



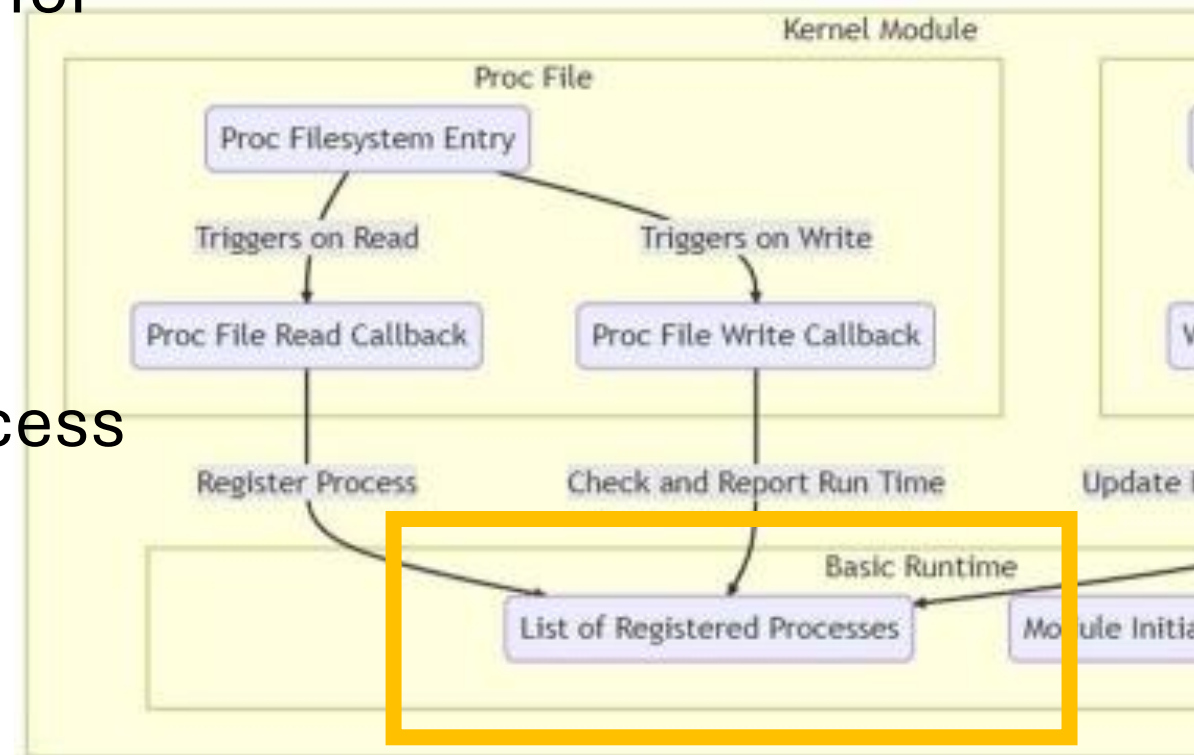
# Proc Filesystem Entry

- Not regular files, does not store data in binary format
- Can be read/write as regular files
- Create an entry (e.g. `/proc/mp1/status`) in the proc filesystem
  - `proc_mkdir()`
  - `proc_create()`
- Register a process:
  - `echo "pid" > /proc/mp1/status`
  - Use `fprintf()`, etc.
- Get userspace CPU time:
  - `cat /proc/mp1/status`
  - Use `fgets()`, etc.
  - Should print in the following format:
    - `<PID1>:[space]<CPU time of PID1(decimal)>\n`
    - `<PID2>:[space]<CPU time of PID2(decimal)>\n`



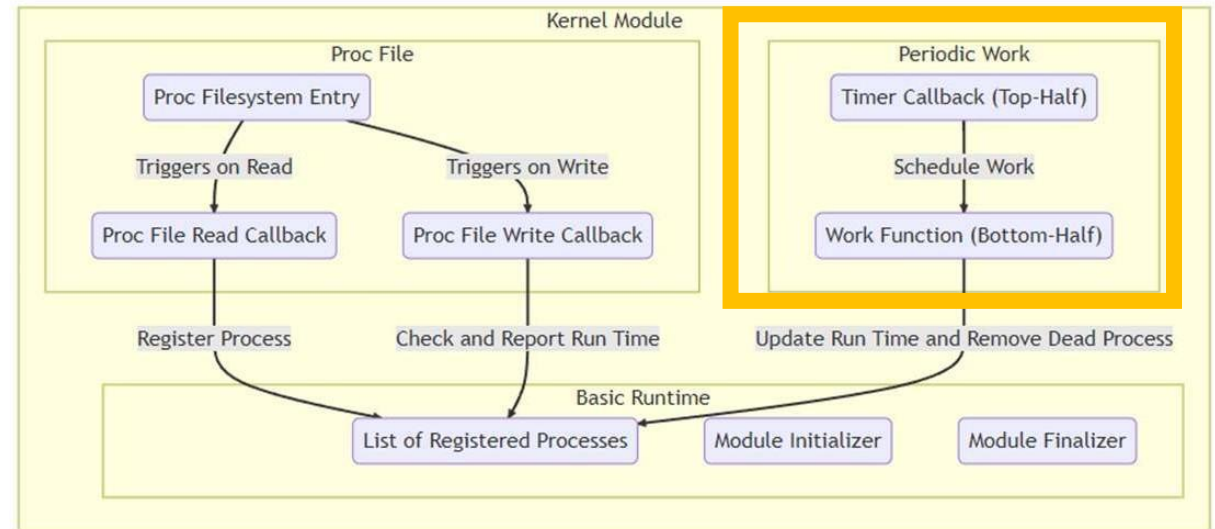
# Store States

- Implement read and write callback for the proc entry
  - `proc_read()`
  - `proc_write()`
- Use kernel linked list to store the information of every registered process
  - APIs in `<linux/list.h>`
- Need to consider concurrency for linked list operations
  - E.g. using a lock



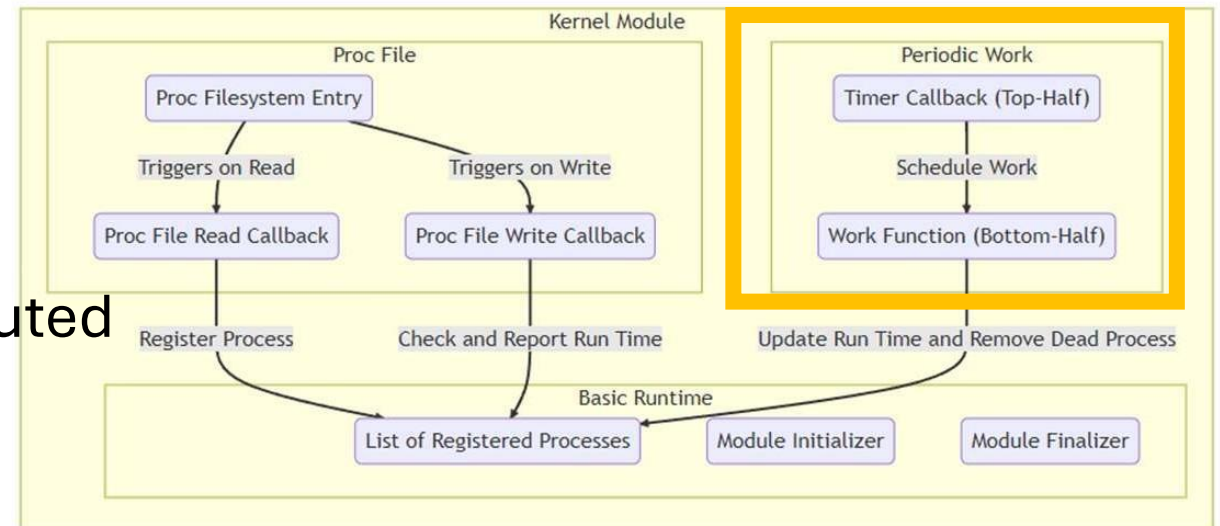
# Periodic work: timer

- Use a kernel timer to perform a task after a preset timeout
  - APIs in `<linux/timer.h>`
- Setup timer
  - `timer_setup(timer, callback, flags)`
- Setup timeout
  - Timeout is represented in jiffy in kernel. Jiffy can be converted between regular time units (s, ms, etc.)
  - `mod_timer(timer, expires)`
- Challenge: timer only fire once



# Periodic work: two-halves

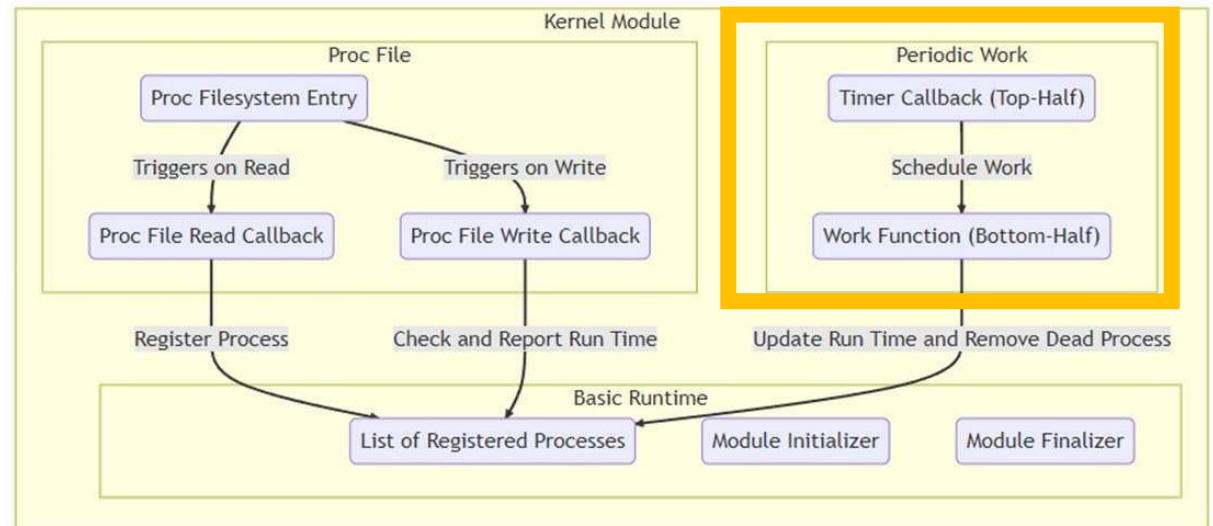
- Not put all work in timer handler
  - Why?
  - Registered list can be long. Better not to not block other timers
- Use a two-halves approach
- Use kernel work queue
  - allow kernel functions to be executed by special kernel threads
  - APIs in `<linux/workqueue.h>`





# Work Queue

- Schedule a function to be run in a work queue
  - `queue_work(work_queue, work)`
  - callback only calls `queue_work()` (Top-Half)
  - work is where we are going to do the actual updates (Bottom-Half)



# Passing Data

- Passing data between kernelspace and userspace
  - E.g. `ssize_t proc_read(struct file *file, char user *buf, size_t size, loff_t *loff)`
  - `buf` here is a userspace address and can't be dereferenced directly in kernel space
  - Use `copy_from_user()/...` to copy to a kernel buffer
  - Same for `copy_to_user()/...`
- Free/deallocate any memory/objects before exiting the kernelmodule
  - Dynamic allocated memory using `kmalloc()` must be freed using `kfree()`
  - Objects such as timer/work\_queue must be destroyed
  - Proc FS entry must be removed

# Debug and Submission

- Debug
  - Use `printk()` to print to the kernel log
    - View the kernel log using `dmesg` (e.g. `dmesg | less`)
    - Works on any platform
  - Sufficient for MP1 (from my experience)
- Use `gdb`
  - A bit tricky to load the symbol table for kernel module. You can ask Cathy/Peizhe how to do that
- Submission
  - Push your code to your GitHub repo before ddl