

Kernel Debugging 101

Siyuan Chai



A bit about myself



- Fifth year PhD student working with Prof. Tianyin Xu
 - schai.me, siyuanc3@illinois.edu
- TA for 423 in Fall 2022, 2023, 2024
- Work on operating systems, memory systems
 - Building OS framework for emerging address translation schemes
- Interned at NVIDIA, Meta, and Google
 - Work on GPU memory, AI systems etc.
- For fun
 - Badminton, frisbee, soccer, hit the gym

Outline



- How to understand error message?
- What are the helpful kernel tools?
 - printk
 - GDB
 - Memory debugging
- Other common errors
 - Kernel version mismatch

Warmup: Version Mismatch



- `root@q:~/cs423/mp1-siyuanchai1999# insmod mp1.ko`
- `[216.683949][T220] mp1: disagrees about version of symbol module_layout`

Warmup: Version Mismatch



- Compile mp1 module on host (UTM/Vmware/WSL)
 - Run with ubuntu 24
 - Default kernel version 6.2+
- Launch cs423-q in Linux 5.15 folder
 - VM runs with 5.15 version
 - Symbol layout are different across versions
- How to fix?
 - Compile inside cs423-q VM
 - Update your Makefile to run on host
 - `make -C /path/to/your/kernel/source M=$(PWD) modules`

Kernel Error Message



```
[ 1783.059199][ T896] Kernel panic - not syncing: stack-protector: Kernel stack is corrupted in:
mp1_read+0x12a/0x130 [mp1]
[ 1783.059904][ T896] CPU: 0 PID: 896 Comm: cat Tainted: G          0      5.15.165 #2
[ 1783.060386][ T896] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.15.0-1
04/01/2014
[ 1783.060943][ T896] Call Trace:
[ 1783.061141][ T896]  <TASK>
[ 1783.061319][ T896]  dump_stack_lvl+0x34/0x48
[ 1783.061600][ T896]  panic+0xfb/0x2b9
[ 1783.061842][ T896]  ? mp1_read+0x12a/0x130 [mp1]
[ 1783.062135][ T896]  __stack_chk_fail+0x10/0x10
[ 1783.062418][ T896]  mp1_read+0x12a/0x130 [mp1]
[ 1783.062701][ T896]  proc_reg_read+0x4d/0x90
[ 1783.062972][ T896]  vfs_read+0x8d/0x190
[ 1783.063222][ T896]  ksys_read+0x5a/0xee
```

Kernel Error Message



```
[ 1783.065475][ T896] RIP: 0033:0x7fedc434a7e2
[ 1783.065770][ T896] Code: c0 e9 b2 fe ff ff 50 48 8d 3d 8a b4 0c 00 e8 a5 1d 02 00 0f 1f 44 00
00 f3 0f 1e fa 64 8b 04 25 18 00 00 00 85 c0 75 10 0f 05 <48> 3d 00 f0 ff ff 77 56 c3 0f 1f 44 00
00 48 83 ec 28 48 89 54 24
[ 1783.066948][ T896] RSP: 002b:00007ffde1040308 EFLAGS: 00000246 ORIG_RAX: 0000000000000000
[ 1783.067463][ T896] RAX: ffffffffda RBX: 0000000000020000 RCX: 00007fedc434a7e2
[ 1783.067937][ T896] RDX: 0000000000020000 RSI: 00007fedc4212000 RDI: 0000000000000003
[ 1783.068398][ T896] RBP: 00007fedc4212000 R08: 00007fedc4211010 R09: 00007fedc4211010
[ 1783.068856][ T896] R10: 0000000000000022 R11: 0000000000000246 R12: 0000000000022000
[ 1783.069321][ T896] R13: 0000000000000003 R14: 0000000000020000 R15: 0000000000020000
[ 1783.069783][ T896] </TASK>
[ 1783.070046][ T896] Kernel Offset: disabled
[ 1783.070312][ T896] ---[ end Kernel panic - not syncing: stack-protector: Kernel stack is
corrupted in: mp1_read+0x12a/0x130 [mp1] ]---
```



Locate buggy code from stack trace

- Add EXTRA_CFLAGS += -g to Makefile
 - Recompile kernel modules
- `eu-addr2line -e module.ko function:offset`
- `eu-addr2line -e vmlinux function:offset`
 - Apply to kernel functions

Connect with GDB to debug



- Kernel configs
 - Turn on CONFIG_GDB_SCRIPTS, CONFIG_KGDB
 - Turn *off* CONFIG_DEBUG_INFO_REDUCED, CONFIG_RANDOMIZE_BASE
 - Recompile kernel and make scripts_gdb
- `../qemu-script/cs423-q`
 - QEMU flag to setup gdb server at port 1234
- (on UTM/WSL/VMware) `gdb vmlinux`
 - (inside gdb) `target remote localhost:1234`

Connect with GDB to debug



- Inside QEMU VM
 - insmod mp1.ko
 - lx-symbols ../mp1-githubID/
 - GDB terminal will show loading ../mp1-githubID/mp1.ko
 - b mp1_read # set break point
 - Watch *(int *)0xdeadbeef # memory watch point
 - Int for 4 bytes
 - finish # get out of current function
 - p VAR # print variable

GDB FAQ



- lx-symbols not defined
 - Missing make scripts_gdb
 - Add this line to ~/.gdbinit
`add-auto-load-safe-path
path/to/kernel_source/scripts/gdb/vmlinux-gdb.py`
- Fancy gdb support
 - <https://github.com/hugsy/gef>

Check memory leaking with kmemleak



- Turn on CONFIG_DEBUG_KMEMLEAK
- `mount -t debugfs nodev /sys/kernel/debug/`
- `cat /sys/kernel/debug/kmemleak`
- `echo scan > /sys/kernel/debug/kmemleak`
- Reference: <https://docs.kernel.org/dev-tools/kmemleak.html>

Quick Summary



- eu-addr2line: locate your code from back trace
- printk: quick and intuitive
 - Requires a few iterations to get the correct point
 - Requires recompiling, can be time consuming if hacking kernel code
- GDB
 - Observe pretty much all the code behavior
 - Hard to capture historical info