



CS 423

Operating System Design

<https://cs423-uiuc.github.io>

Tianyin Xu

tyxu@illinois.edu

* Thanks Adam Bates for the slides.

Scheduling



- A forever topic in Computer Systems and Life
 - Uniprocessor: 100 threads in the ready queue – which one to run next?
 - Multiprocessor: 400 threads in the ready queues of four cores – which one to run next on which core?
 - Cluster: 1000 MapReduce jobs – which one to run on which machine and on which core?
 - Datacenters: 10000 user request – which one to run on which datacenter on which cluster on which machine?

More complexity



- Jobs/requests are not created equal.
 - Some are more important than the others
- Jobs/requests could have deadlines
 - Finishing late means nothing but wasting resources.
- Jobs/requests have constraints
 - Affinity is important – same node and same PCIe switch for GPUs
- Workloads could be very different.

Scheduling



- Always an active research topic
 - Everyone wants run more jobs with less resources
- In this class, we are going to focus on the simplest setup – a uniprocessor

What Are Scheduling Goals?



- What are the goals of a scheduler?
- Scheduling Goals:
 - Generate illusion of concurrency
 - Maximize resource utilization (e.g., mix CPU and I/O bound processes appropriately)
 - Meet needs of both I/O-bound and CPU-bound processes
 - Give I/O-bound processes better interactive response
 - Do not starve CPU-bound processes
 - Support Real-Time (RT) applications



Definitions



- **Task/Job**
 - Something that needs CPU time: a thread associated with a process or with the kernel...
 - ... a user request, e.g., mouse click, web request, shell command, ...
- **Latency/response time**
 - How long does a task take to complete?
- **Throughput**
 - How many tasks can be done per unit of time?

Definitions



- **Overhead**
 - How much extra work is done by the scheduler?
- **Fairness**
 - How equal is the performance received by different users?
- **Predictability**
 - How consistent is the performance over time?
- **Starvation**
 - A task 'never' receives the resources it needs to complete
 - Not very fair : - (

Definitions




- **Workload**
 - Set of tasks for system to perform
- **Work-conserving**
 - Resource is used whenever there is a task to run
 - For non-preemptive schedulers, work-conserving is not always better

Definitions



- **Non-preemptive scheduling:**

- The running process keeps the CPU until it **voluntarily** gives up the CPU
 - process exits
 - switches to blocked state
 - 1 and 4 only (no 3)



- **Preemptive scheduling:**

- The running process can be interrupted and must release the CPU (can be **forced** to give up CPU)



- Scheduling algorithm
 - takes a workload as input
 - decides which tasks to do first
 - Performance metric (throughput, latency) as output
 - Only preemptive, work-conserving schedulers to be considered

First In First Out (FIFO)



- Schedule tasks in the order they arrive
 - Continue running them until they complete or give up the processor
- On what workloads would FIFO be particularly bad?

Shortest Job First (SJF)




- Always do the task that has the shortest remaining amount of work to do
 - Often called Shortest Remaining Time First (SRTF)
- Suppose we have five tasks arrive one right after each other, but the first one is much longer than the others
 - Which completes first in FIFO? Next?
 - Which completes first in SJF? Next?

FIFO vs. SJF




Tasks

FIFO



Tasks

SJF




Time

Round Robin (RR)




- Each task gets resource for a fixed period of time (time quantum)
 - If task doesn't complete, it goes back in line
- Characteristics of scheduler change depending on the time quantum size
 - What if time quantum is too short?
 - One instruction?
 - What if time quantum is too long?
 - Infinite?

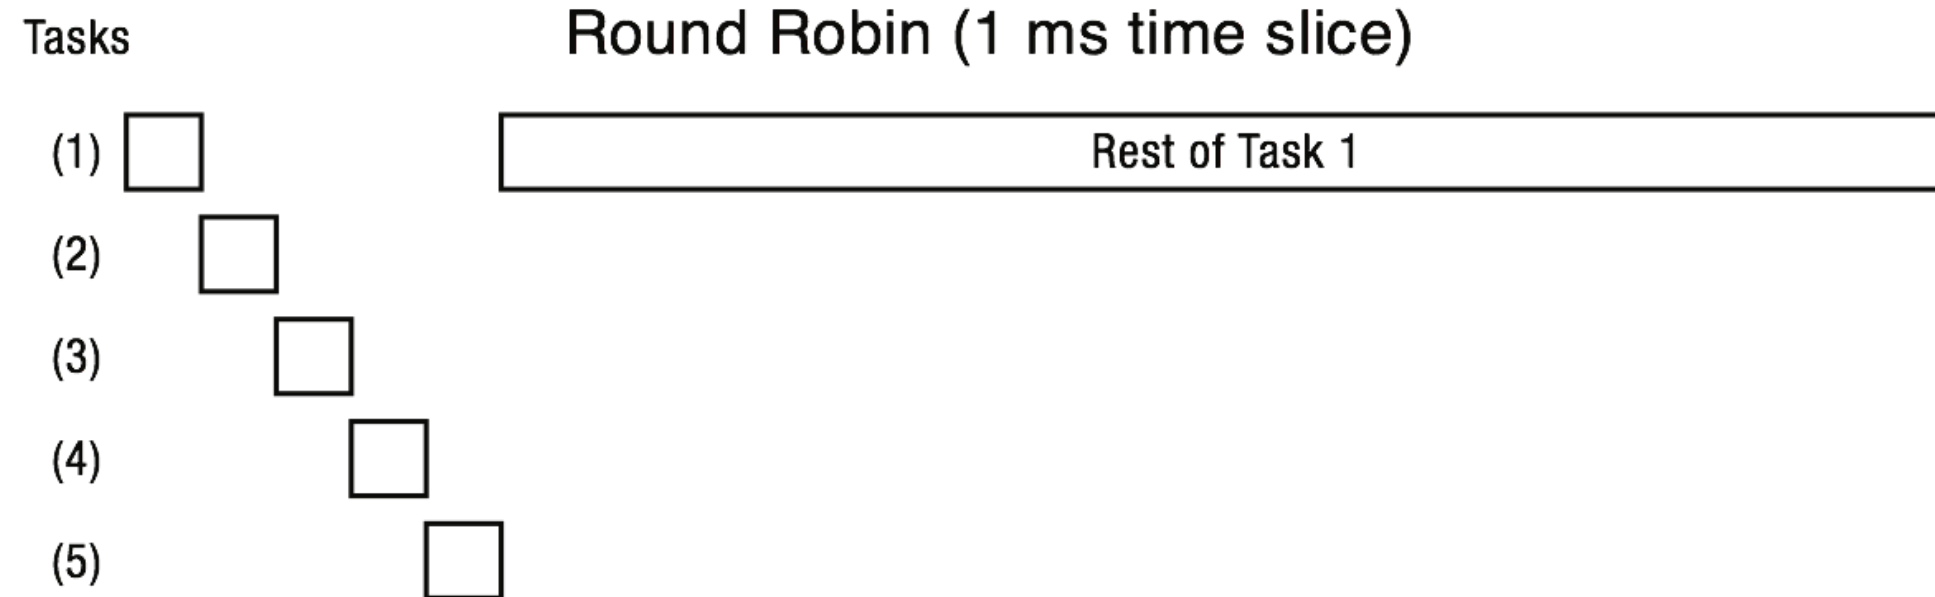
Round Robin




Round Robin



Round Robin



Round Robin





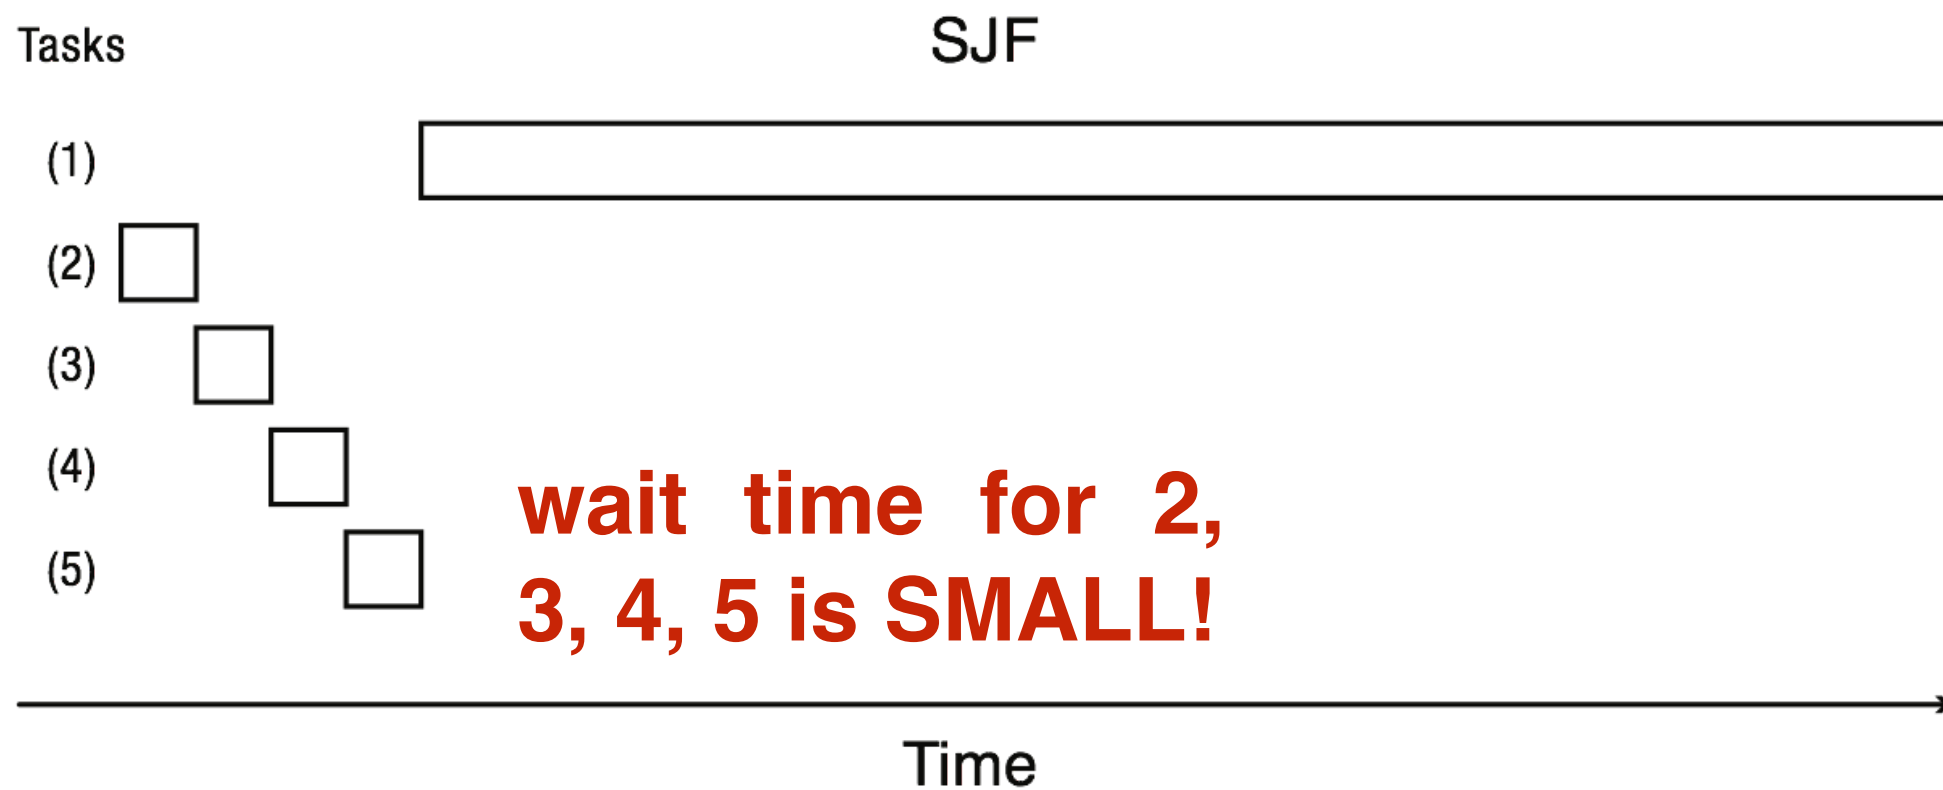
- Basic scheduling algorithms
 - FIFO (FCFS)
 - Shortest job first
 - Round Robin



- Basic scheduling algorithms
 - FIFO (FCFS)
 - Shortest job first
 - Round Robin

- What is an optimal algorithm in the sense of maximizing the number of jobs finished (i.e., minimizing average response time)?

FIFO vs. SJF


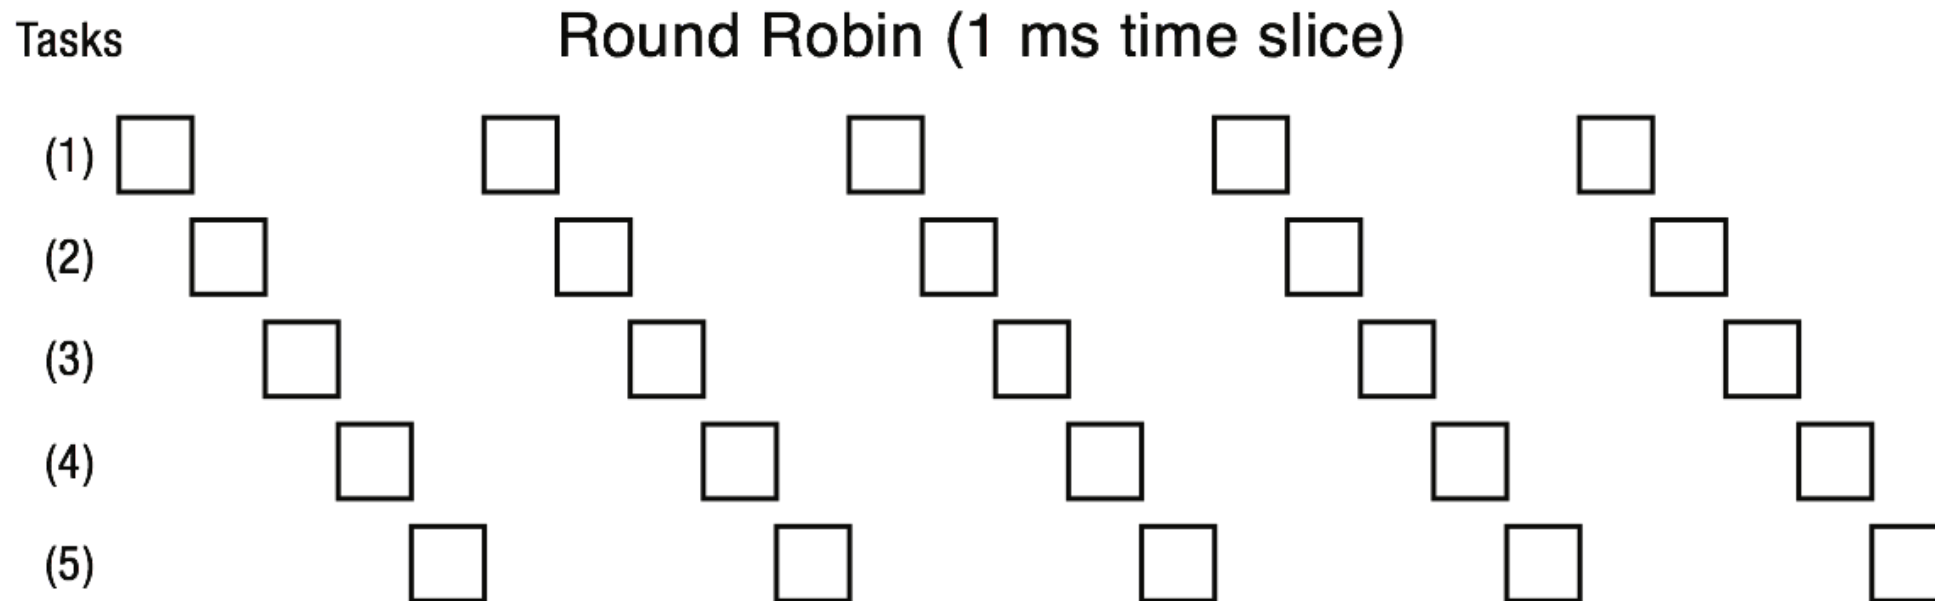




- Basic scheduling algorithms
 - FIFO (FCFS)
 - Shortest job first
 - Round Robin

- Assuming zero-cost to time slicing, is Round Robin always better than FIFO?

RR v. FIFO (fixed size tasks)



Starvation, Sample Bias



- Suppose you want to compare two scheduling algorithms
 - Create some infinite sequence of arriving tasks
 - Start measuring
 - Stop at some point
 - Compute average response time as the average for completed tasks between start and stop
- Is this valid or invalid?

Sample Bias Solutions



- Measure for long enough that # of completed tasks >> # of uncompleted tasks
 - For both systems!
- Start and stop system in idle periods
 - Idle period: no work to do
 - If algorithms are work-conserving, both will complete the same tasks

Round Robin = Fairness?



Is Round Robin the fairest possible algorithm?

What is fair?

- FIFO?
- Equal share of the CPU?
- What if some tasks don't need their full share?
- Minimize worst case divergence?
- Time task would take if no one else was running
- Time task takes under scheduling algorithm

Fairness needs to be defined.



- 4 kids share a cake.
 - Each gets 25% of the cake.
 - Quite fair!
- There is one little kids and the kid can only eat 10% of the cake.
 - We either force her to eat the 25% -- to be fair
 - Or we give 15% remaining to the other 3 kids.
 - **Min-max fairness**

Max-Min Fairness



- The ***least*** demanding one will get its fair share ***first***
- After this, the ***next least*** demanding one will get its fair share ***first***
- And so on...

Max-Min Fairness



- Kid 1: 20%
- Kid 2: 26%
- Kid 3: 40%
- Kid 4: 50%
- 100% -> 25% each kid
 - 20% -> 5% left -> 1.666666% to the other three
 - 25%
 - 25%
 - 25%

Max-Min Fairness



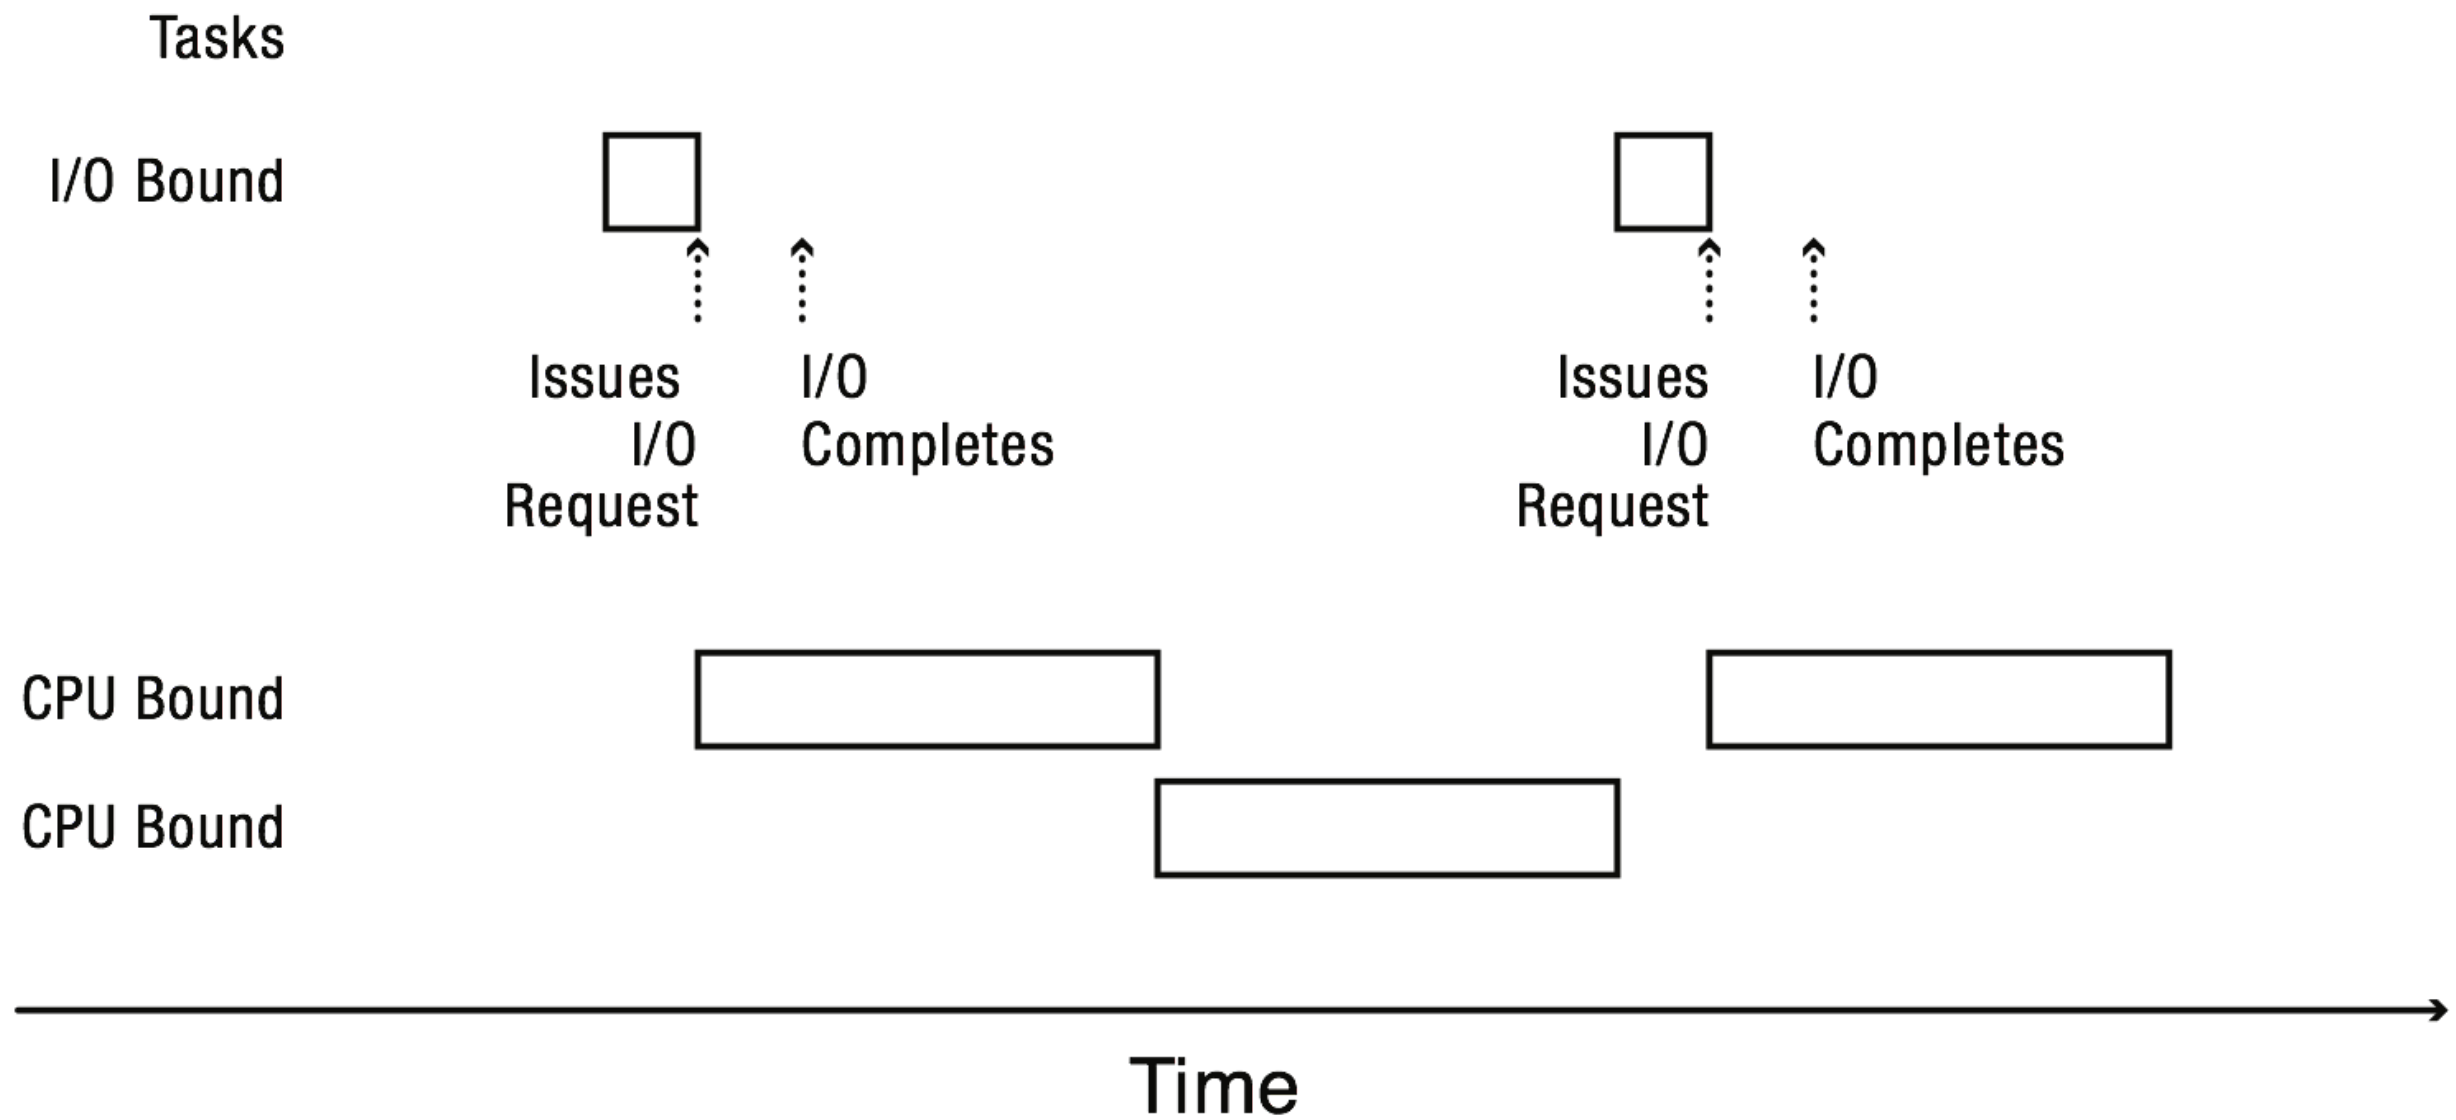
- Kid 1: 20%
 - Kid 2: 26%
 - Kid 3: 40%
 - Kid 4: 50%
-
- 100% -> 25% each kid
 - ~~20%~~
 - 26%
 - 27%
 - 27%

Max-Min Fairness



- How do we balance a mixture of repeating tasks?
 - Some I/O bound, need only a little CPU
 - Some compute bound, can use as much CPU as they are assigned
- One approach: maximize the minimum allocation given to a task
 - If any task needs less than an equal share, schedule the smallest of these first
 - Split the remaining time using max-min
 - If all remaining tasks need at least equal share, split evenly

Mixed Workloads??



Multi-Level Feedback Queue




- Set of Round Robin queues
 - Each queue has a separate priority
- High priority queues have short time slices
 - Low priority queues have long time slices
- Scheduler picks first thread in highest priority queue
- Tasks start in highest priority queue
 - If time slice expires, task drops one level

Multi-Level Feedback Queue



- **Goals:**
 - Responsiveness
 - Low overhead
 - Starvation freedom
 - Some tasks are high/low priority
 - Fairness (among equal priority tasks)
- **Not perfect at any of them!**
 - Used in Linux (and probably Windows, MacOS)

Multi-Level Feedback Queue



Why is MLFQ a good design?



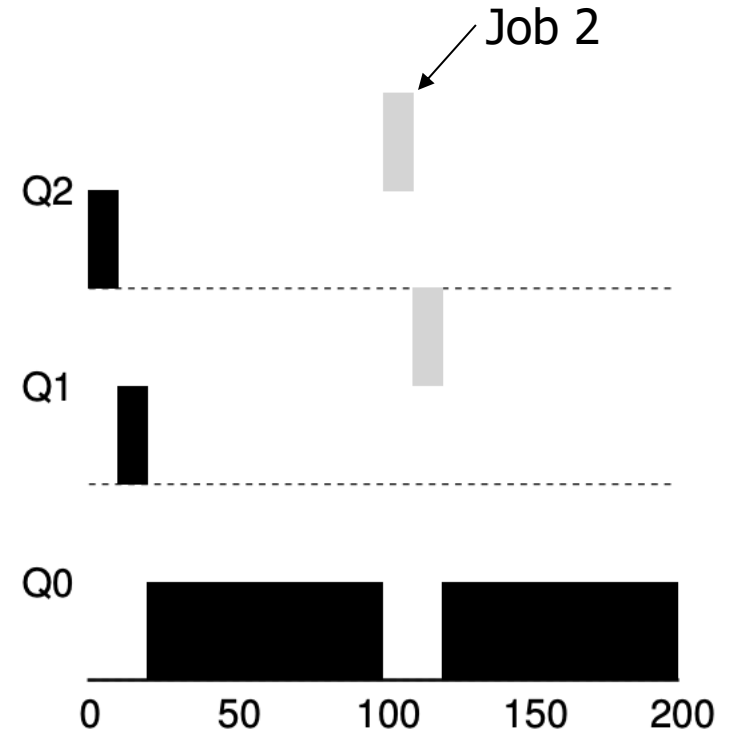

- How to design a scheduler that both minimizes response time for interactive jobs while also minimizing turnaround time without a priori knowledge of job length?
- Yes, SJF – the assumption is to know which is the “shortest..”
 - It’s just very hard to know in advance.
 - Sometimes processes/threads could try to game (we will see an example).


Why is MLFQ a good design?



- The Key Idea
 - Dynamically adjusting the priority level based on observing the behavior of the processes/threads
- Basic Design
 - When a job enters the system, it is placed at the highest priority (the topmost queue).
 - If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue).
 - If a job gives up the CPU before the time slice is up, it stays at the same priority level.

Basic Design







- because it doesn't know whether a job will be a short job or a long-running job, it first assumes it might be a short job, thus giving the job high priority. If it actually is a short job, it will run quickly and complete; if it is not a short job, it will slowly move down the queues, and thus soon prove itself to be a long-running more batch-like process.

Limitations?




- Starvation
- A process changing its characteristics
- Gaming the scheduler



Priority Boost




- After some time period S , move all the jobs in the system to the topmost queue



Better Accounting



- Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue).



Sounds perfect?



- How many queues should there be?
- How big should the time slice be per queue?
- How often should priority be boosted in order to avoid starvation and account for changes in behavior?

Summary



- FIFO is simple and minimizes overhead.
- If tasks are variable in size, then FIFO can have very poor average response time.
- If tasks are equal in size, FIFO is optimal in terms of average response time.
- Considering only the processor, SJF is optimal in terms of average response time.
- SJF is pessimal in terms of variance in response time.

Summary



- If tasks are variable in size, Round Robin approximates SJF.
- If tasks are equal in size, Round Robin will have very poor average response time.
- Tasks that intermix processor and I/O benefit from SJF and can do poorly under Round Robin.



- Max-Min fairness can improve response time for I/O-bound tasks.
- Round Robin and Max-Min fairness both avoid starvation.
- By manipulating the assignment of tasks to priority queues, an MFQ scheduler can achieve a balance between responsiveness, low overhead, and fairness.
- Is MFQ optimally fair??