

Safe Kernel Extensions

Guest Lecture for CS 423

Jinghao Jia

Tenured Grad Student

About me


- I was a Ph.D. student at UIUC working on safe kernel extensions
 - Advised by Professor Tianyin Xu from UIUC and Professor Dan Williams from Virginia Tech
 - Graduated this spring
- Have been a TA for CS423 for 3 semesters

Kernel extensibility is an essential OS capability

- Customizing kernel functionality for **diverse needs**
 - Supporting specialized OS features and policies
- No **complexity** added to core kernel code
 - Modularization produces maintainable code
- No need to perform **disruptive** kernel reboots
 - No effect on service availability


Traditional kernel extensibility: loadable modules

User




Kernel


Traditional kernel extensibility: loadable modules




Traditional kernel extensibility: loadable modules




Traditional kernel extensibility: loadable modules




Loadable kernel modules are inherently unsafe!




Loadable kernel modules are inherently unsafe!




Rust-based kernel modules does not help



Rust-based kernel modules does not help



Rust-based kernel modules does not help



eBPF extensions are taking over the OS kernel





Figure credit: Brendan Gregg

eBPF extensions are taking over the OS kernel

BMC: Accelerating Memcached using Safe In-kernel Caching and Pre-stack Processing

Yoann Ghigoff, *Orange Labs, Sorbonne Université, Inria, LIP6*; Julien Sopena,
Sorbonne Université, LIP6; Kahina Lazri, *Orange Labs*; Antoine Blin, *Gandi*;
Gilles Muller, *Inria*

XRP: In-Kernel Storage Functions with eBPF

Yuhong Zhong, Haoyu Li, Yu Jian Wu, Ioannis Zarkadas, Jeffrey Tao,
Evan Mesterhazy, Michael Makris, and Junfeng Yang, *Columbia University*;
Amy Tai, *Google*; Ryan Stutsman, *University of Utah*;
Asaf Cidon, *Columbia University*

DINT: Fast In-Kernel Distributed Transactions with eBPF

Yang Zhou, *Harvard University*; Xingyu Xiang, *Peking University*;
Matthew Kiley, *Harvard University*; Sowmya Dharamipragada, *Cornell University*;
Minlan Yu, *Harvard University*

Electrode: Accelerating Distributed Protocols with eBPF

Yang Zhou, *Harvard University*; Zezhou Wang, *Peking University*;
Sowmya Dharamipragada, *Cornell University*; Minlan Yu, *Harvard University*

eBPF extensions are taking over the OS kernel

BMC: Accelerating Memcached using Safe In-kernel Caching and Pre-stack Processing

Yoann Ghigoff, *Orange Labs, Sorbonne Université, Inria, LIP6*; Julien Sopena, *Sorbonne Université, LIP6*; Kahina Lazri, *Orange Labs*; Antoine Blin, *Gandi*; Gilles Muller, *Inria*

XRP: In-Kernel Storage Functions with eBPF

Yuhong Zhong, Haoyu Li, Yu Jian Wu, Ioannis Zarkadas, Jeffrey Tao, Evan Mesterhazy, Michael Makris, and Junfeng Yang, *Columbia University*; Amy Tai, *Google*; Ryan Stutsman, *University of Utah*; Asaf Cidon, *Columbia University*

DINT: Fast In-Kernel Distributed Transactions with eBPF

Yang Zhou, *Harvard University*; Xingyu Xiang, *Peking University*; Matthew Kiley, *Harvard University*; Sowmya Dharamipragada, *Cornell University*; Minlan Yu, *Harvard University*

Electrode: Accelerating Distributed Protocols with eBPF

Yang Zhou, *Harvard University*; Zezhou Wang, *Peking University*; Sowmya Dharamipragada, *Cornell University*; Minlan Yu, *Harvard University*



eBPF in CPU Scheduler


Hao Luo <haoluo@google.com>
Barret Rhoden <brho@google.com>

Towards Programmable Memory Management with eBPF

Presented by Kaiyang Zhao <kaiyang2@cs.cmu.edu>


Basic principle of eBPF: Safety verification

User




Kernel


Basic principle of eBPF: Safety verification



Basic principle of eBPF: Safety verification




Basic principle of eBPF: Safety verification




Symbolic execution on
all possible code paths


Basic principle of eBPF: Safety verification



Basic principle of eBPF: Safety verification




Basic principle of eBPF: Safety verification



Symbolic execution on
all possible code paths

Basic principle of eBPF: Safety verification



Safety at the cost of usability



```
/* cilium/bpf/bpf_sock.c */
int __sock4_post_bind(struct bpf_sock *ctx)
{
    struct lb4_service *svc;
    struct lb4_key key = ...;
    ...
    svc = __lb4_lookup_service(&key);
    if (!svc) {
        key.dport = bpf_htons(ctx->src_port);
        svc = sock4_nodeport_wildcard_lookup(&key, ...);
    }
    ...
}
```

Safety at the cost of usability



```
/* cilium/bpf/bpf_sock.c */
int __sock4_post_bind(struct bpf_sock *ctx)
{
    struct lb4_service *svc;
    struct lb4_key key = ...;
    ...
    svc = __lb4_lookup_service(&key);
    if (!svc) {
        key.dport = bpf_htons(ctx->src_port);
        svc = sock4_nodeport_wildcard_lookup(&key, ...);
    }
    ...
}
```

Look up the service
from a map using key

Safety at the cost of usability



```
/* cilium/bpf/bpf_sock.c */
int __sock4_post_bind(struct bpf_sock *ctx)
{
    struct lb4_service *svc;
    struct lb4_key key = ...;
    ...
    svc = __lb4_lookup_service(&key);
    if (!svc) {
        key.dport = bpf_htons(ctx->src_port);
        svc = sock4_nodeport_wildcard_lookup(&key,
    }
    ...
}
```

Look up the service
from a map using key

Redo a wildcard lookup if
not found in the last round

Safety at the cost of usability



```
/* cilium/bpf/bpf_sock.c */  
int __sock4_post_bind(struct bpf_sock *ctx)  
{  
    struct lb4_service *svc;  
    struct lb4 key key = ...;
```



This simple code does not pass the eBPF verifier!

```
    key.dport = bpf_htons(ctx->src_port);  
    svc = sock4_nodeport_wildcard_lookup(&key, ...);  
}  
...  
}
```

Safety at the cost of usability



```
32: (85) call bpf_map_lookup_elem#1
33: (15) if r0 == 0x0 goto pc+2"
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0)
; R6=ctx(id=0,off=0,imm=0)
34: (69) r1 = *(u16 *)(r0 +4)"
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0)
; R6=ctx(id=0,off=0,imm=0)
35: (55) if r1 != 0x0 goto pc+51
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0) R1=inv0
; R6=ctx(id=0,off=0,imm=0)
36: (69) r2 = *(u16 *)(r6 +44)
invalid bpf_context access off=44 size=2
```

Verifier log

Safety at the cost of usability



```
32: (85) call bpf_map_lookup_elem#1
33: (15) if r0 == 0x0 goto pc+2"
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0)
; R6=ctx(id=0,off=0,imm=0)
34: (69) r1 = *(u16 *)(r0 +4)"
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0)
; R6=ctx(id=0,off=0,imm=0)
35: (55) if r1 != 0x0 goto pc+51
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0) R1=inv0
; R6=ctx(id=0,off=0,imm=0)
36: (69) r2 = *(u16 *)(r6 +44)
invalid bpf_context access off=44 size=2
```

Which source-code line do these instructions map to?

Verifier log

Safety at the cost of usability



```
32: (85) call bpf_map_lookup_elem#1
33: (15) if r0 == 0x0 goto pc+2"
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0)
; R6=ctx(id=0,off=0,imm=0)
34: (69) r1 = *(u16 *)(r0 +4)"
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0)
; R6=ctx(id=0,off=0,imm=0)
35: (55) if r1 != 0x0 goto pc+51
; R0=map_value(id=0,off=0,ks=12,vs=12,imm=0) R1=inv0
; R6=ctx(id=0,off=0,imm=0)
36: (69) r2 = *(u16 *)(r6 +44)
invalid bpf_context access off=44 size=2
```

Which source-code line do these instructions map to?

Why is it an invalid access?

Verifier log

Root cause: verifier does not understand compiler

- The program uses `bpf_htons()` to convert the endianness of the `src_port` field in the context.
 - `src_port` is defined as a 32-bit int, while `bpf_htons()` only performs operations on the upper 16 bits
- The compiler optimizes the code to only load the upper 16 bits
- The verifier checks context field accesses based on its size
 - Expect a 32-bit load on `src_port`, but only sees a 16-bit load
 - Reject the extension program with size mismatch error

Workarounds


```
+static __always_inline __maybe_unused __be16
+ctx_src_port(struct bpf_sock *ctx)
+{
+. volatile __u32 sport = ctx->src_port;
+ return (__be16)bpf_htons(sport);
+}
+
...
if (!svc) {
- key.dport = bpf_htons(ctx->src_port);
+ key.dport = ctx_src_port(ctx);
    svc = sock4_nodeport_wildcard_lookup(&key, ...);
}
```

Workarounds

```
+static __always_inline __maybe_unused __be16  
+ctx_src_port(struct bpf_sock *ctx)  
{  
+. volatile __u32 sport = ctx->src_port;  
+ return (__be16)bpf_htons(sport);  
}  
+  
...  
if (!svc) {  
- key.dport = bpf_htons(ctx->src_port);  
+ key.dport = ctx_src_port(ctx);  
    svc = sock4_nodeport_wildcard_lookup(&key, ...);  
}
```


Used volatile to force a 32-bit load from the compiler

The language-verifier gap




Symbolic execution on
all possible code paths

The language-verifier gap




Symbolic execution on
all possible code paths

The language-verifier gap



The language-verifier gap



The language-verifier gap causes many problems

Workaround	Count
Refactoring extension code into smaller ones	27
Hinting compilers to generate verifier-friendly code	22
Tweaking code to assist verification	15
Dealing with verifier bugs	9
Reinventing the wheels	1



Closing the language-verifier gap

- Running kernel extensions *safely* without a verifier
 - Key challenge: **how to ensure safety?**

Closing the language-verifier gap

- Running kernel extensions *safely* without a verifier
 - Key challenge: **how to ensure safety?**
- Insight: **Language-based safety + runtime mechanism**
 - Rust as the safe language (safe Rust only)
 - Runtime safety checks for other safety properties
 - e.g., termination and stack safety

Rex: Safe, usable Rust kernel extensions

User

Kernel


Rex: Safe, usable Rust kernel extensions

User




Kernel

Rex: Safe, usable Rust kernel extensions




Rex: Safe, usable Rust kernel extensions


User
Kernel




Rex: Safe, usable Rust kernel extensions



Rex: Safe, usable Rust kernel extensions




Rex: Safe, usable Rust kernel extensions




Roadmap

- Rex kernel crate: Compile-time safety
 - Memory safety
 - Extended type safety
 - Safe resource management
- Rex Runtime safety
 - Safe exception handling
 - Kernel stack safety
 - Program Termination



Roadmap

- Rex kernel crate: Compile-time safety
 - Memory safety
 - Extended type safety
 - Safe resource management
- Rex Runtime safety
 - Safe exception handling
 - Kernel stack safety
 - Program Termination



Memory safety

- Rex enforces extensions to access kernel memory safely
 - Memory owned by the Rex extension
 - Safe Rust already ensures type safety (no unsafe memory access)
 - Safe kernel interface provided by the Rex kernel crate
 - Memory owned by the kernel
 - Kernel objects with static size are handled by Rust type system
 - Rust slices for dynamic pointers (runtime checks)
 - In principle similar to dynptrs in eBPF

Extended type safety

- There is the desire of transforming a byte stream to typed data
 - Common in networking use cases
 - **Inevitably requires unsafe code** under the Rust type system

Extended type safety

- There is the desire of transforming a byte stream to typed data
 - Common in networking use cases
 - **Inevitably requires unsafe code** under the Rust type system

```
void *payload = ctx->data;
/* extract ip header */
struct iphdr *ip = payload;
__u8 protocol = ip->protocol;
```

Extended type safety

- There is the desire of transforming a byte stream to typed data
 - Common in networking use cases
 - **Inevitably requires unsafe code** under the Rust type system

```
let payload: &[u8] = ctx.data;
// extract ip header
let ip = unsafe { &*(payload.as_ptr() as *const iphdr) };
let protocol = ip.protocol;
```



Extended type safety

- There is the desire of transforming a byte stream to typed data
 - Common in networking use cases
 - **Inevitably requires unsafe code** under the Rust type system
- eBPF considers the cast safe under two rules
 - #1 Not making a pointer by casting a scalar value
 - #2 New value fitting in the memory boundary
- Rex follows the same rules and extends the Rust type safety

Extended type safety

- Idea: the compiler has full knowledge of the types

Extended type safety

- Idea: the compiler has full knowledge of the types
- Create an **auto trait**: `rex::NoRef`
 - Default implemented on a type if all fields are `rex::NoRef`
 - Negative implementation on reference and pointer types

```
pub unsafe auto trait NoRef {}  
impl<T: ?Sized> !NoRef for &T {}  
impl<T: ?Sized> !NoRef for &mut T {}  
impl<T: ?Sized> !NoRef for *const T {}  
impl<T: ?Sized> !NoRef for *mut T {}
```

Extended type safety

- Idea: the compiler has full knowledge of the types
- Create an **auto trait**: `rex::NoRef`
 - Default implemented on a type if all fields are `rex::NoRef`
 - Negative implementation on reference and pointer types
 - Compiler auto-implements only on types w/o **any** refs/pointers

Extended type safety

- Idea: the compiler has full knowledge of the types
- Create an **auto trait**: `rex::NoRef`
 - Default implemented on a type if all fields are `rex::NoRef`
 - Negative implementation on reference and pointer types
 - Compiler auto-implements only on types w/o **any** refs/pointers

```
pub struct my_hdr {  
    request_id: u16,  
    seq_num: u16,  
    num_dgram: u16,  
    other: NonNull<u64>,  
}
```

Extended type safety

- Idea: the compiler has full knowledge of the types
- Create an **auto trait**: `rex::NoRef`
 - Default implemented on a type if all fields are `rex::NoRef`
 - Negative implementation on reference and pointer types
 - Compiler auto-implements only on types w/o **any** refs/pointers

```
pub struct my_hdr {  
    request_id: u16,  
    seq_num: u16,  
    num_dgram: u16,  
    other: NonNull<u64>,  
}
```

NonNull contains a *const u64

Extended type safety

- Idea: the compiler has full knowledge of the types
- Create an **auto trait**: `rex::NoRef`
 - Default implemented on a type if all fields are `rex::NoRef`
 - Negative implementation on reference and pointer types
 - Compiler auto-implements only on types w/o **any** refs/pointers

```
pub struct my_hdr {  
    request_id: u16,  
    seq_num: u16,  
    num_dgram: u16,  
    other: NonNull<u64>,  
}
```

NonNull contains a `*const u64`

`Rex::NoRef` is not implemented
for `NonNull` and `my_hdr`

Extended type safety

- Idea: the compiler has full knowledge of the types
- Create an **auto trait**: `rex::NoRef`
 - Default implemented on a type if all fields are `rex::NoRef`
 - Negative implementation on reference and pointer types
 - Compiler auto-implements only on types w/o **any** refs/pointers
- Require type to be `rex::NoRef` to convert from bytes (#1)
- Check memory bounds at runtime (#2)

```
fn from_bytes<T: NoRef>(bytes: &[u8]) -> &T {  
    assert!(data.len() >= mem::size_of::<T>());  
    ...  
}
```

Safe resource management

- Rex leverages Resource-Acquisition-Is-Initialization (RAII)
 - Ties lifetime of acquirable resources with Rust wrapper types
 - **Acquire** resource in object constructor
 - e.g., calling `spin_lock()` constructs a `SpinlockGuard` object
 - **Release** resource in the destructor (via `Drop` trait in Rust)
 - e.g. `SpinlockGuard` implements the `Drop` trait and release the lock
- Compiler inserts a drop call when objects go out of scope
 - No need to **explicitly** manage resources

Safe resource management

- Certain APIs from the core library interferes with RAII

Safe resource management

- Certain APIs from the core library interferes with RAII
 - `core::mem::forget`
 - `core::mem::ManuallyDrop`
 - `core::intrinsics::abort`

Safe resource management

- Certain APIs from the core library interferes with RAII
 - `core::mem::forget`
 - `core::mem::ManuallyDrop`
 - `core::intrinsics::abort`
- 
- Prevents execution of destructors**

Safe resource management


- Certain APIs from the core library interferes with RAII
 - `core::mem::forget`
 - `core::mem::ManuallyDrop`
 - `core::intrinsics::abort`
-  **Prevents execution of destructors**
- ↳ Additionally **crashes** the kernel with an illegal instruction

Safe resource management

- Certain APIs from the core library interferes with RAII
 - `core::mem::forget`
 - `core::mem::ManuallyDrop`
 - `core::intrinsics::abort`
 - Rex forbids the usage of such language items
 - Reject via the disallowed methods/types linter configuration
- 
- Prevents execution of destructors**
- ↳ Additionally **crashes** the kernel with an illegal instruction

Roadmap

- Rex kernel crate: Compile-time safety
 - Memory safety
 - Extended type safety
 - Safe resource management
- Rex Runtime safety
 - Safe exception handling
 - Kernel stack safety
 - Program Termination



Safe exception handling

- Certain safety properties of Rust are checked at runtime
 - Violations (e.g., out-of-bound access) trigger **Rust panics**
- Rust performs Itanium-ABI exception handling in user space
 - Unwind each stack frame and executes cleanup code
 - Not suitable in context of safe kernel extensions
 - Stack unwinding **cannot** fail (causing kernel crash or resource leaks)
 - Executing user-defined destructors is **not safe**
- Rex supports safe exception handling with two components
 - **Graceful exit**
 - **Resource cleanup**

Safe exception handling (graceful exit)

```
// In-kernel dispatcher function
rex_dispatcher_func:
// save old stack pointer
mov %rsp, PER_CPU_VAR(rex_old_sp)
...
// invoke the REX program
call *%rdx

rex_exit:
// reset to old stack pointer
mov PER_CPU_VAR(rex_old_sp), %rsp
...
ret
```

Safe exception handling (graceful exit)


```
// In-kernel dispatcher function
rex_dispatcher_func:
// save old stack pointer
mov %rsp, PER_CPU_VAR(rex_old_sp)
...
// invoke the REX program
call *%rdx

rex_exit:
// reset to old stack pointer
mov PER_CPU_VAR(rex_old_sp), %rsp
...
ret
```

Save the old stack pointer
before program invocation

Safe exception handling (graceful exit)

```
// In-kernel dispatcher function
rex_dispatcher_func:
// save old stack pointer
mov %rsp, PER_CPU_VAR(rex_old_sp)
...
// invoke the REX program
call *%rdx .....  
  
rex_exit:
// reset to old stack pointer
mov PER_CPU_VAR(rex_old_sp), %rsp
...
ret
```



The diagram illustrates the flow of control between two components. On the left, a large blue-bordered box contains assembly code for an 'In-kernel dispatcher function'. On the right, a smaller blue-bordered box contains assembly code for a 'REX program'. A green dotted arrow points from the end of the dispatcher's 'call *%rdx' instruction to the start of the REX program's code. This visualizes how the dispatcher invokes the REX program and then exits back to its caller.

```
// Rex program
rex_prog1:
...
```


Safe exception handling (graceful exit)

```
// In-kernel dispatcher function
rex_dispatcher_func:
// save old stack pointer
mov %rsp, PER_CPU_VAR(rex_old_sp)
...
// invoke the REX program
call *%rdx .....  
  
rex_exit:
// reset to old stack pointer
mov PER_CPU_VAR(rex_old_sp), %rsp
...
ret
```



Safe exception handling (graceful exit)

```
// In-kernel dispatcher function
rex_dispatcher_func:
// save old stack pointer
mov %rsp, PER_CPU_VAR(rex_old_sp)
...
// invoke the REX program
call *%rdx .....  
  
rex_exit:
// reset to old stack pointer
mov PER_CPU_VAR(rex_old_sp), %rsp
...
ret
```



Safe exception handling (graceful exit)

```
// In-kernel dispatcher function  
rex_dispatcher_func:  
// save old stack pointer  
mov %rsp, PER_CPU_VAR(rex_old_sp)  
...  
  
// invoke the REX program  
call *%rdx .....  
  
rex_exit:  
// reset to old stack pointer  
mov PER_CPU_VAR(rex_old_sp), %rsp  
...  
ret
```



Safe exception handling (graceful exit)

```
// In-kernel dispatcher function  
rex_dispatcher_func:  
// save old stack pointer  
mov %rsp, PER_CPU_VAR(rex_old_sp)  
...  
  
// invoke the REX program  
call *%rdx .....  
  
rex_exit: <-----  
// reset to old stack pointer  
mov PER_CPU_VAR(rex_old_sp), %rsp  
...  
ret
```



Safe exception handling (graceful exit)

```
// In-kernel dispatcher function
rex_dispatcher_func:
// save old stack pointer
mov %rsp, PER_CPU_VAR(rex_old_sp)
...
// invoke the REX program
call *%rdx
...
rex_exit: <----- // reset to old stack pointer
mov PER_CPU_VAR(rex_old_sp), %rsp
...
ret
```



Restore the old stack
pointer to unwind stack

jmp rex_exit

Safe exception handling (graceful exit)

- Integrate with kernel stack trace for better user experience
 - Dump stack trace of Rex panics to the kernel console
 - De-mangle Rex symbols and register them into kallsyms on load
 - Enable frame-pointers for the kernel stack unwinder

Safe exception handling (graceful exit)

- Integrate with kernel stack trace for better user experience
 - Dump stack trace of Rex panics to the kernel console
 - De-mangle Rex symbols and register them into kallsyms on load
 - Enable frame-pointers for the kernel stack unwinder

```
Panic from Rex prog: called `Option::unwrap()` on a `None` value
```

```
Call Trace:
```

```
rex_landingpad+0x64/0xb0
rex_prog_4168211f00000000::rust_begin_unwind+0x15a/0x1a0
rex_prog_4168211f00000000::core::panicking::panic_fmt+0x9/0x10
rex_prog_4168211f00000000::core::panicking::panic+0x53/0x60
rex_prog_4168211f00000000::core::option::unwrap_failed+0x9/0x10
rex_prog_4168211f00000000::err_injector+0x8f/0xa0
rex_dispatcher_func+0x32/0x32
```

Safe exception handling (resource cleanup)

- Safe exception handling **must clean up** acquired resources
 - Itanium EH uses DWARF to identify acquired resources
- Insight: **extensions can only acquire resources via helpers**
 - Only these resources need to be released
 - Record resources allocated via helper functions in a per-CPU buffer
 - Upon panic, iterate the resources and perform cleanup

Kernel stack safety

- Safe kernel extension should **never overflow the stack**
 - Unlike user space, kernel stack does not grow
 - Stack check from the eBPF verifier can be tricked via BPF tail calls [1]
- Rex divides stack protection into two cases
 - Extension programs that have no indirect/recursive calls
 - Static call graph is a DAG
 - Compute total stack usage statically based on the call graph
 - Extension programs that have indirect/recursive calls
 - Check stack usage dynamically before each function call
 - A failed check triggers a Rust panic and will be gracefully handled

Safe termination

- A non-terminating extension could hold CPU for a long time
 - Prevents other tasks/interrupts from being executed
- Rex uses a dynamic timeout mechanism for termination
 - Set up an hrtimer per CPU that triggers periodically
 - Timer interrupt (hardirq) can preempt programs in softirq or task
 - If timeout, set program instruction pointer to the timeout handler
- **Defer** termination if the program is in a helper call
 - Preventing inconsistent states (e.g., reference count, lock)

Implementation

- Rex is supported on the latest stable Linux kernel and Rust
 - Linux-6.17 + Rust-1.91.0
- Rex Kernel crate provides the kernel interface
 - Kernel symbol bindings are implemented manually
 - Emitted as relocations in the compiled PIE binary
 - Kernel data-type bindings are automated by bindgen
 - Structs, constants, and type aliases
 - Program-type-specific code
 - Currently supports kprobe, perf-event, tracepoint, xdp, and tc

Implementation

- Kernel support
 - Program loading and symbol resolution
 - Map the program binary into the kernel address space
 - Patch GOT relocations of referenced kernel symbols
 - In-kernel Rex runtime
 - Stack unwinder and termination timer handler
- Compiler support
 - Rex pass in LLVM backend of rustc
 - Perform Rex-specific instrumentations (e.g., stack checks)
 - Generate "extern C" program entry points (actively moving to proc-macros)

Usability evaluation

- Heuristic evaluation
 - No language-verifier gap anymore
- Dogfooding
 - Used Rex to implement the BPF Memcached Cache (BMC)
 - Much cleaner, simpler code
 - 326 lines of Rust code vs. 513 lines of C code

Case study: BMC cache invalidation

Original eBPF-BMC

```
// Searches for SET command in payload
for (unsigned int off = 0;
    off < BMC_MAX_PACKET_LENGTH &&
    payload + off + 1 <= data_end;
    off++) {
    if (set_found == 0 &&
        payload[off] == 's' &&
        payload + off + 3 <= data_end &&
        payload[off + 1] == 'e' &&
        payload[off + 2] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

Case study: BMC cache invalidation

Original eBPF-BMC

```
// Searches for SET command in payload
for (unsigned int off = 0;
    off < BMC_MAX_PACKET_LENGTH &&
    payload + off + 1 <= data_end;
    off++) {
    if (set_found == 0 &&
        payload[off] == 's' &&
        payload + off + 3 <= data_end &&
        payload[off + 1] == 'e' &&
        payload[off + 2] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

Additional limit to fit in verifier's complexity limit

Case study: BMC cache invalidation

Original eBPF-BMC

```
// Searches for SET command in payload
for (unsigned int off = 0;
     off < BMC_MAX_PACKET_LENGTH &&
     payload + off + 1 <= data_end;
     off++) {
    if (set_found == 0 &&
        payload[off] == 's' &&
        payload + off + 3 <= data_end &&
        payload[off + 1] == 'e' &&
        payload[off + 2] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

Additional limit to fit in verifier's complexity limit

Boilerplate to explicitly check end of payload

Case study: BMC cache invalidation

Original eBPF-BMC

```
// Searches for SET command in payload
for (unsigned int off = 0;
     off < BMC_MAX_PACKET_LENGTH &&
     payload + off + 1 <= data_end;
     off++) {
    if (set_found == 0 &&
        payload[off] == 's' &&
        payload + off + 3 <= data_end &&
        payload[off + 1] == 'e' &&
        payload[off + 2] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

Additional limit to fit in verifier's complexity limit

Boilerplate to explicitly check end of payload

Verbose logic to match SET command in payload

Case study: BMC cache invalidation

Original eBPF-BMC

```
// Searches for SET command in payload
for (unsigned int off = 0;
    off < BMC_MAX_PACKET_LENGTH &&
    payload + off + 1 <= data_end;
    off++) {
    if (set_found == 0 &&
        payload[off] == 's' &&
        payload + off + 3 <= data_end &&
        payload[off + 1] == 'e' &&
        payload[off + 2] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

Rex-BMC

```
// Searches for SET command in payload
let set_iter = payload
    .windows(4)
    .enumerate()
    .filter_map(|(i, v)|
        if v == b"set " {
            Some(i)
        } else {
            None
        }
    );
    ...
}
```

Case study: BMC cache invalidation

Original eBPF-BMC

```
// ...
for {
    No extra code to handle complexity limit from verifier
    payload + off + 1 <= data_end;
    off++) {
    if (set_found == 0 &&
        payload[off] == 's' &&
        payload + off + 3 <= data_end &&
        payload[off + 1] == 'e' &&
        payload[off + 2] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

Rex-BMC

```
// Searches for SET command in payload
let set_iter = payload
    .windows(4)
    .enumerate()
    .filter_map(|(i, v)|
        if v == b"set " {
            Some(i)
        } else {
            None
        }
    );
    ...
}
```

Case study: BMC cache invalidation

Original eBPF-BMC

```
// ...
for {
    payload + off + 1 <= data_end;
    ...
} {
    if (payload[off] == 'e' &&
        payload[off + 1] == 't') {
        off += 3;
        set_found = 1;
    }
    ...
}
```

No extra code to handle complexity limit from verifier

Rust slices already provides bound checks in its methods

Rex-BMC

```
// Searches for SET command in payload
let set_iter = payload
    .windows(4)
    .enumerate()
    .filter_map(|(i, v)|
        if v == b"set " {
            Some(i)
        } else {
            None
        }
    );
    ...
}
```

Case study: BMC cache invalidation

Original eBPF-BMC

```
// ...  
for (off = 0; off < data_end; off += 4) {  
    if (payload[off] == 'e' &&  
        payload[off + 1] == 't' &&  
        payload[off + 2] == 's' &&  
        payload[off + 3] == 't') {  
        // No extra code to handle  
        // complexity limit from verifier  
        if (payload + off + 1 <= data_end;  
            && payload[off + 1] != '\0') {  
            // Rust slices already provides  
            // bound checks in its methods  
            payload[off + 1] = '\0';  
            payload[off + 2] = '\0';  
            payload[off + 3] = '\0';  
            payload[off + 4] = '\0';  
            // Easy match of SET command  
            // using Rust byte slice  
            if (payload[off + 1] == 'e' &&  
                payload[off + 2] == 't' &&  
                payload[off + 3] == 's' &&  
                payload[off + 4] == 't') {  
                // ...  
            }  
        }  
    }  
}
```

Rex-BMC

```
// Searches for SET command in payload  
let set_iter = payload  
.windows(4)  
.enumerate()  
.filter_map(|(i, v)|  
    if v == b"set " {  
        Some(i)  
    } else {  
        None  
    }  
);  
...
```

Case study: BMC cache invalidation

Original eBPF-BMC


```
// ...  
for (off = 0; off < data_end; off += 4) {  
    if (payload[off] == 'e' &&  
        payload[off + 1] == 't' &&  
        payload[off + 2] == 's' &&  
        payload[off + 3] == 't') {  
        // No extra code to handle  
        // complexity limit from verifier  
        if (payload + off + 1 <= data_end;  
            && payload[off + 1] != '\0') {  
            // Rust slices already provides  
            // bound checks in its methods  
            payload[off + 1] == 'e' &&  
            payload[off + 2] == 't' &&  
            payload[off + 3] == 's') {  
                // Easy match of SET command  
                // using Rust byte slice  
                if (payload[off + 4] == ' ') {  
                    // ...  
                }  
            }  
        }  
    }  
}
```

Rex-BMC


```
// Searches for SET command in payload  
let set_iter = payload  
.windows(4)  
.enumerate()  
.filter_map(|(i, v)|  
    if v == b"set " {  
        Some(i)  
    } else {  
        None  
    }  
);  
...
```

Rex-BMC is much *cleaner* and *simpler*


Performance evaluation




Performance evaluation




Performance evaluation



Performance evaluation



Performance evaluation



Conclusion

- Static verification in kernel extensions leads to poor usability
- Insight of Rex: Language-based safety + runtime mechanism
- Rex delivers usability without losing safety and performance
- Code: <https://github.com/rex-rs/rex>

