



UNIVERSITY OF  
**ILLINOIS**  
URBANA-CHAMPAIGN

# MP2: Rate-Monotonic CPU Scheduling

Peizhe Liu

CS 423: Operating System Design

Fall 2025

## Important Dates

- **MP2 is released today.**
- **The due date is 10/28.**
- **MP2 will likely cost you several days (probably more than MP1).**
- **NO LATE SUBMISSION.** We found that GitHub allows it, but we won't.

## Goals

- **Learn the basics of real-time CPU scheduling.**
- **Develop a Rate-Monotonic Scheduler for Linux.**
- **Utilize more kernel APIs, like scheduler, slab allocator, kthread.**
- **Test your scheduler with a user app.**


## What you need

- Your MP0 environment.
- VSCode+clangd setup (strongly recommended).
- Instructions on the course website.
- Accept your assignment on GitHub classroom and start right away.



## Periodic Tasks Model

- Real time systems often require response time and predictability.
  - Surveillance camera that captures a frame per 30ms
- Liu and Layland Periodic Task Model to provide that timing guarantee.
- Every tasks carries a processing time and period.
- Every period, the task must run once, and must finish before the next period starts.



## Rate-Monotonic Scheduler


- An algorithm to implement the Periodic Tasks Model (methods VS rules).
- Static priority: the shorter the period, the higher the priority.
- Preemptive: higher priority task will preempt lower priority task.
- Utilization bound:  $\sum_{i \in T} \frac{C_i}{P_i} \leq 0.693$ 
  - T: all tasks in the system, including the incoming one
  - C: processing time.
  - P: period.

## MP2 Overview

- **Kernel module: implements RMS.**
  - Communicates with your user application via proc file write.
  - Prints list of applications via proc file read.
  - Handles the real scheduling via Linux scheduler APIs.
- **User application: simulates a RT application for your RMS.**
  - Registers itself to your kernel module.
  - May run multiple instances.

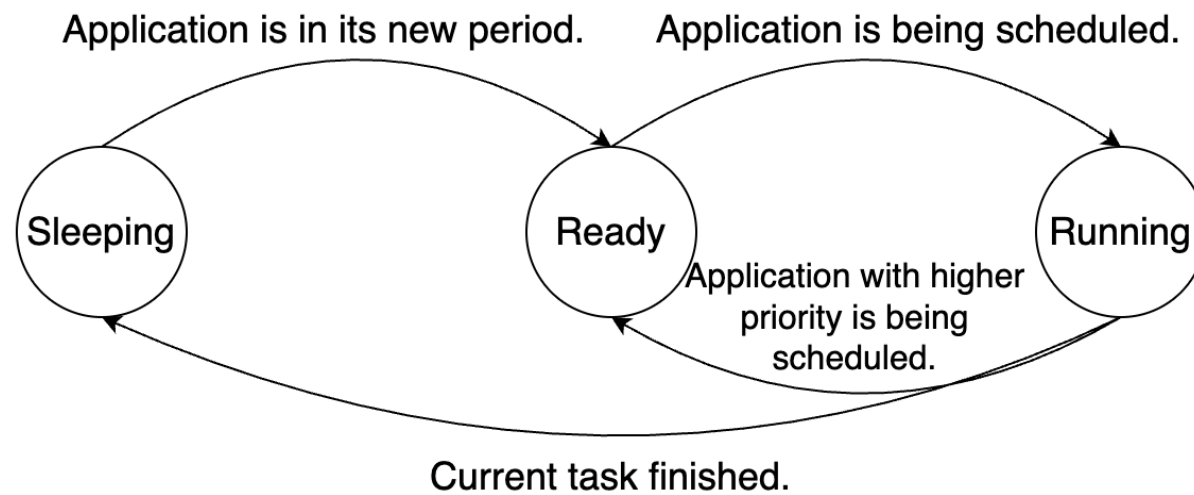
## RMS States

- **Ready:** application is currently in its new period. Ready to be scheduled.
- **Running:** application is currently running.
- **Sleeping:** application has finished running in its current period.



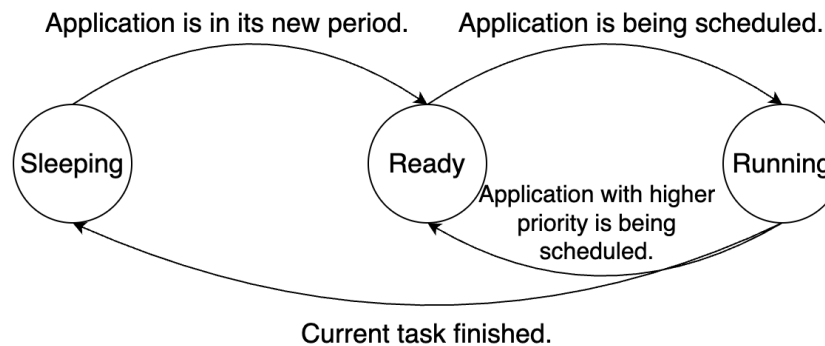


## State Transitions



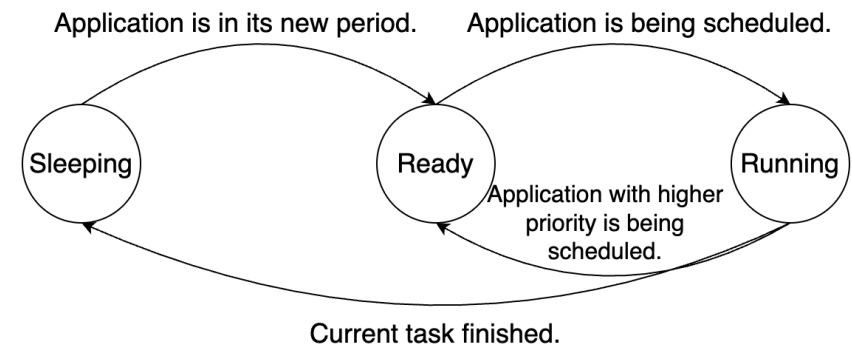
## Implement RMS

- New period: kernel timer
- Being scheduled/preempt: dispatcher kthread
- Task finished: userapp yield via proc file write



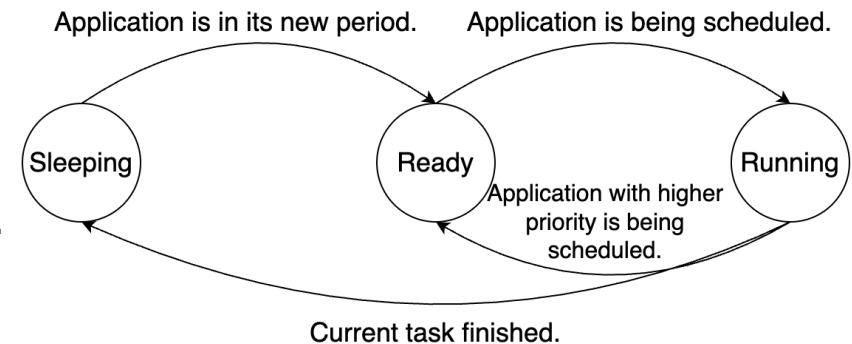
## Kernel Timer

- **Every application got its own timer.**
  - Set to its period time.
- **When timer fired:**
  - State transits from **SLEEPING** to **READY**.
  - Wake up the dispatcher kthread.
  - *Technically, it cannot fire at **READY** or **RUNNING** state...*



## Dispatcher KThread

- **Wake up by timer or yield.**
  - One app got its new period, or app finished running.
  - i.e.: Ready for schedule new application.
- **From applications in Ready state, pick the lowest period one and run it.**
  - Implements priority property of RMS.
  - State transits from READY to RUNNING.
- **Preempt the running one if necessary.**
  - Implements preemption property of RMS.
  - State transites from RUNNING to READY.



## Work with Linux Scheduler

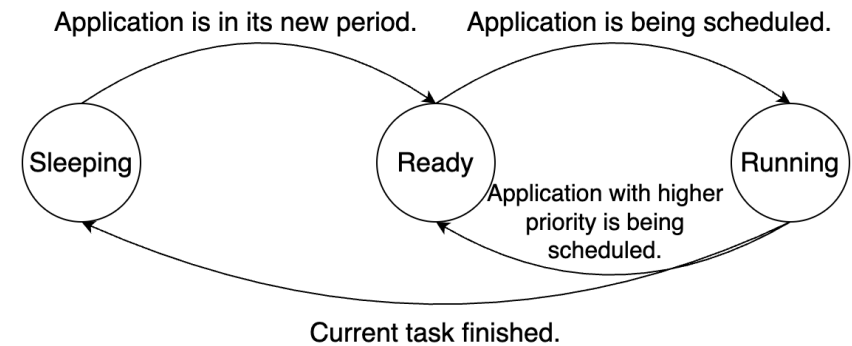
- "Bypass" scheduler, to put an application in sleep or wake it up immediately.
- Utilize kernel scheduling policy...

```
struct sched_attr attr;  
attr.sched_policy = SCHED_NORMAL;  
attr.sched_priority = 0;  
sched_setattr_nocheck(task, &attr);
```

```
#include <uapi/linux/sched/types.h>  
  
struct sched_attr attr;  
wake_up_process(task);  
attr.sched_policy = SCHED_FIFO;  
attr.sched_priority = 99;  
sched_setattr_nocheck(task, &attr);
```

## UserApp Yield

- **UserApp will yield when its task was done.**
  - This is done via `proc file write`.
- **When UserApp yields:**
  - State transits from **RUNNING** to **SLEEPING**.
  - Wake up the dispatcher kthread.
  - *Also technically, it cannot yield at SLEEPING or READY state...*



## That's MP2 RMS!

- Old APIs: timer, proc file write.
- New APIs: kthread, scheduler.
- We directly give you the code how to work with scheduler.

## Proc File Write

- **Register:** RT application registers itself with its PID, Period, Processing Time.

R,PID,PERIOD,COMPUTATION

- **Yield:** RT application has finished its period, yield the CPU.

Y,PID

- **De-register:** RT application has finished and de-registers itself.

D,PID



## Proc File Read

- **Print all registered applications with its PID, period, and process time.**
- **Reuse your MP1 code, certainly.**

```
<pid 1>: <period 1>, <processing time 1>  
<pid 2>: <period 2>, <processing time 2>  
...  
<pid n>: <period n>, <processing time n>
```

## Register and De-Register

- **Data structure: custom PCB and linked list.**
  - Locking can be pretty challenging. Let's see...
  - Remember, clean up after you use.
- **Utilization bound-based admission control:**  $\sum_{i \in T} \frac{C_i}{P_i} \leq 0.693$ 
  - Caution! FP arithmetic is very expensive in kernel!
  - Kernel does not save FP registers during context switch.

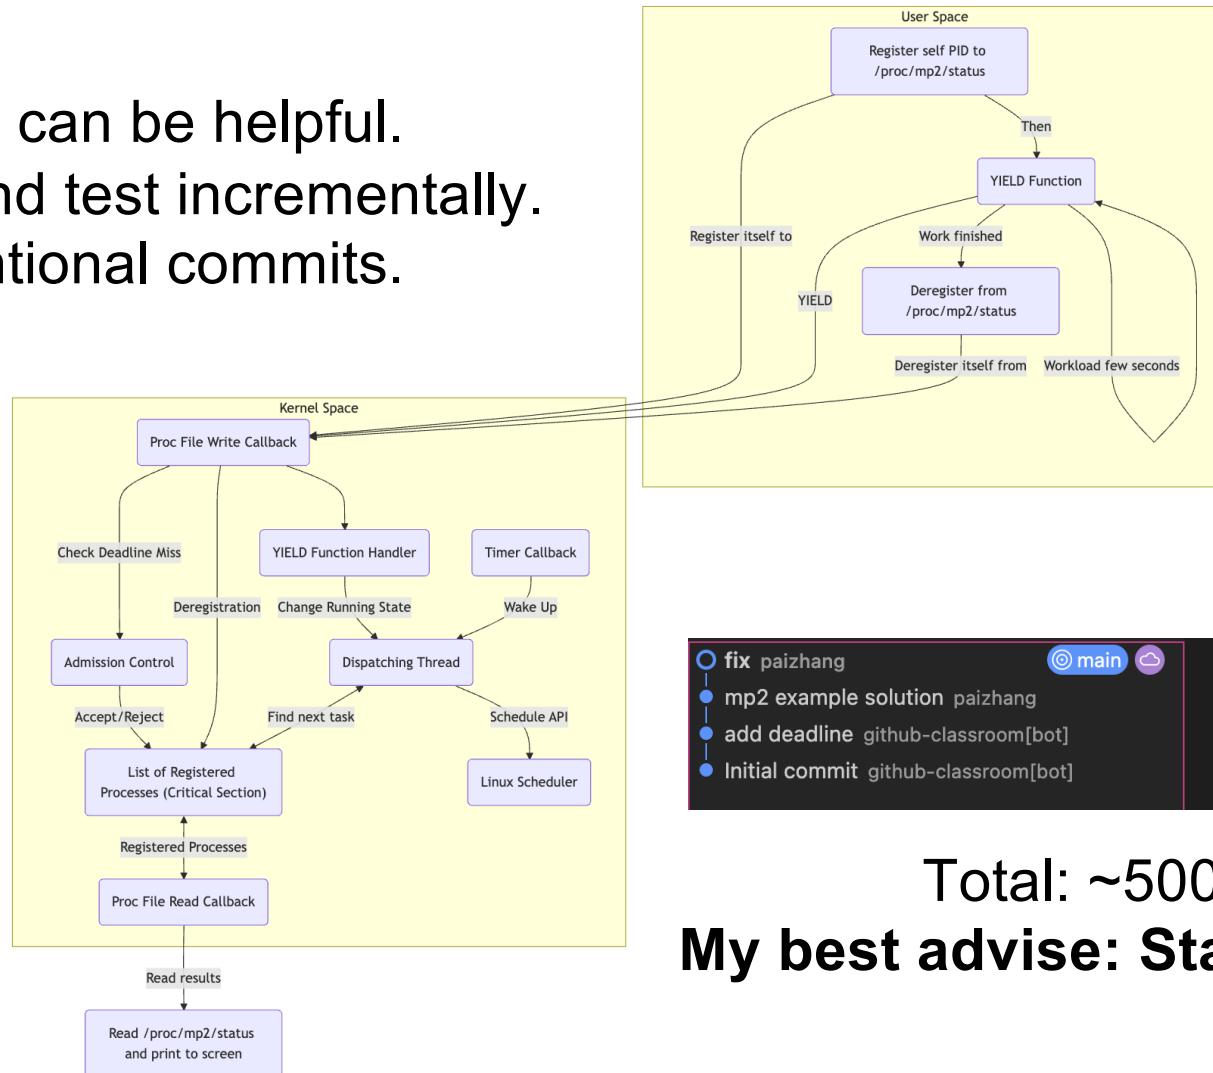
## User Application

- Register itself to the scheduler (double check to ensure it was registered).
- Yield immediately (signal the scheduler it is ready).
- Work loop:
  - Do some “real time work”
  - Yield
- Processing time should match with actual time.

```
void main(void)
{
    // Interact with Proc filesystem
    REGISTER(pid, period, processing_time);
    // Read ProcFS: Verify the process was admitted
    list = READ(ProcFS);
    if (!process in the list) exit(1);
    // setup everything needed for RT loop
    t0 = clock_gettime();
    // Proc filesystem
    YIELD(PID);
    // this is the real-time loop
    while (exist jobs)
    {
        wakeup_time = clock_gettime() - t0;
        // factorial computation
        do_job();
        process_time = clock_gettime() - wakeup_time;
        YIELD(PID);
    }
    // Interact with ProcFS
    DEREGISTER(PID);
}
```

Pro tips:

1. Good tools can be helpful.
2. Develop and test incrementally.
3. Try conventional commits.



```
fix paizhang
mp2 example solution paizhang
add deadline github-classroom[bot]
Initial commit github-classroom[bot]
```

```
507
508 module_init(mp2_init);
509 module_exit(mp2_exit);
510
```

Total: ~500 lines  
**My best advise: Start right away!**