



UNIVERSITY OF
ILLINOIS
URBANA-CHAMPAIGN

MP3: Virtual Memory Page Fault Profiler

Peizhe Liu

CS 423: Operating System Design

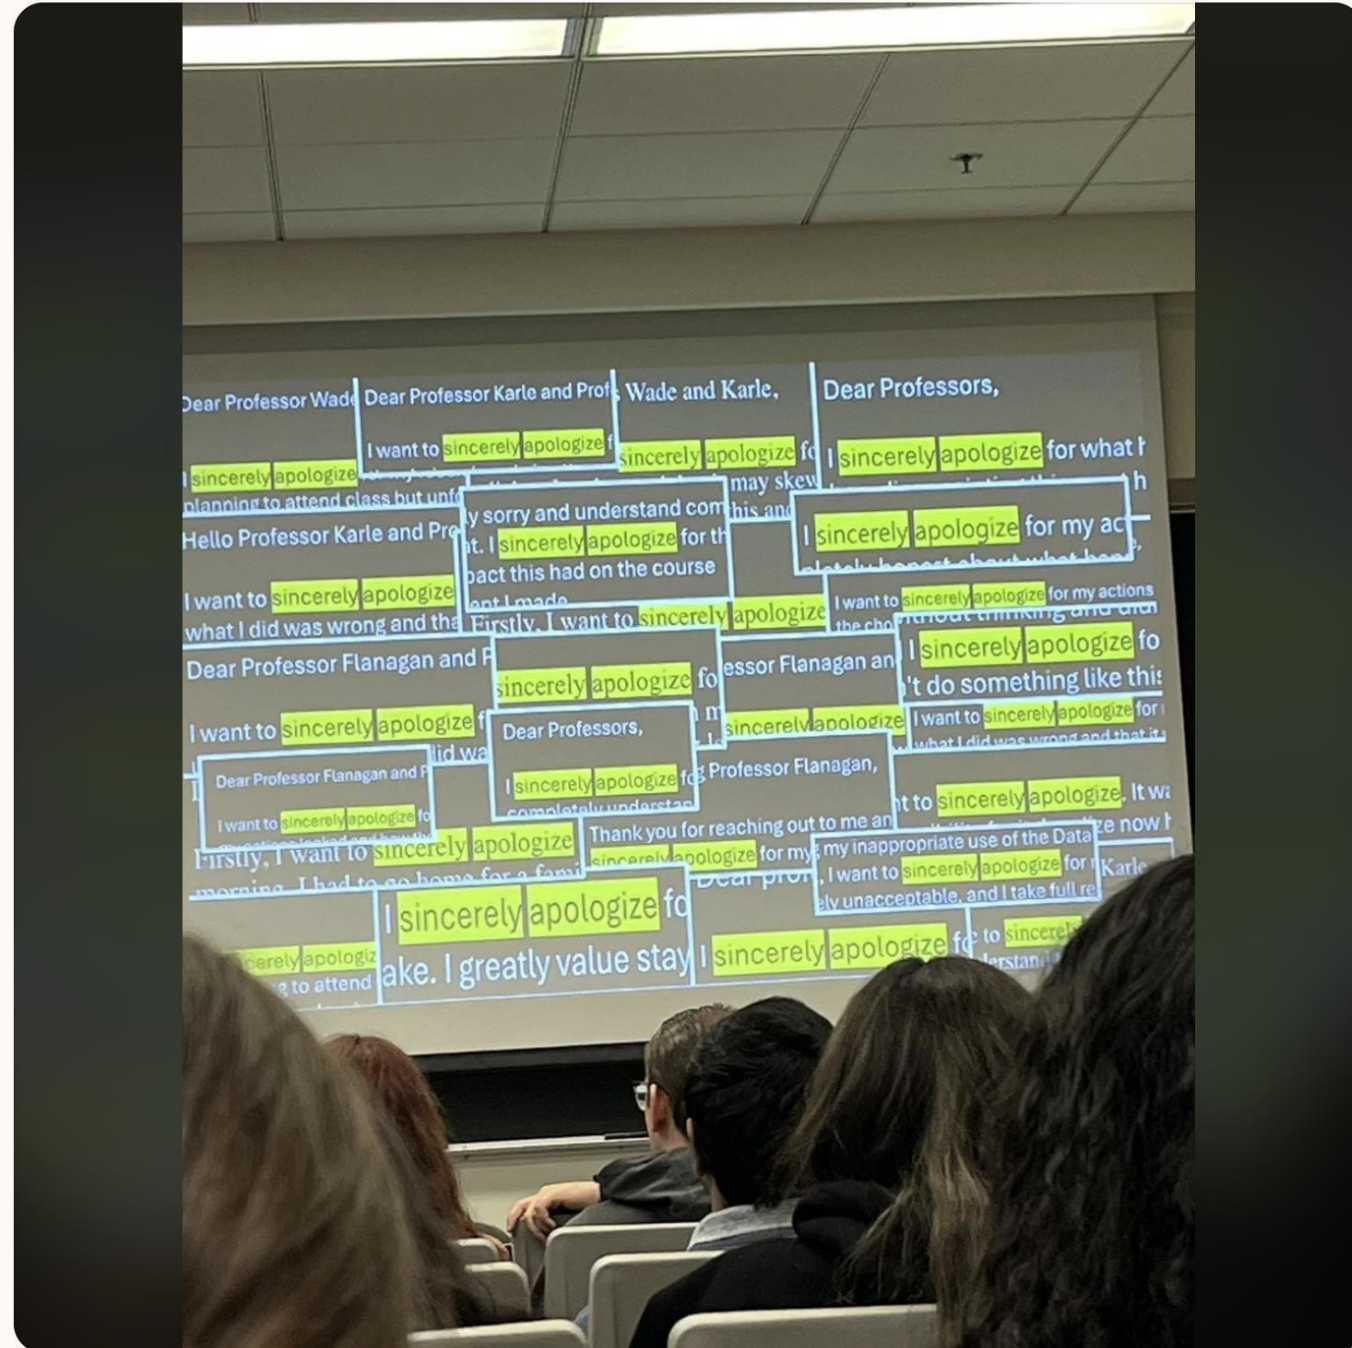
Fall 2025

Important Dates

- **MP3 is released today.**
- **The due date is 11/20.**
- **MP3 will likely cost you several days.**
- **NO LATE SUBMISSION.** We found that GitHub allows it, but we won't.

Important Reminders

- You must complete this MP individually.
- Copy & Paste from whatever AI is a severe academic integrity violation. (Student Code 1-402)



Goals

- **Understand Linux paging management.**
- **Develop a kernel PF profiler tool.**
- **Work with your profile tool to analyze PF and its effect on CPU usage.**
- **Utilize more kernel APIs, like char device, vmalloc, paging.**

What You Need

- Your MP0 environment.
- VSCode+clangd setup (strongly recommended).
- Instructions on the course website.
 - <https://github.com/cs423-uiuc/cs423-uiuc.github.io/blob/master/fall25/mps/MP3.md>
- Accept your assignment on GitHub classroom and start right away.
 - <https://classroom.github.com/a/V5g52O1A>



Major and Minor PF

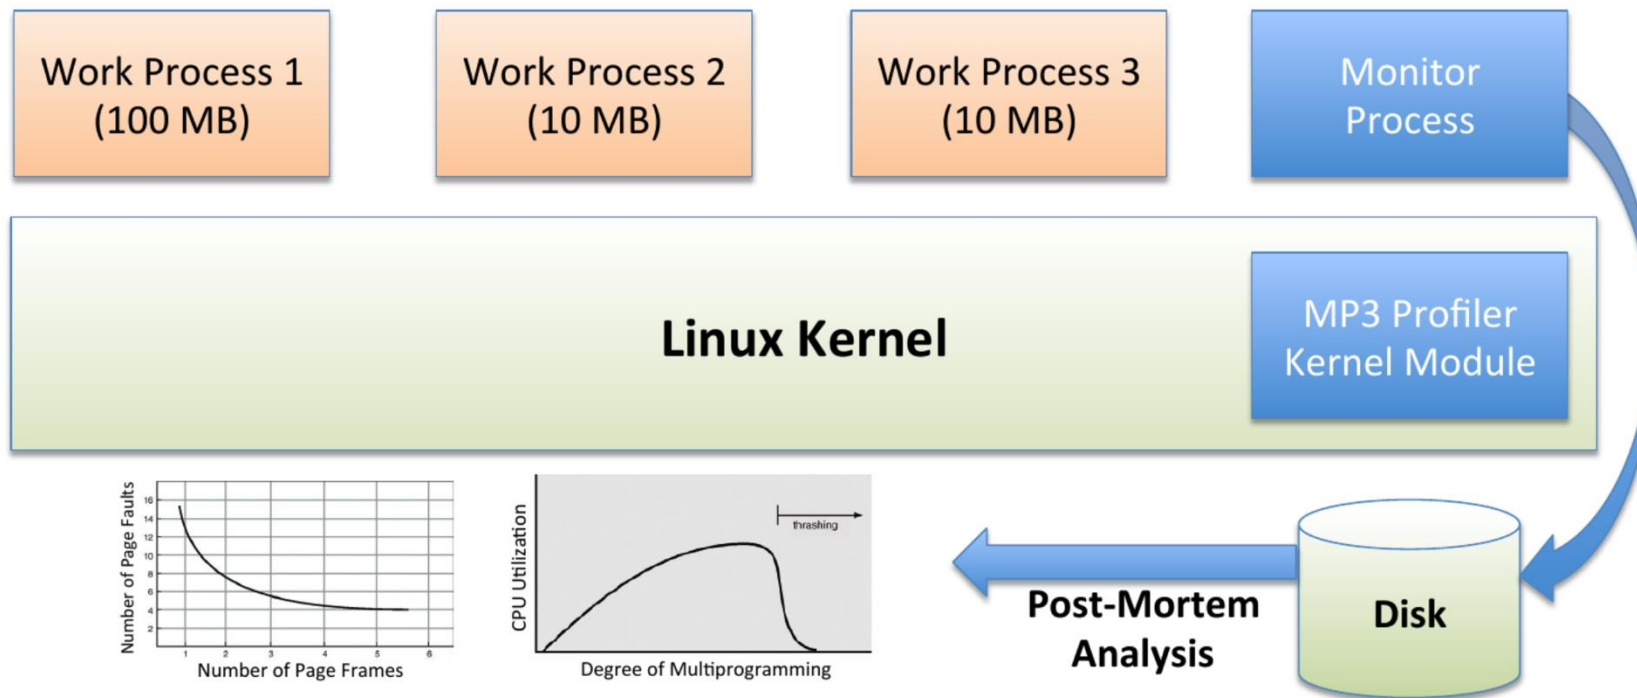
- **What is a Page Fault?**
- **What is a Minor Page Fault?**
- **What is a Major Page Fault?**
- **How do various PFs affect CPU usage?**

MP3 Profiler

- Can profile multiple programs.
- Provides: jiffies, minor PF count, major PF count, and CPU usage data.
- Read out the data using a monitor.
- Analyze the PF and how it affects CPU usage.

MP3 Profiler

- Can profile multiple work programs (provided).
- Captures: jiffies, minor PF count, major PF count, and CPU usage data.
- Save the statistics using a monitor program (provided).



Register, De-Register, and Process List

- **Register:** worker registers itself with its PID via proc write. `R <PID>`
- **De-register:** worker has finished and de-registers itself via proc write. `U <PID>`
- **MP1-style process list** via proc read.

Profiling PF

- Use a delayed workqueue
- Profile at 20 Hz (we provide a function):

```
struct kbuf_sample {  
    unsigned long jiffies;  
    unsigned long minfault;  
    unsigned long majfault;  
    unsigned long cpu;  
};
```

- However, these statistics are only available in kernel (read via procfs?).
- We want to avoid overheads crossing the kernel boundary.

Buffer Sharing

- **We can create and populate a buffer in kernel.**
- **Map the buffer back to user.**
- **So, the buffer is shared between kernel and user.**

Creating the Buffer

- **Must be at least 512KB (we force you to use `vmalloc()`).**
- **Must not be swapped to hard disk (reserved bit).**
- **Kernel populates it, while monitor program reads it.**

Populating the Buffer

- We make the buffer a queue.
- Must contains 12000 samples.
- For each sample:


```
struct kbuf_sample {  
    unsigned long jiffies;  
    unsigned long minfault;  
    unsigned long majfault;  
    unsigned long cpu;  
};
```

Mapping the Buffer

- Implement `mmap`
 - `procfs`, which everyone is very familiar with, does not support it!
- Implement a character device.
- Remap buffer pages to user page table.
 - Does kernel still have access?
 - Be careful that pages are not continuous!

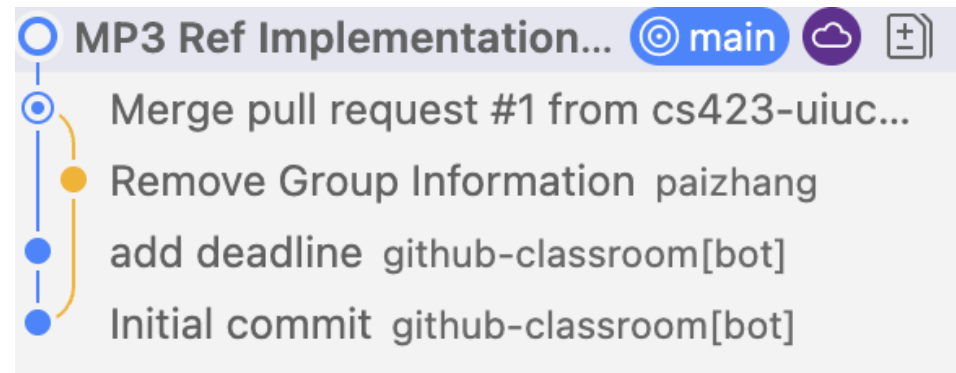
Case Study

- We provide the monitor and work program.
- You will be required to do some case studies.
- CPU Utilization:
$$U_T = \frac{\text{cpu time}_T}{\text{wall time}} = \frac{\text{stime}_T + \text{utime}_T}{\text{jiffies}}$$



Pro Tips

- Good tools can be very helpful.
- Try conventional commits.
- Source code is your best teacher.
- Start right away!



```
413     module_init(mp3_init);  
414     module_exit(mp3_exit);  
415
```

Total: ~400 lines