



# CS 423

# Operating System Design:

# OS Support for Containers

Tianyin Xu

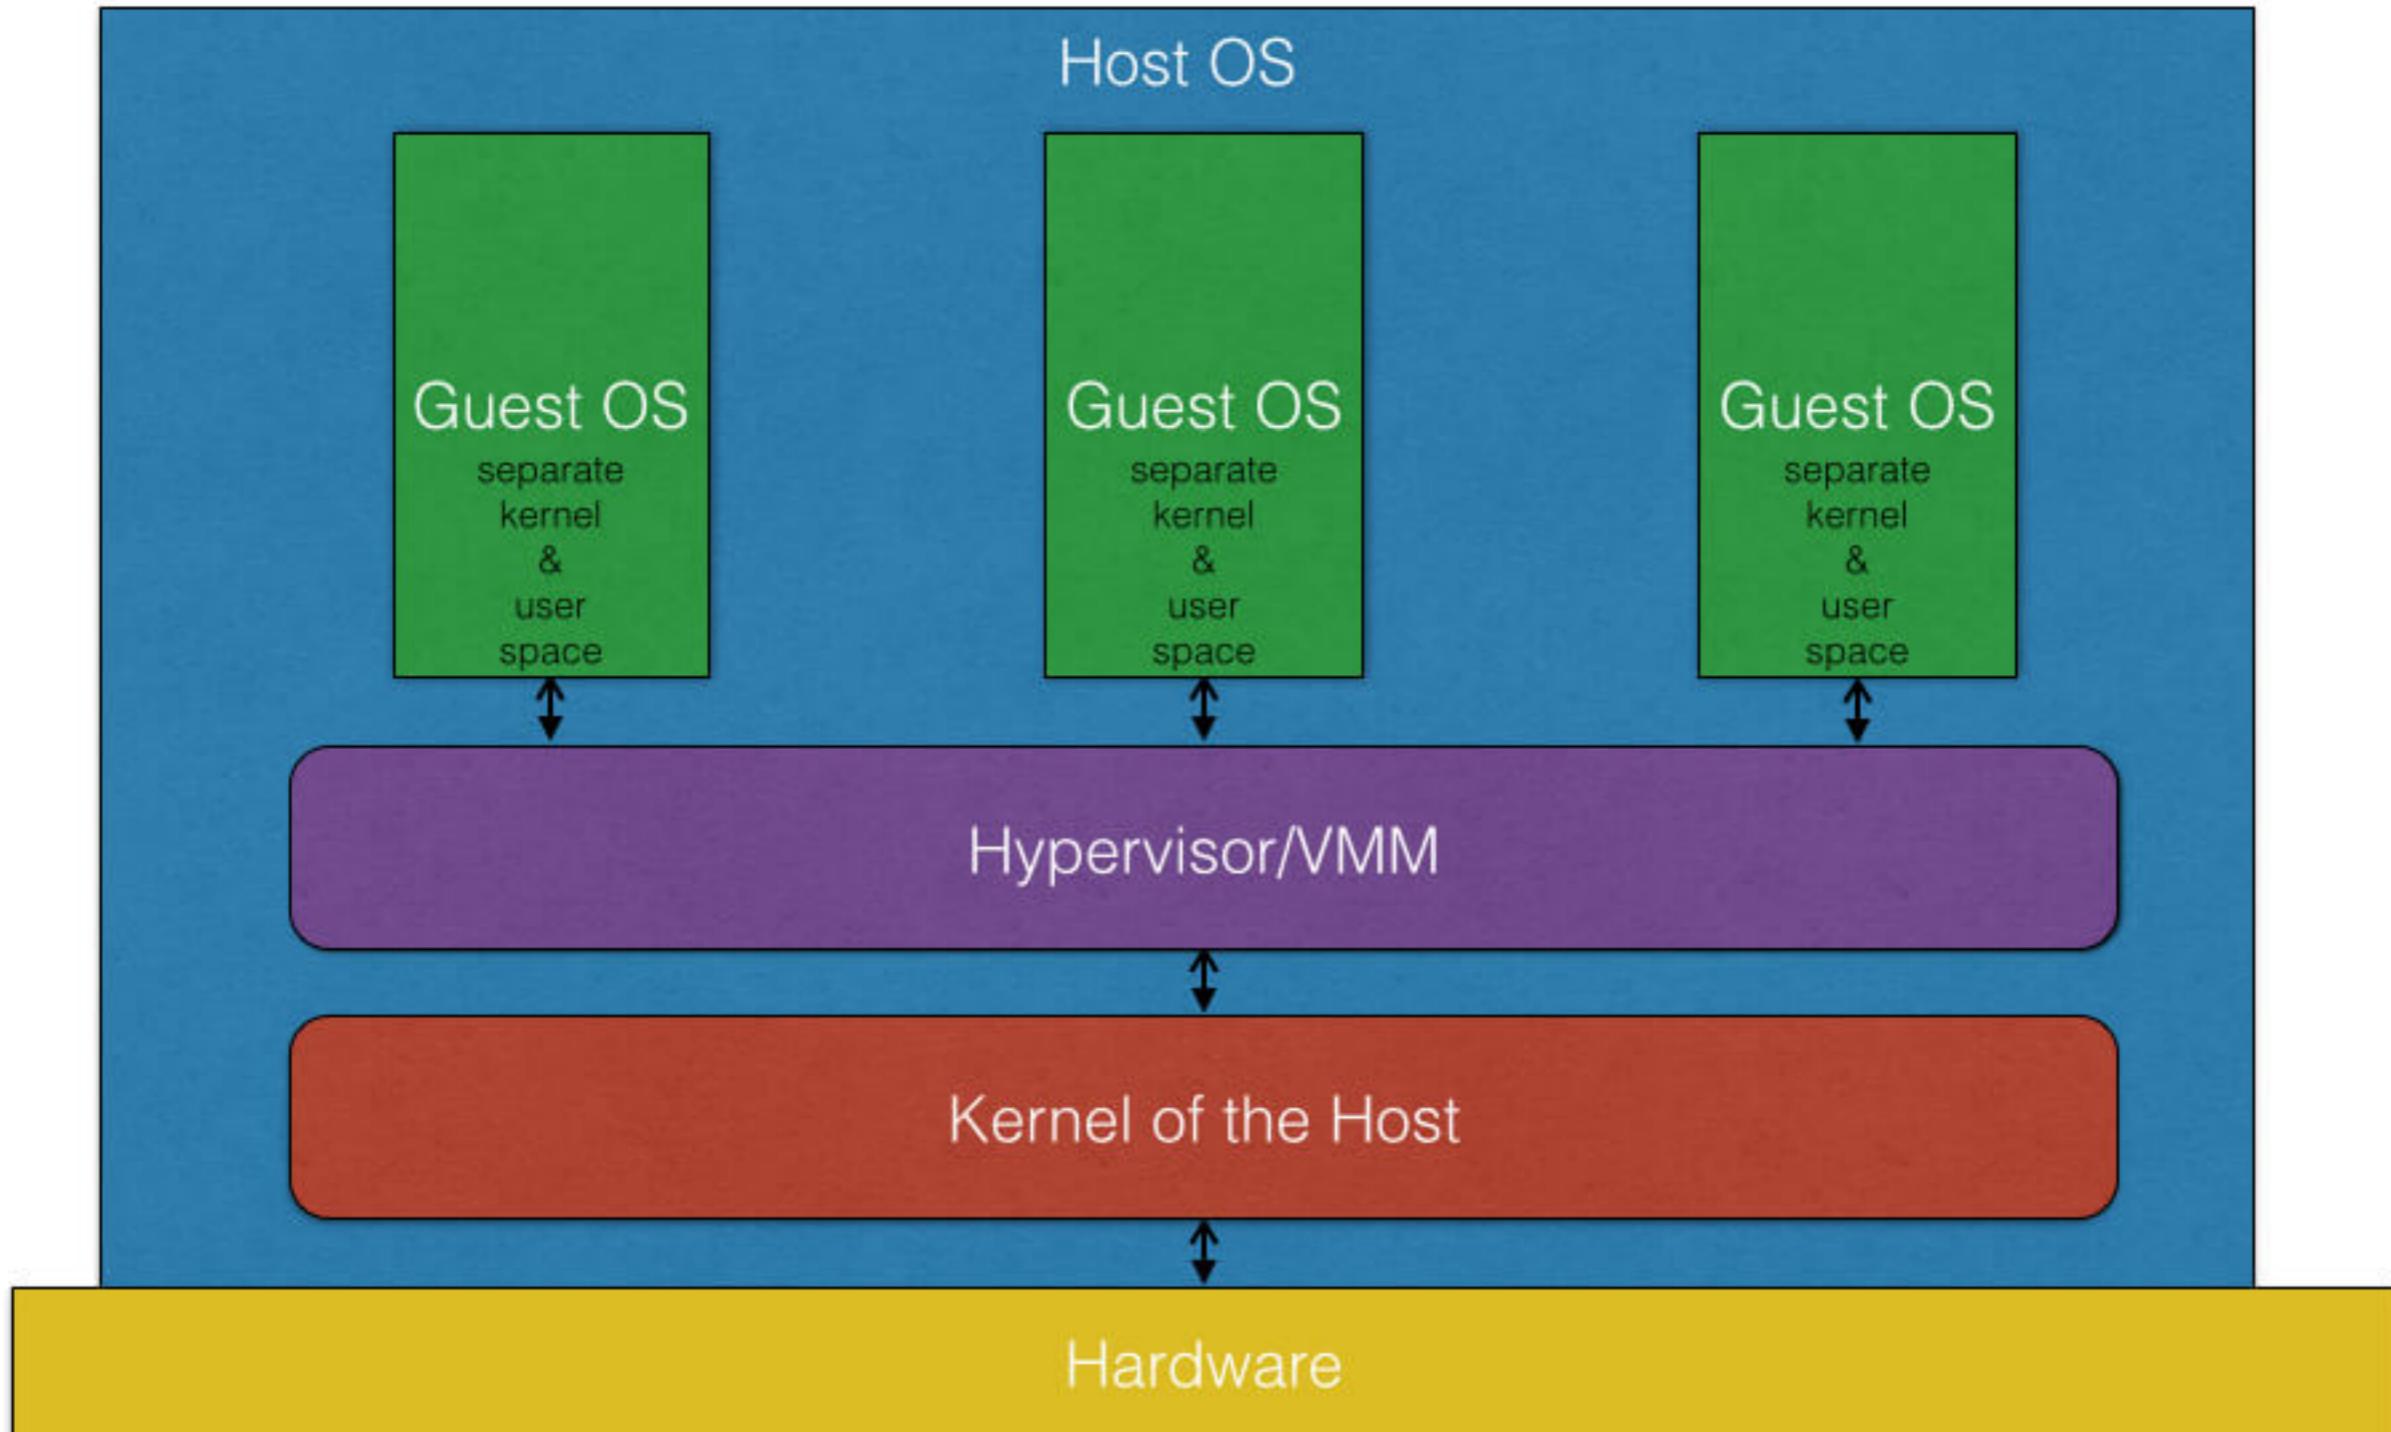
Thanks for Adam Bates and Julie Evans

# Cloud Computing (Gen I)



- Dominated by Infrastructure-as-a-Service clouds (and storage services)
- Big winner was Amazon EC2
- Hypervisors that virtualized the hardware-software interface
- Customers were responsible for provisioning the software stack from the kernel up

# Hypervisors



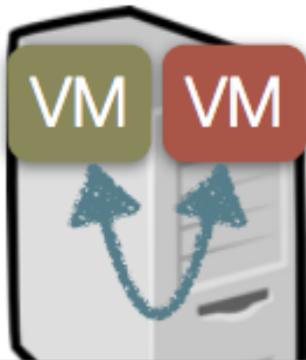


# Hypervisors

- Strong isolation between different customer's virtual machines
- VMM is ‘small’ compared to the kernel... less LoC means less bugs means (~)more security.

# Hypervisors

- ‘Practical’ attacks on IaaS clouds relied on side channels to detect co-location between attacker and victim VM
- E.g., we could correlate the performance of a shared resource
  - network RTT’s, cache performance
- After co-resident, make inferences about victim’s activities



n/w pings or  
covert-channels

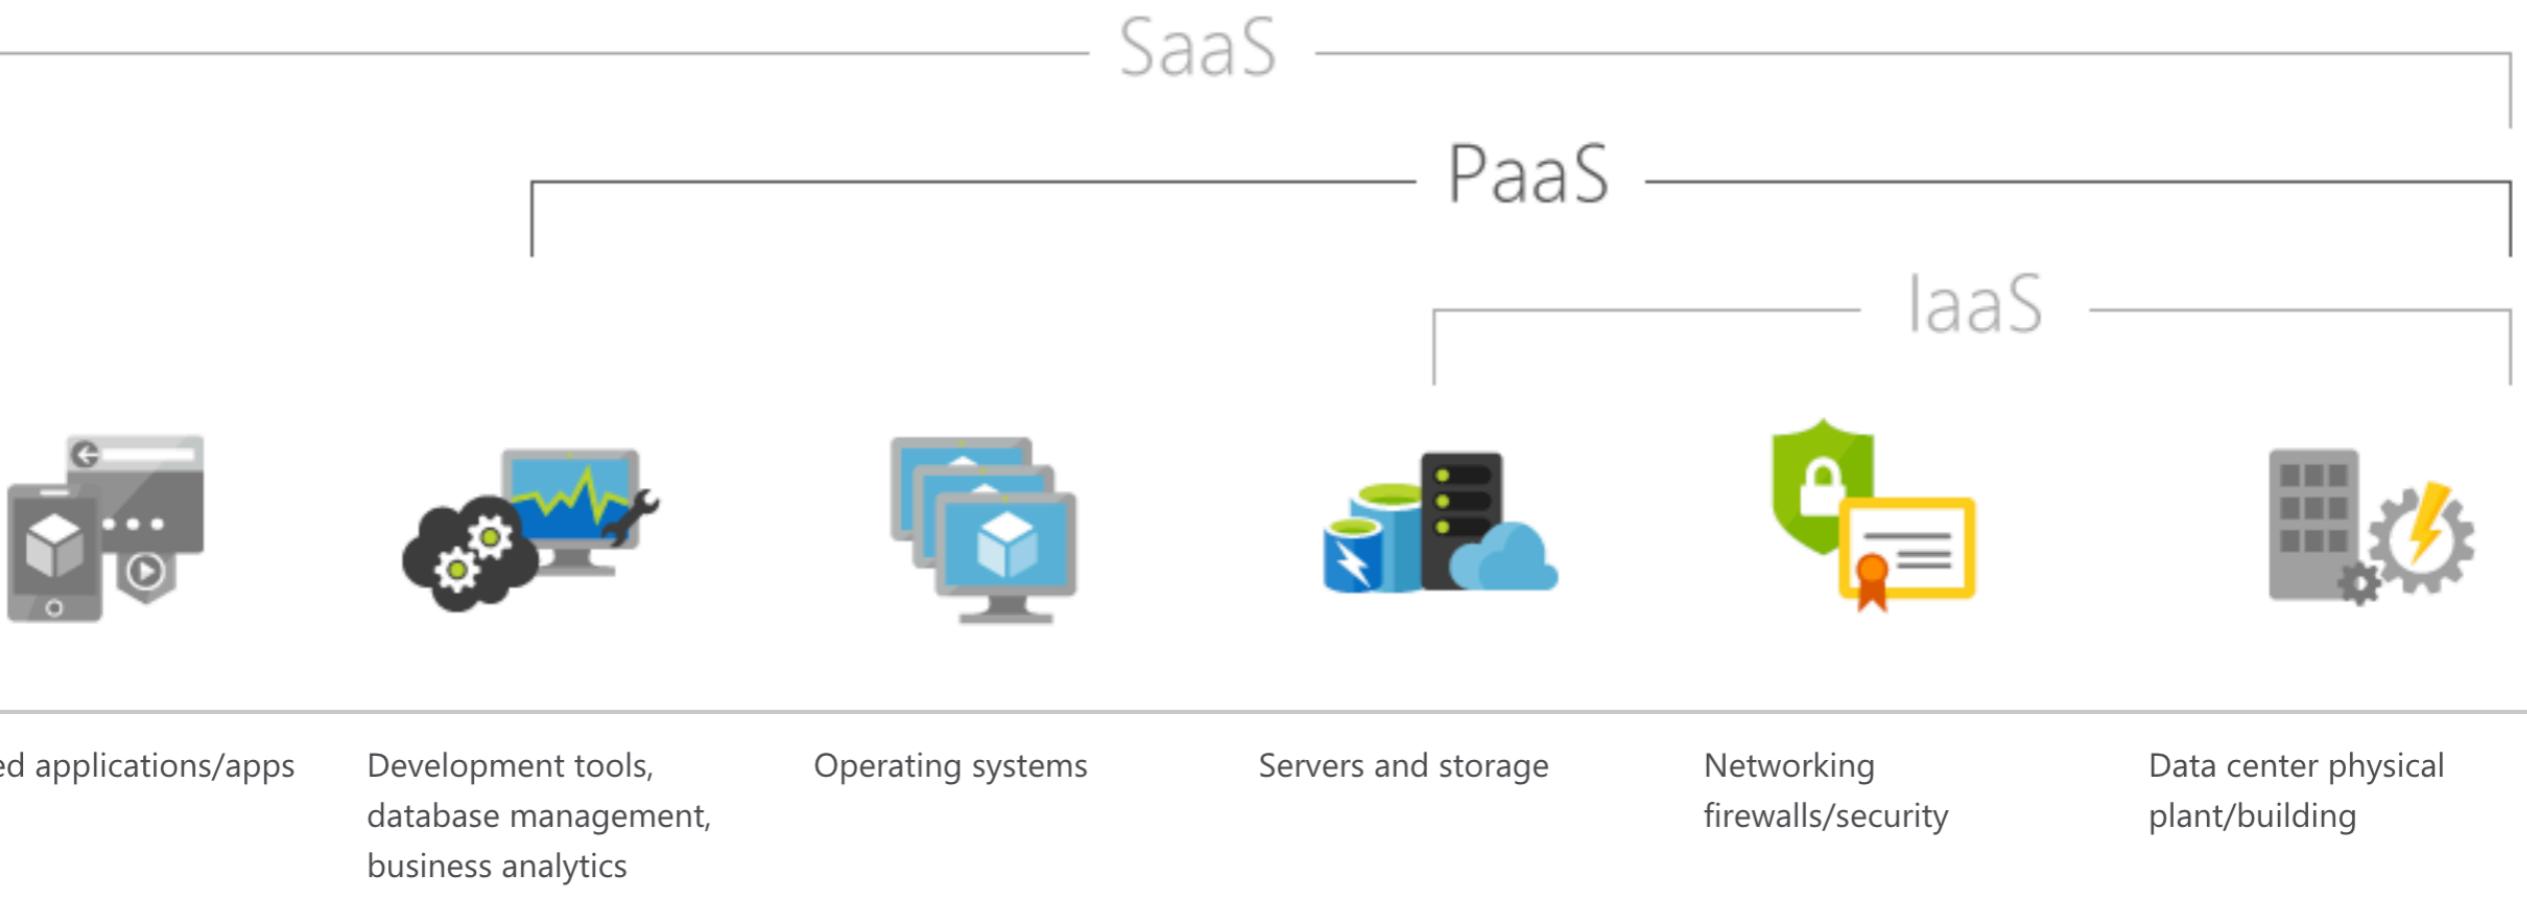


# Hypervisors

- Strong isolation between different customer's virtual machines
- VMM is ‘small’ compared to the kernel... less LoC means less bugs means (~)more security.
- High degree of flexibility... but did most customers really need it?

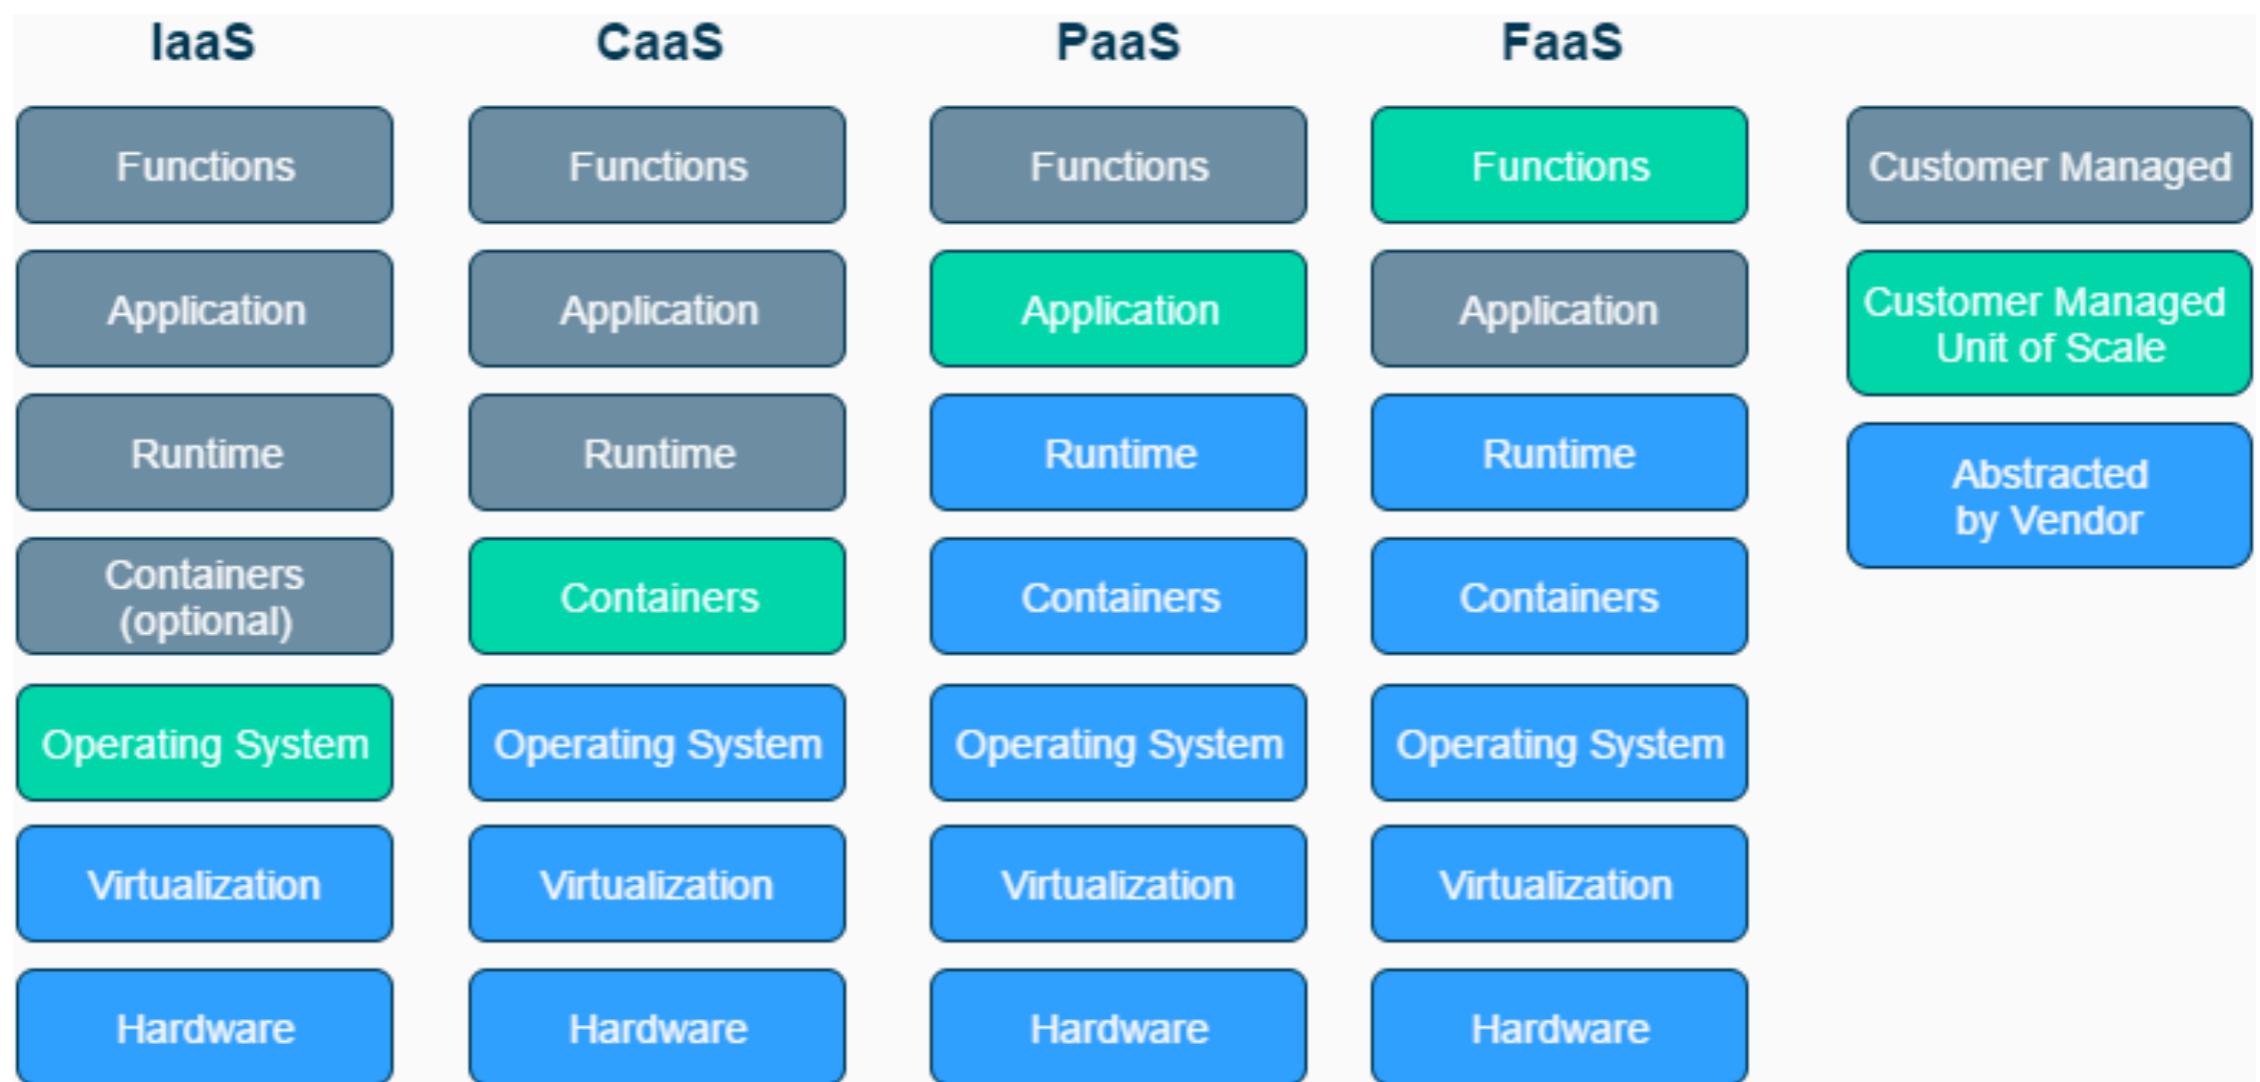
# Cloud Computing (Gen 2)

- PaaS: Platform as a Service
- SaaS: Software as a Service



# New Gen of Cloud Computing

- **FaaS: Function as a Service**



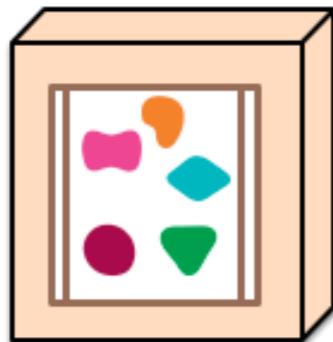
# New Gen of Cloud Computing

- ## Microservices

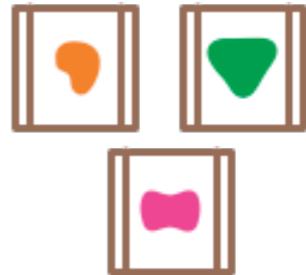
*A monolithic application puts all its functionality into a single process...*



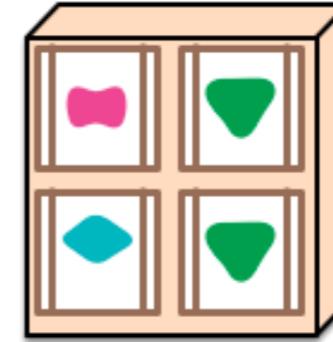
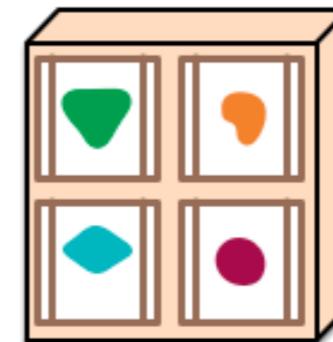
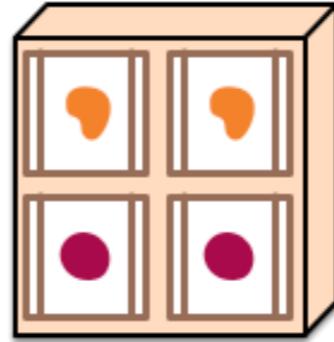
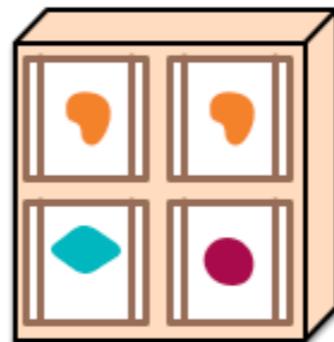
*... and scales by replicating the monolith on multiple servers*



*A microservices architecture puts each element of functionality into a separate service...*



*... and scales by distributing these services across servers, replicating as needed.*

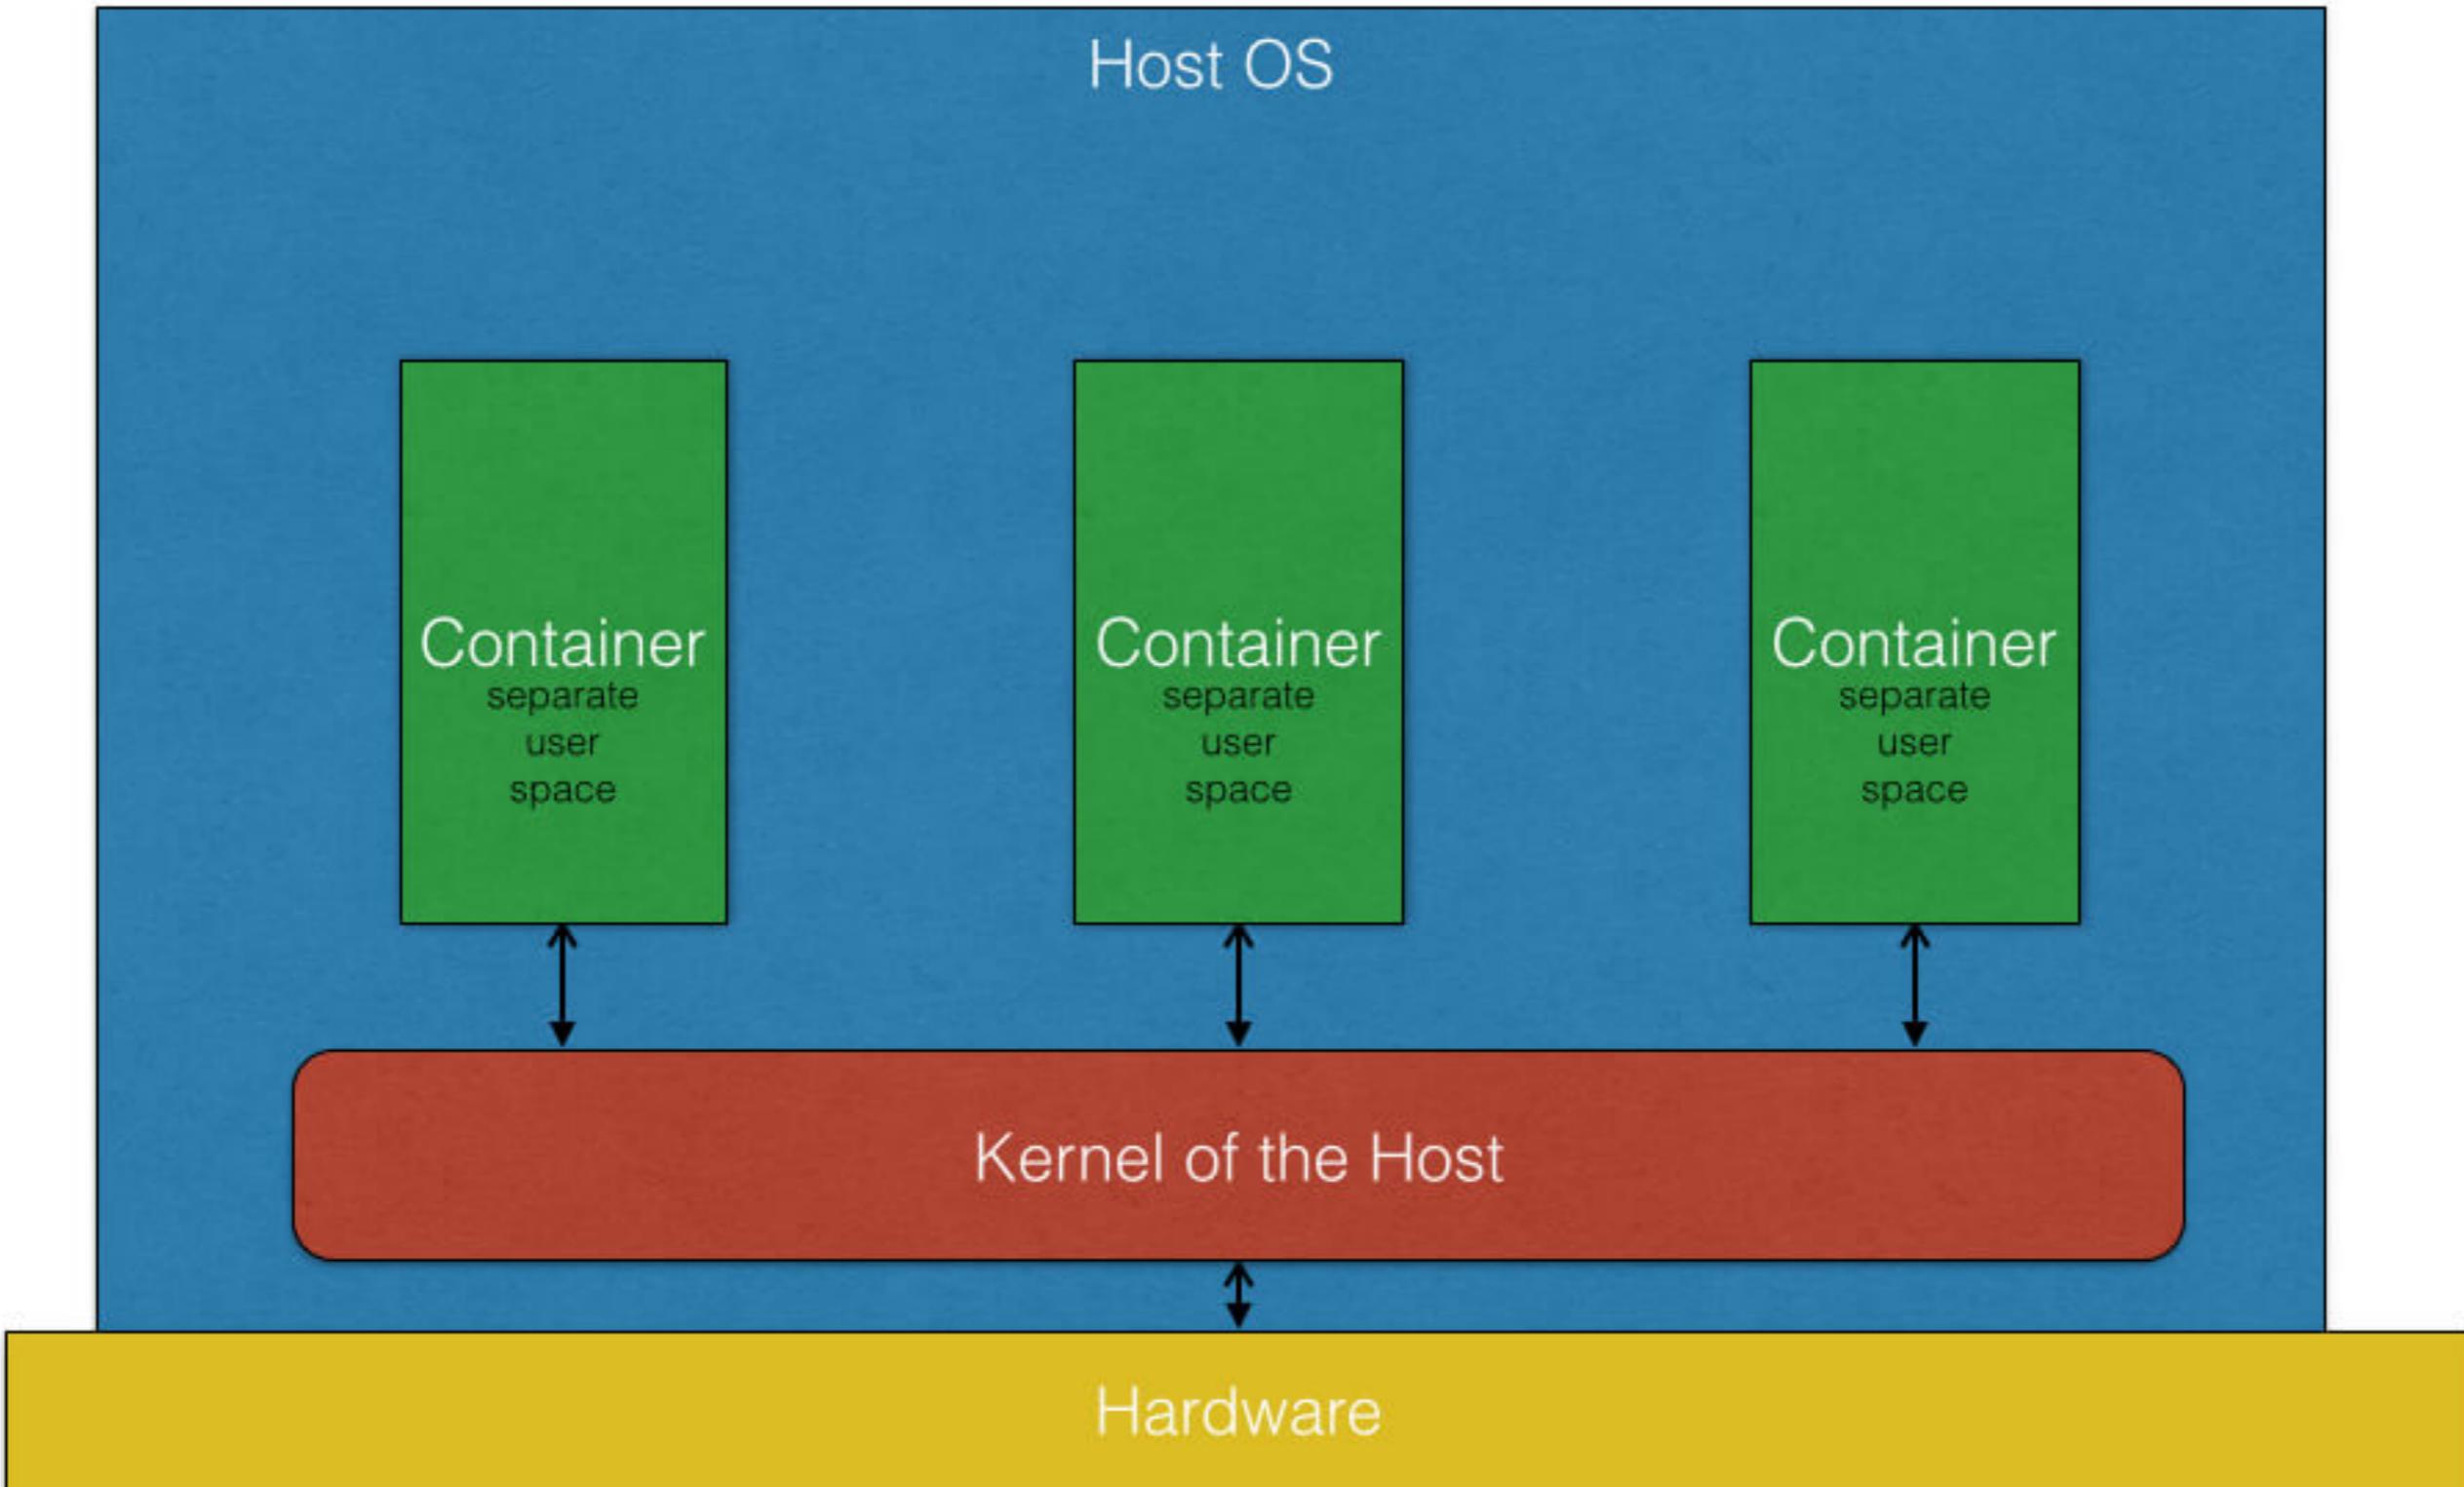




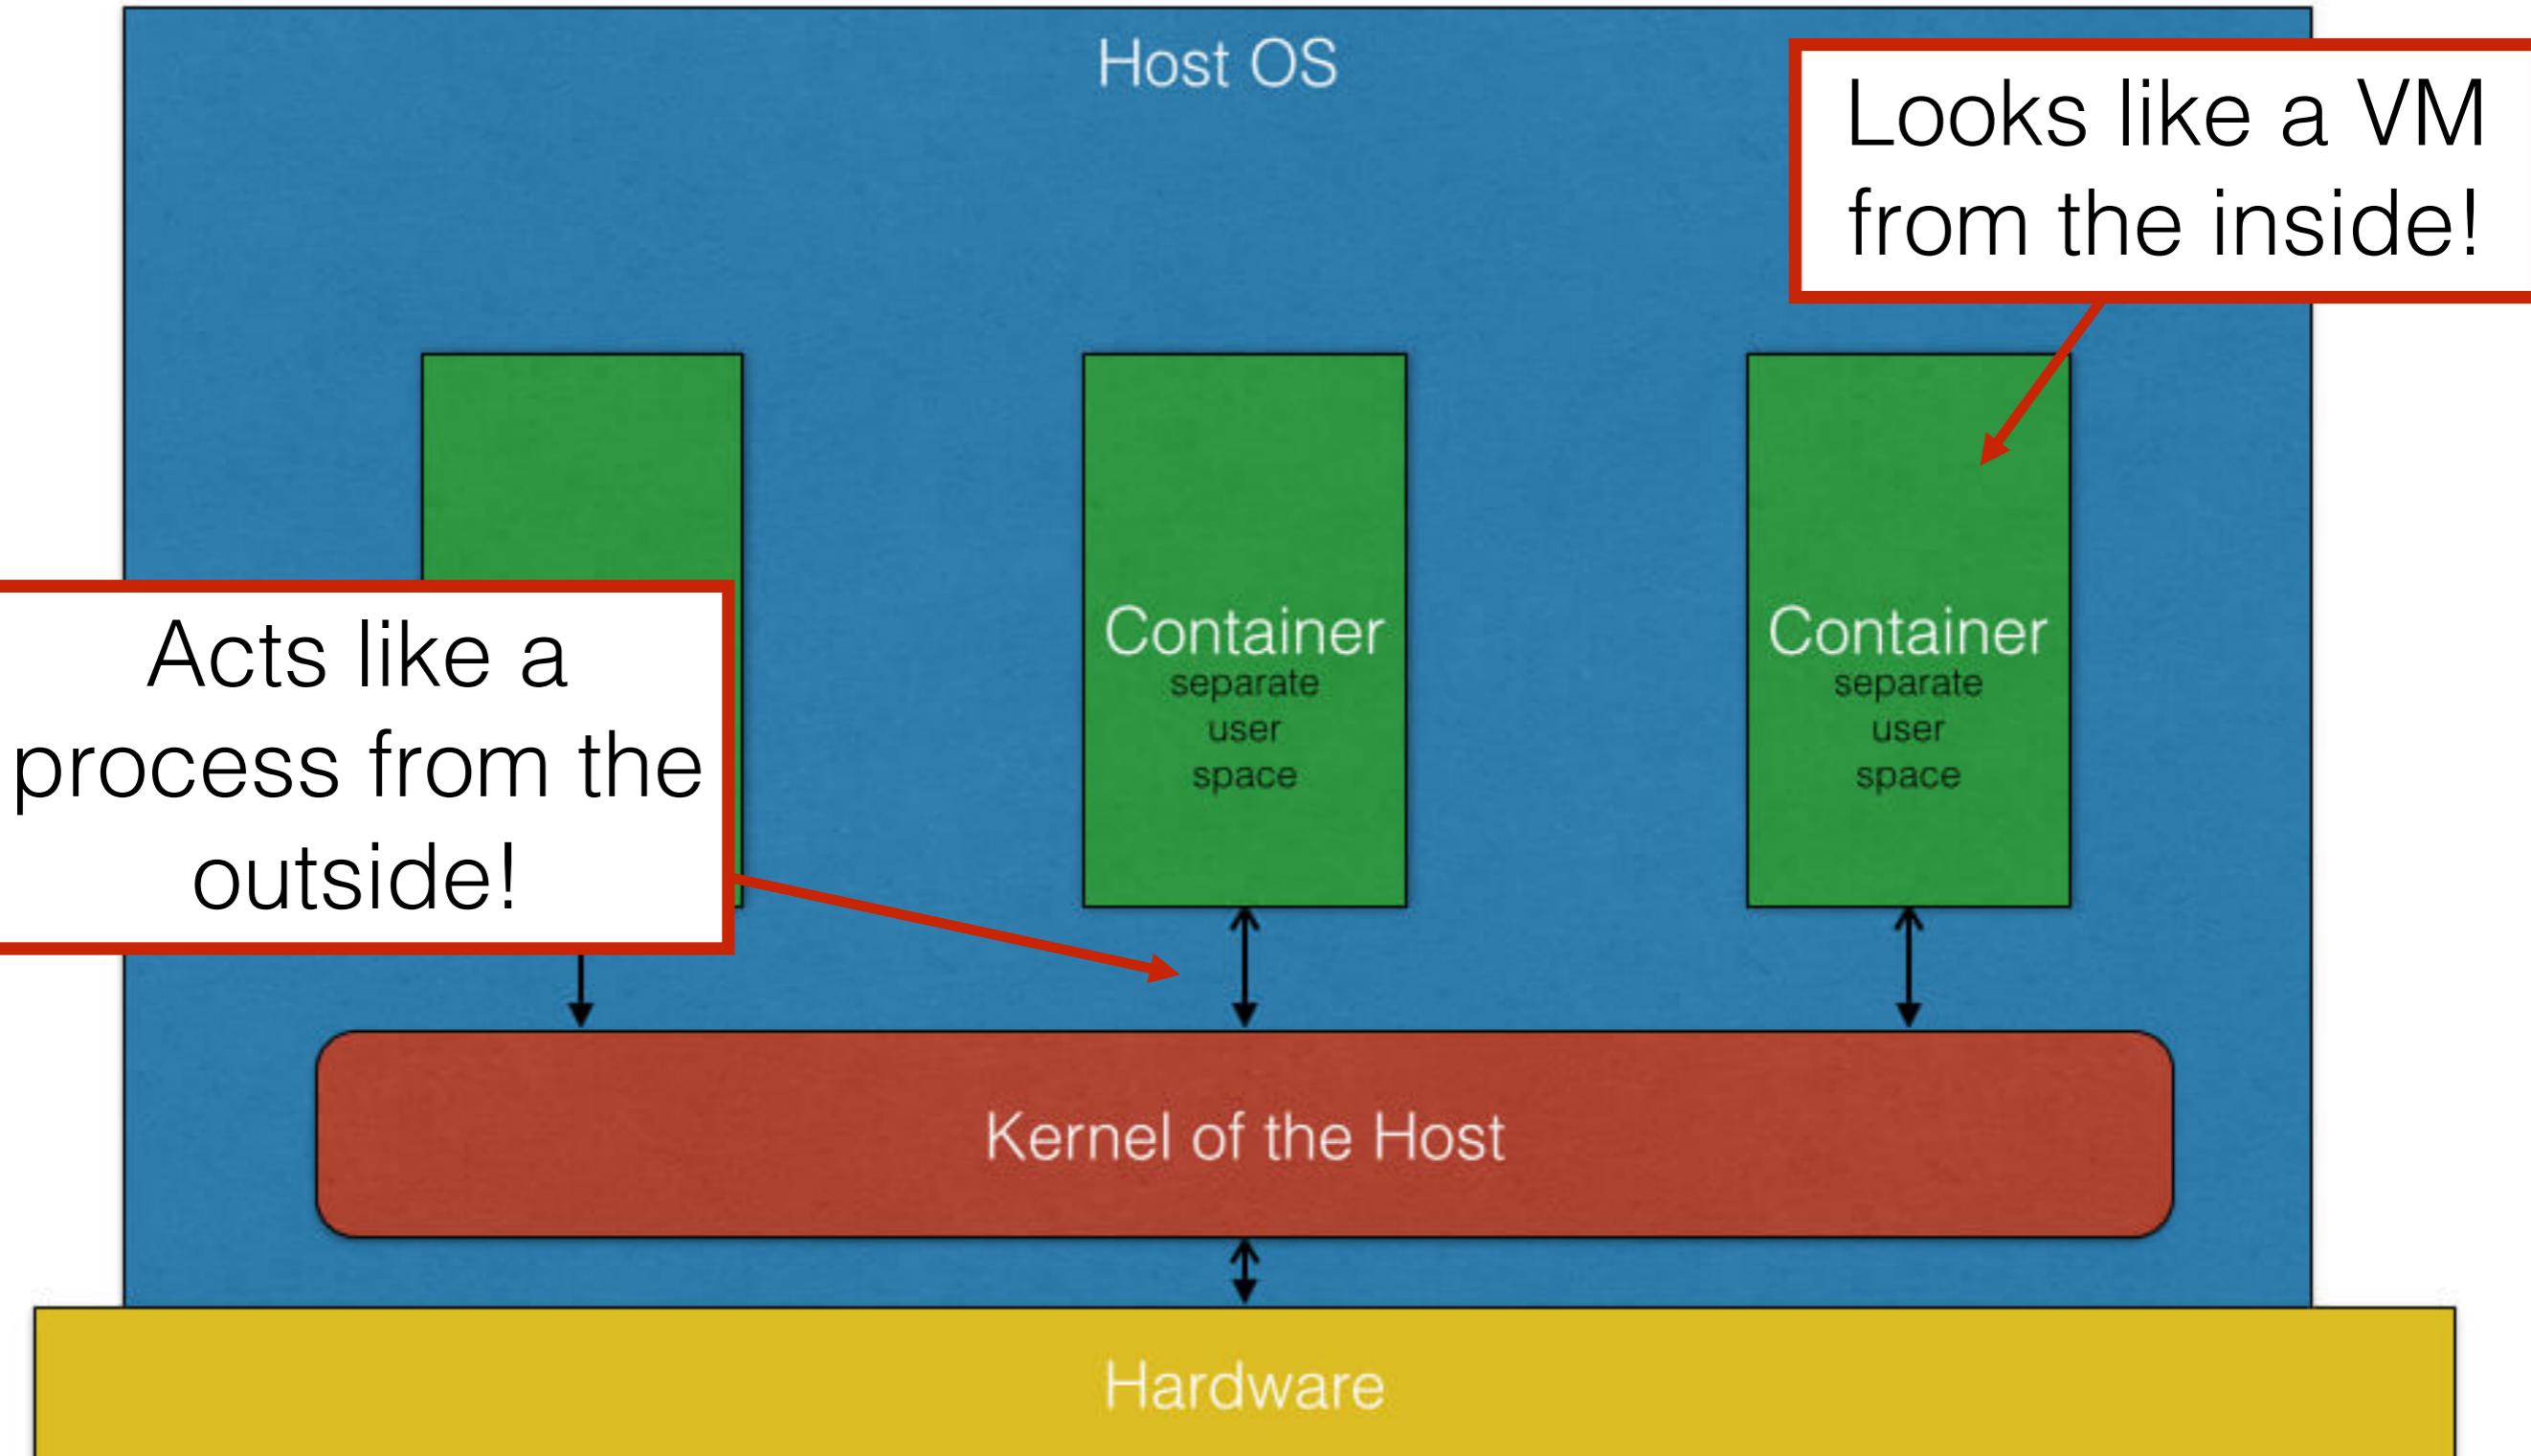
# Enter Containers

- Rather than virtualize both user space and kernel space... why not just ‘virtualize’ user space?
- Meets the needs of most customers, who don’t require significant customization of the OS.
- Sometimes called ‘OS virtualization,’ which is highly misleading given our existing taxonomy of virtualization techniques
- Running natively on host, containers enjoy bare metal performance without reliance on advanced virtualization support from hardware.

# Enter Containers



# Enter Containers

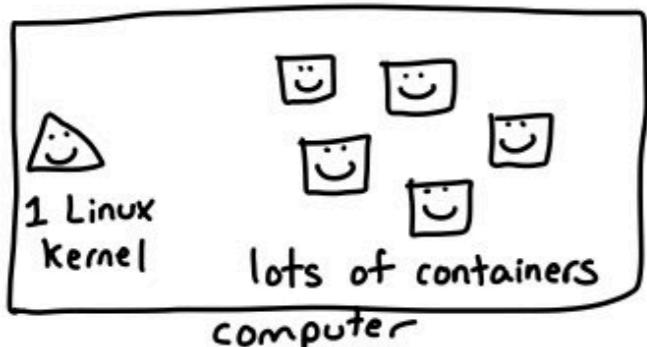


# Containers are processes

JULIA EVANS  
@bork

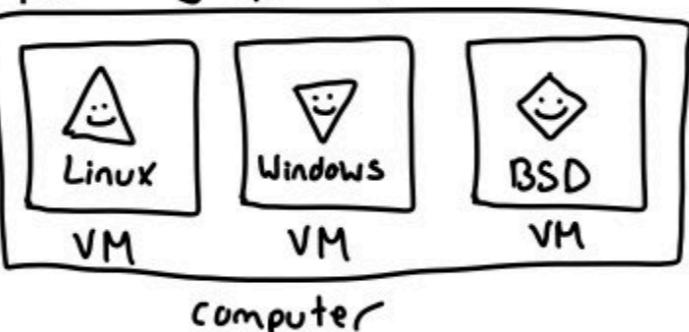
## containers vs VMs

a container is a group of processes



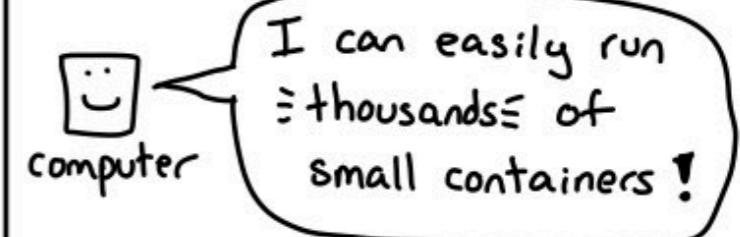
a virtual machine is a fake computer

each one has its own operating system!



containers use less RAM

This is because they share a single Linux kernel.

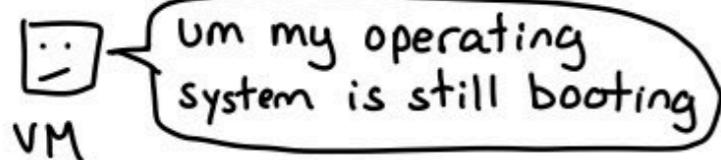


Containers start faster

because they're processes and process start fast !

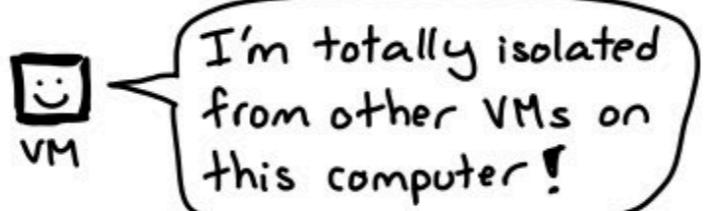


container

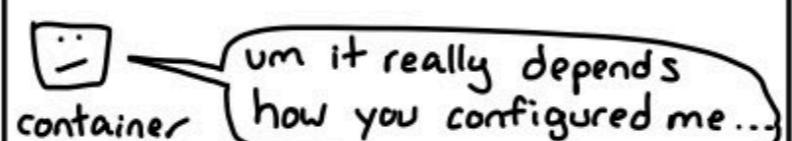


VM

Containers are more complicated to secure

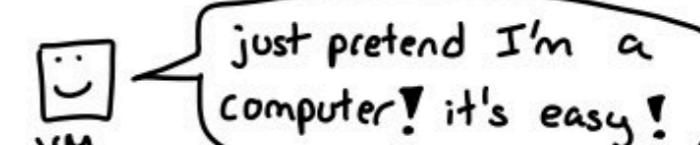


VM

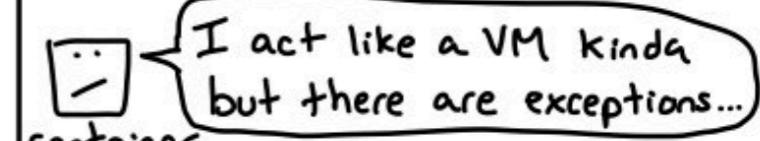


container

it's harder to figure out what you can do in a container



VM



container

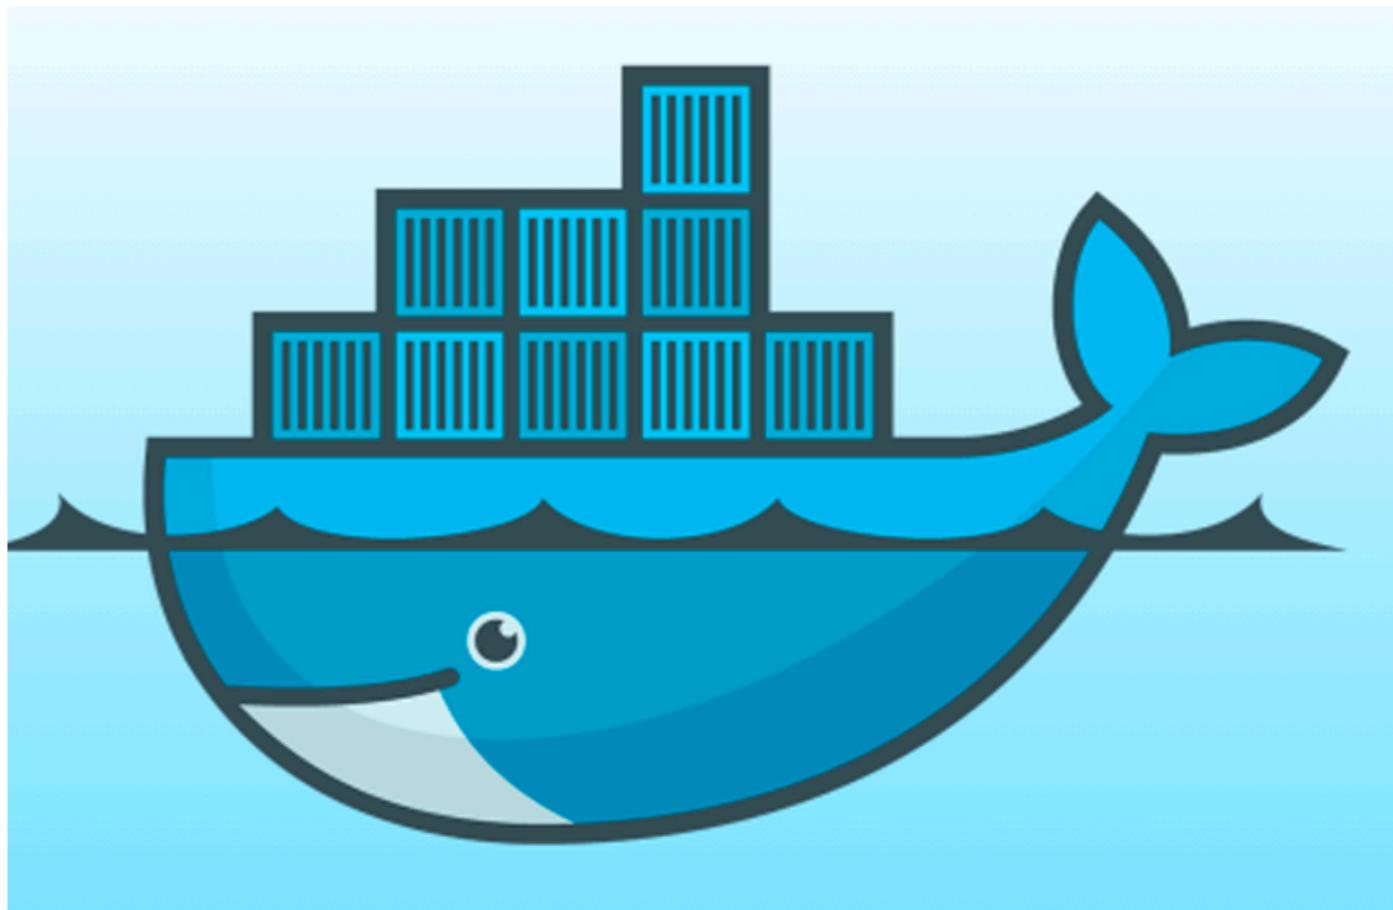


# "Container is an old idea"

- You didn't heard of it cause it was not called "containers."
  - Linux containers
  - BSD Jails
  - Solaris Zones

# Docker's Big Idea

- Build, Ship, and Run App, Anywhere
- Debug your app, not your environment --  
Securely build and share any application,  
anywhere



# Docker's Big Idea

JULIA EVANS  
@bork

## the big idea: include EVERY dependency

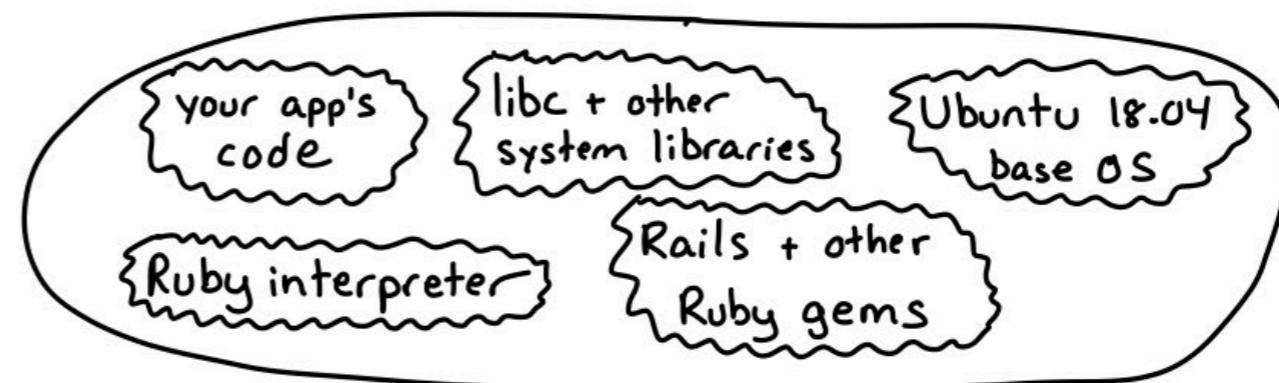
containers package  
EVERY dependency  
together



to make sure this  
program will run on  
your laptop, I'm going  
to send you every single  
file on my computer  
exaggeration but  
it's the basic idea

a container image is a tarball of a filesystem

Here's what's in a typical Rails app's container:



how images are built

0. start with a base os
1. install program+dependencies
2. configure it how you want
3. make a tarball of the  
WHOLE FILESYSTEM

(this is what 'docker build' does)

Running an image

1. download the tarball
2. unpack it into a directory
3. Run a program and pretend  
that directory is its  
whole filesystem

(this is what 'docker run' does)

images let you "install"  
programs really easily



wow, I can get a  
Postgres test database  
running in 45 seconds!



# OS Support for Containers

- **Linux Containers (LXC):**
  - **chroot**
  - **namespace**
    - **PID, Network, User, IPC, uts, mount**
  - **cgroups for HW isolation**
  - **Security profiles and policies**
    - **Apparmor, SELinux, Seccomp**

# containers = chroot on steroids

- chroot changes the apparent root directory for a given process and all of its children
- An old idea! POSIX call dating back to 1979
- Not intended to defend against privileged attackers... they still have root access and can do all sorts of things to break out (like chroot'ing again)
- Hiding the true root FS isolates a lot; in \*nix, file abstraction used extensively
- Does not completely hide processes, network, etc., though!



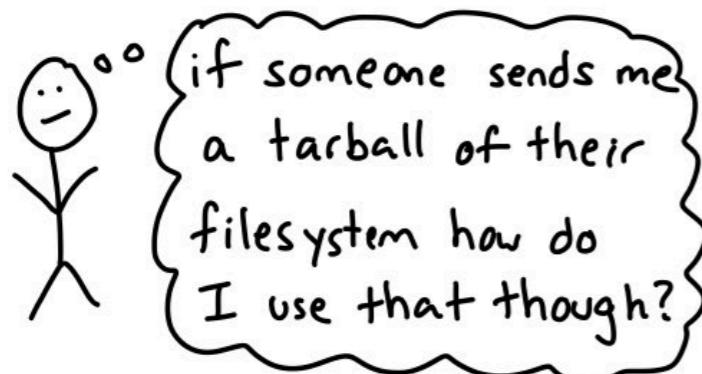
# Chroot

JULIA EVANS  
@b0rk

## chroot

a container image is a tarball of a filesystem

(or several tarballs: 1 per layer)



chroot: trick a program into thinking it has a different root directory

```
$ ls /path/to/container_filesystem
bin/ etc/ usr/ var/
$ sudo chroot /path/to/container_filesystem /bin/bash
(inside chroot now)
$ ls /
bin/ etc/ usr/ var/
```

that's our new fake root directory!  
we tricked ls!

very basic way to "run" a Redis container

tarball of filesystem with Redis installed

```
$ mkdir redis; cd redis
$ tar -xzf redis.tar
$ chroot $PWD /usr/bin/redis
# done! redis is running!
```

problems with just using chroot

- no CPU/memory limits
- other running processes are still visible
- can't use the same network port as another process
- LOTS of security issues

Docker uses pivot-root + extra isolation features to run containers

pivot-root is like chroot but harder to escape from



# Namespaces

- The key feature enabling containerization!
- Partition practically all OS functionalities so that different process domains see different things
- Mount (mnt): Controls mount points
- Process ID (pid): Exposes a new set of process IDs distinct from other namespaces (i.e., the hosts)
- Network (net): Dedicated network stack per container; each interface present in exactly one namespace at a time.
- ....



# Namespaces

- The key feature enabling containerization!
- Partition practically all OS functionalities so that different process domains see different things
- Interprocess Comm. (IPC): Isolate processes from various methods of POSIX IPC.
  - e.g., no shared memory between containers!
- UTS: Allows the host to present different host/domain names to different containers.
- There's also a User ID (user) and cgroup namespace



# User Namespace

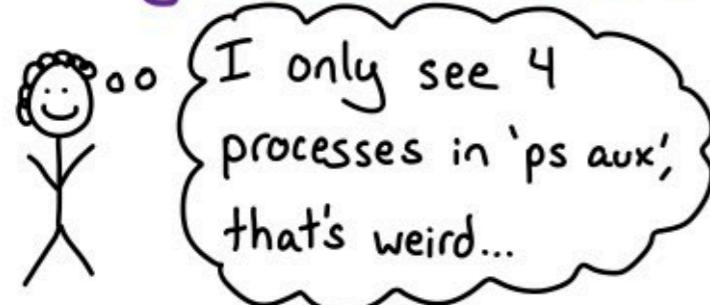
- Like others, can provide a unique UID space to the container.
- More nuanced though — we can map UID 0 inside the container to UID 1000 outside; allows processes inside of container to think they're root.
- Enables containers to perform administration actions, e.g., adding more users, while remaining confined to their namespace.

# Namespace

JULIA EVANS  
@b0rk

## namespaces

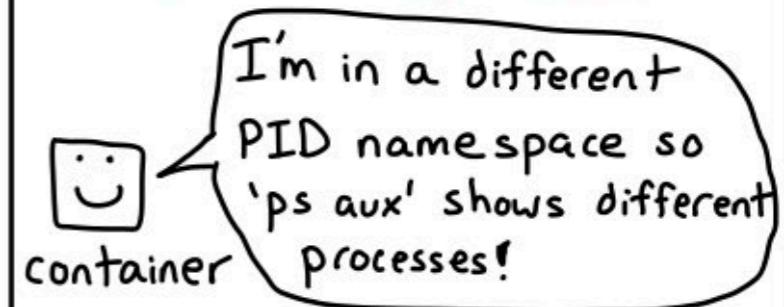
inside a container,  
things look different



Commands that will look different

- ps aux (less processes!)
- mount & df
- netstat -tulpn (different open ports!)
- hostname
- ... and LOTS more

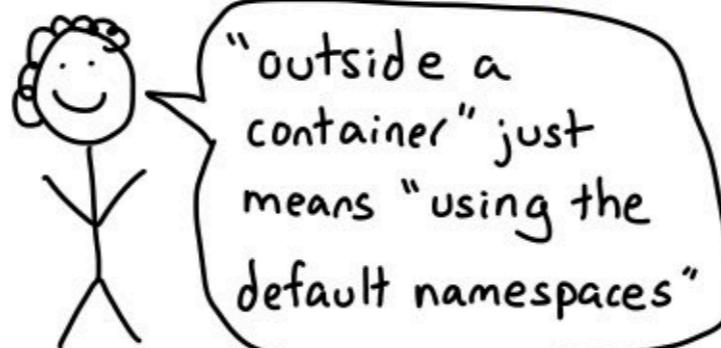
why those commands look different:  
:= namespaces=



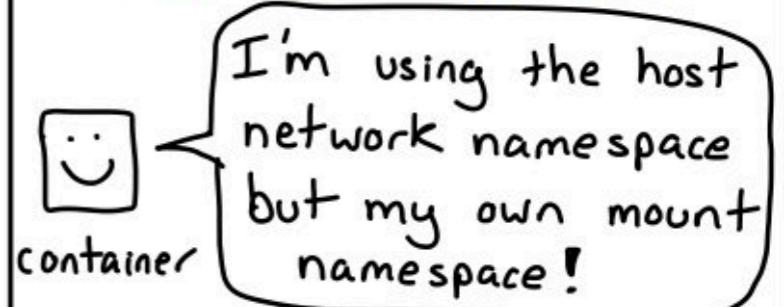
every process has 7 kinds of namespaces



there's a default ("host") namespace



processes can have any combination of namespaces



♥ this? more at [wizardzines.com](http://wizardzines.com)



# cgroups

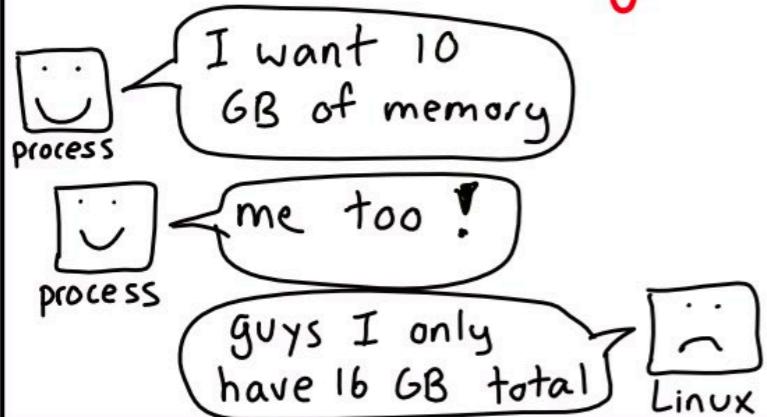
- Limit, track, and isolate utilization of hardware resources including CPU, memory, and disk.
- Important for ensuring QoS between customers! Protects against bad neighbors
- Features:
  - Resource limitation
  - Prioritization
  - Accounting (for billing customers!)
  - Control, e.g., freezing groups
- The cgroup namespace prevents containers from viewing or modifying their own group assignment

# cgroups

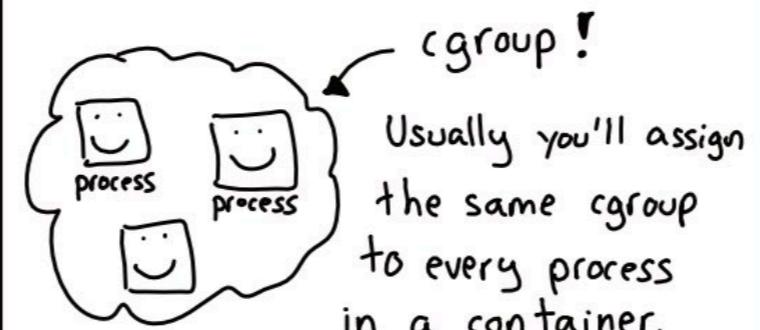
JULIA EVANS  
@b0rk

## cgroups

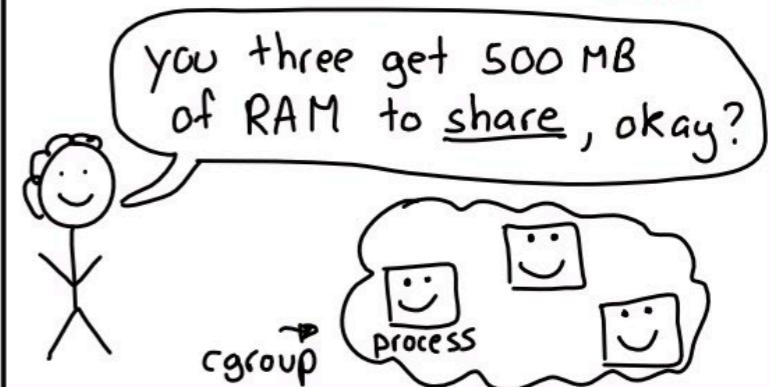
processes can use a lot of memory



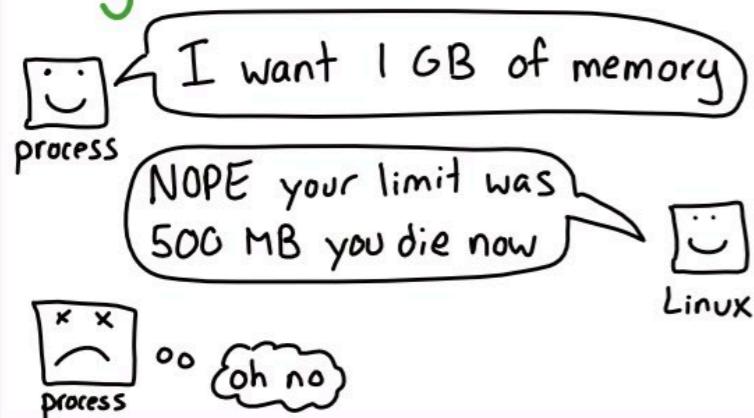
a cgroup is a group of processes



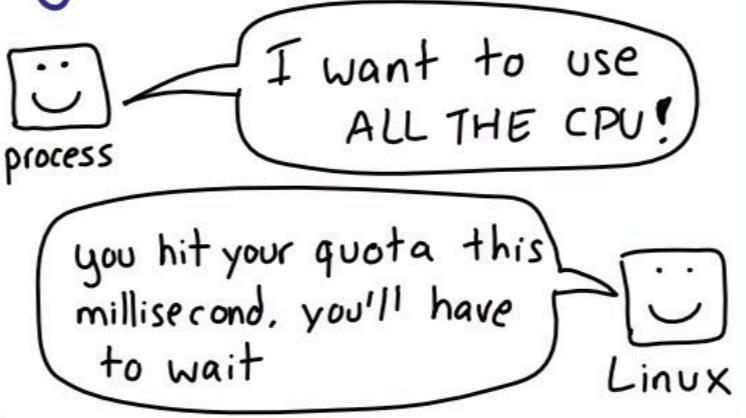
cgroups have memory / CPU limits



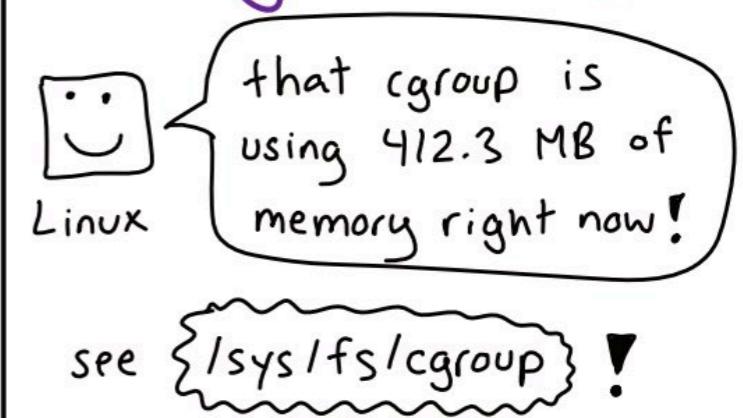
use too much memory:  
get OOM killed



use too much CPU:  
get slowed down



cgroups track  
memory & CPU usage





# Container Security?

*“Containers do not contain.” - Dan Walsh (SELinux contributor)*

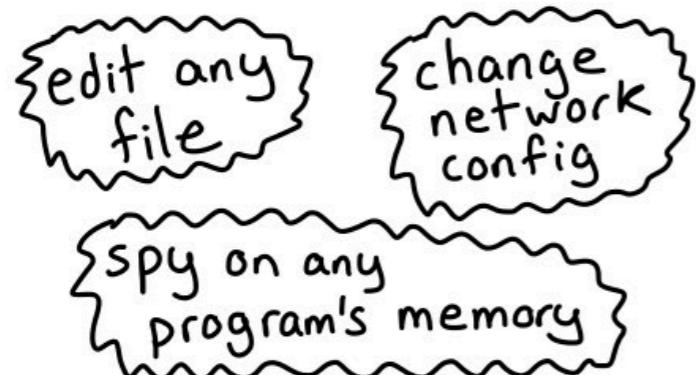
- In a nutshell, it's real hard to prove that every feature of the operating system is namespaced.
  - /sys? /proc? /dev? LKMs? kernel keyrings?
- Root access to any of these enables pwning the host
- Solution? Just don't forget about MAC; at this point SELinux pretty good support for namespace labeling.
- SELinux and Namespaces actually synergize nicely; much easier to express a correct isolation policy over a coarse-grained namespace than, say, individual processes

# Capabilities

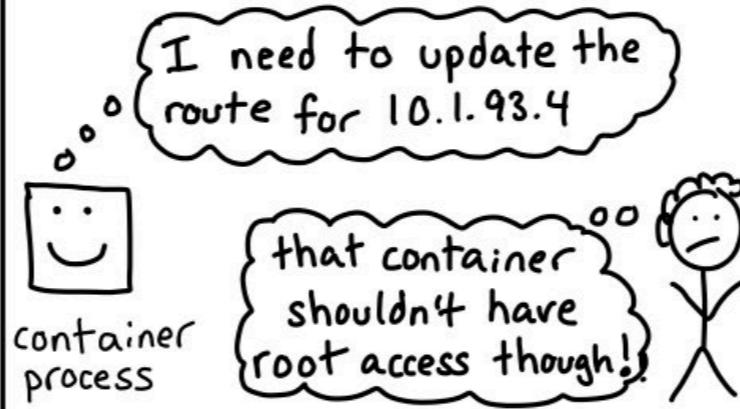
JULIA EVANS  
@b0rk

## capabilities

the root user can do **\*anything\***



sometimes containers need privileged access



**♥capabilities♥**  
let you grant specific permissions



### CAP\_SYS\_ADMIN

basically root access. Try to use a more specific capability!

### CAP\_NET\_ADMIN

for changing network settings

### \$ capsh --print

run this in a container to print its capabilities

### CAP\_SYS\_PTRACE

strace needs this

### CAP\_NET\_RAW

ping needs this to send raw ICMP packets

### \$ getcap /usr/bin/ping

shows which capabilities ping is allowed to use

# Seccomp-bpf

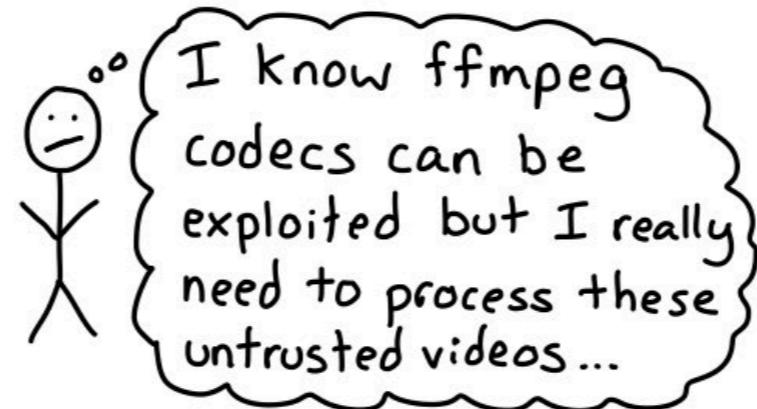
JULIA EVANS  
@b0rk

## seccomp-bpf

all programs use system calls



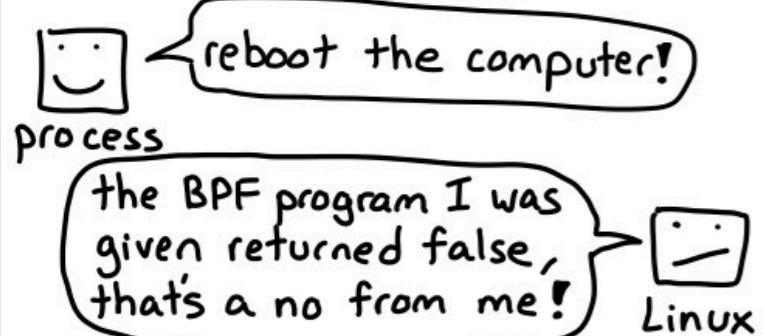
Some programs have security vulnerabilities



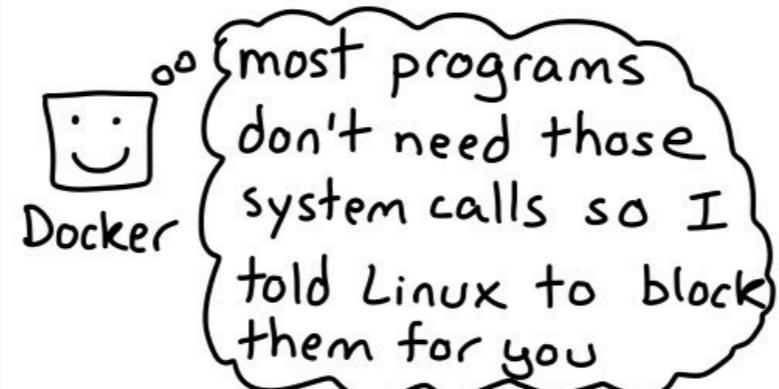
rarely used syscalls can help an attacker



**seccomp-BPF**: make Linux run a tiny program before every system call



Docker blocks dozens of syscalls by default



2 ways to block scary system calls

1. Limit a container's capabilities
2. Use a seccomp-BPF whitelist

Usually people do both!

# Linux Security Modules

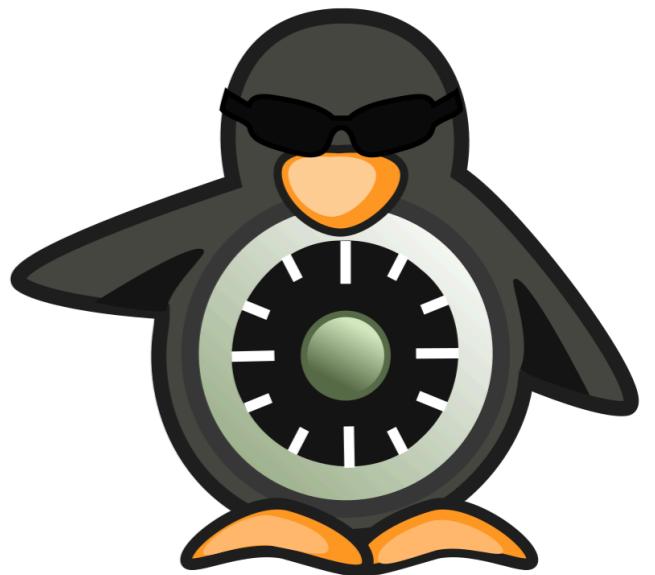
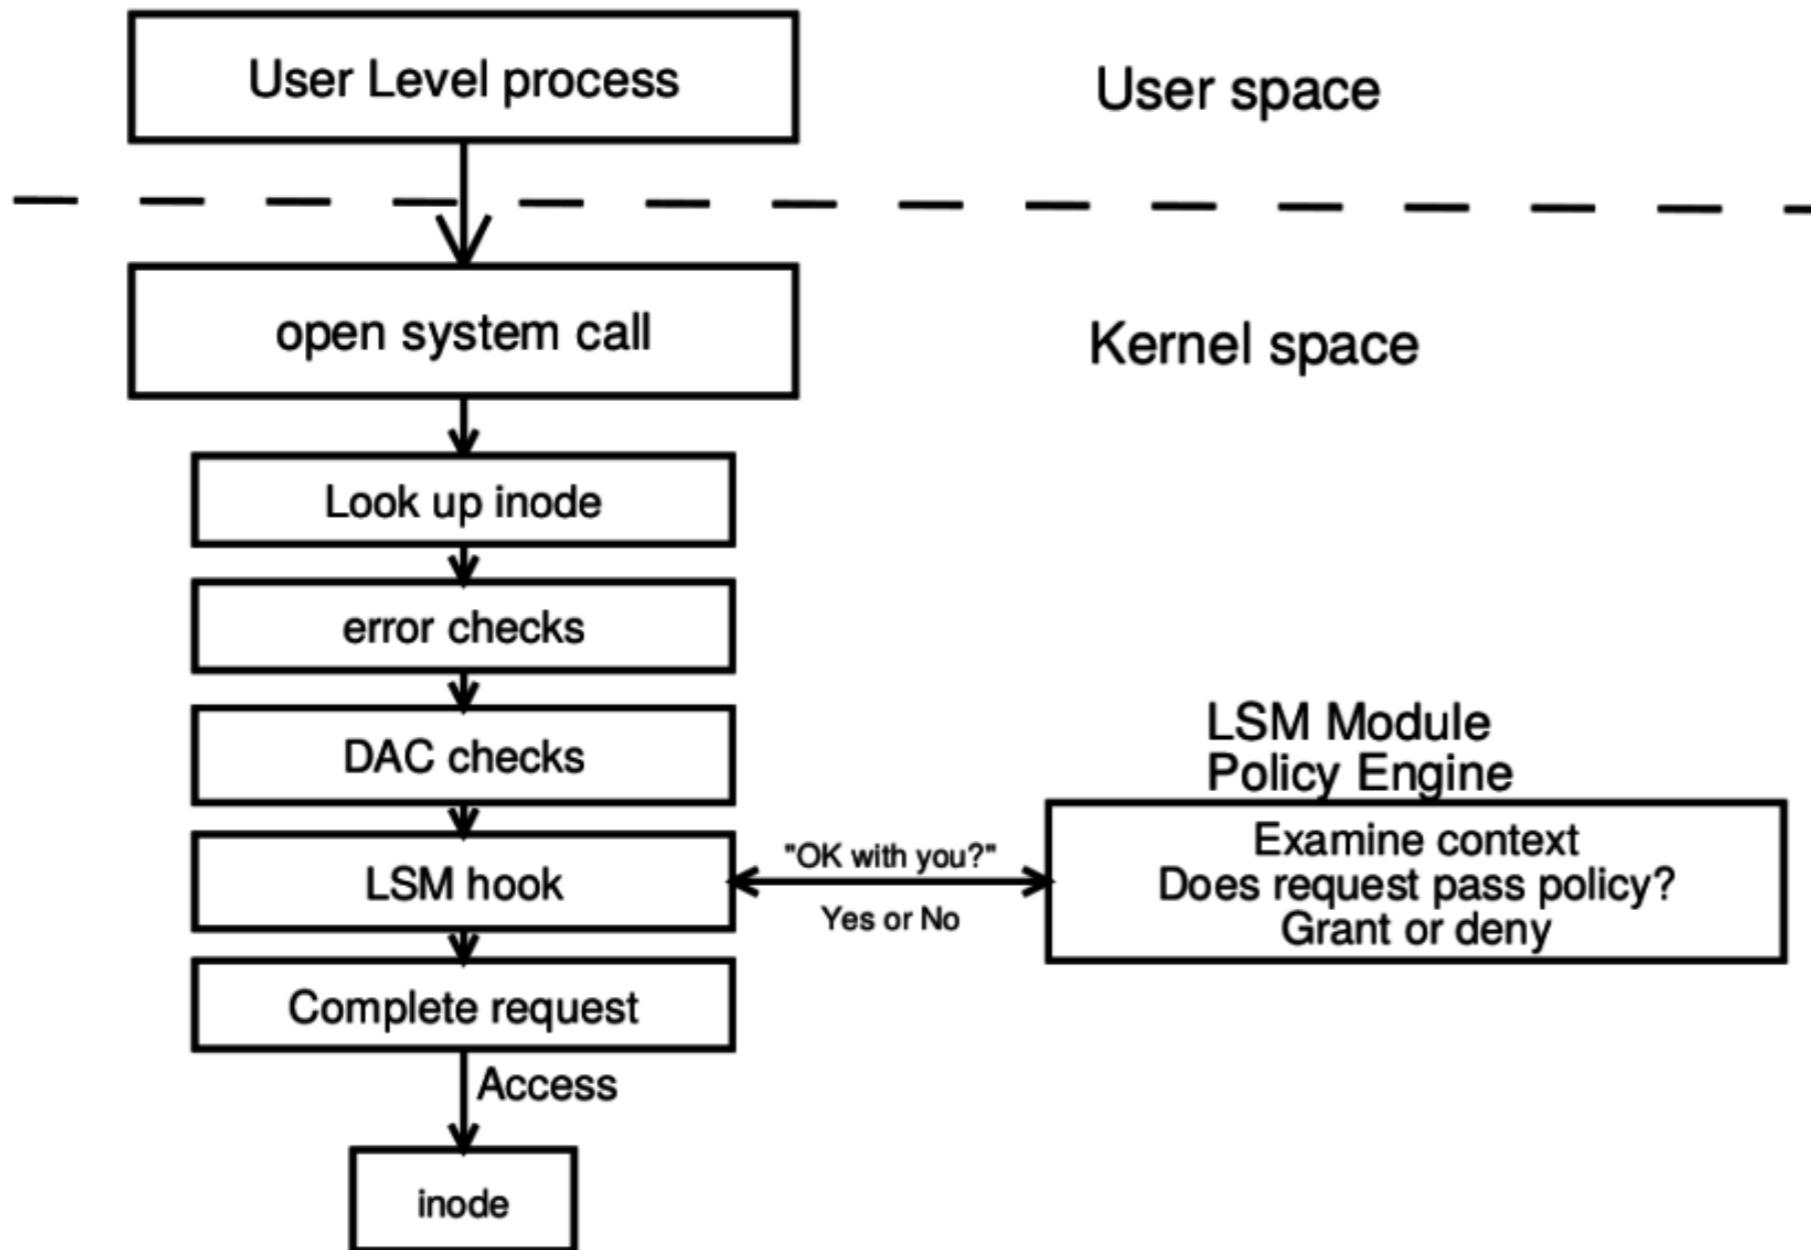


Figure 1: LSM Hook Architecture



# DIY container

SULIA EVANS  
@b0rk

## Containers aren't magic

These 15 lines of bash will start a container running the fish shell. Try it!  
(download this script at [bit.ly/containers-arent-magic](http://bit.ly/containers-arent-magic))

```
wget bit.ly/fish-container -O fish.tar          # 1. download the image
mkdir container-root; cd container-root         #
tar -xf ./fish.tar                             # 2. unpack image into a directory
cgroup_id="cgroup_$(shuf -i 1000-2000 -n 1)" # 3. generate random cgroup name
cgcreate -g "cpu,cpuacct,memory:$cgroup_id"   # 4. make a cgroup &
cgset -r cpu.shares=512 "$cgroup_id"           #      set CPU/memory limits
cgset -r memory.limit_in_bytes=1000000000 \
      "$cgroup_id"                            #
#                                           #
cgexec -g "cpu,cpuacct,memory:$cgroup_id" \    # 5. use the cgroup
      unshare -fmuipn --mount-proc \            # 6. make + use some namespaces
      chroot "$PWD" \                         # 7. change root directory
      /bin/sh -c "
      /bin/mount -t proc proc /proc &&        # 8. use the right /proc
      hostname container-fun-times &&        # 9. change the hostname
      /usr/bin/fish"                           # 10. finally, start fish!
```

# Summary

JULIA EVANS  
@b0rk

## container kernel features

containers are implemented using these Linux kernel features

You can use any of these on their own. When we use them all we call it a "container".

### ♥ pivot-root ♥

set a process's root directory to a directory with the contents of the container image

### ★ cgroups ★

limit memory /CPU usage for a group of processes



only 500 MB of RAM for you!

### ♥ namespaces ♥

allow processes to have their own:

- network      → hostname
- PIDs          → mounts
- users        + more

### ♥ seccomp-bpf ♥

security: prevent dangerous system calls

### ★ overlay filesystems ★

optimization to reduce disk space used by containers which are using the same image

### ★ capabilities ★

security: avoid giving root access



# Takeaways

- Container support has existed in Linux for many years
- Foundations of containerization have been around for decades!
- Automating LXC for portability (i.e., Docker) has revolutionized cloud computing
- Lasting legacy of containers may be enabling the Function-as-a-Service revolution... cloud customers can now pay by the method invocation without any idle costs.