



# CS 423 Operating System Design: Introduction to Linux Kernel Programming (MP2 Walkthrough)

Jack Chen

Some content taken from a previous year's walkthrough by Prof. Adam Bates

# Purpose of MP2



- Understand real time scheduling concepts
- Design a real time schedule module in the Linux kernel
- Learn how to use the kernel scheduling API, timer, procfs
- Test your scheduler by implementation a user level application

# Introduction




- Real-time systems have requirements in terms of response time and predictability
  - Air bag in a car
  - Video surveillance systems
- We will be dealing with periodic tasks
  - Constant period
  - Constant running time
- We will assume tasks are independent

# Periodic Tasks Model



- Liu and Layland [1973] model, each task  $i$  has
  - Period  $P_i$
  - Deadline  $D_i$
  - Runtime  $C_i$



# Rate Monotonic Scheduler (RMS)



- A static scheduler has **complete information** about all the incoming tasks
  - Arrival time
  - Deadline
  - Runtime
  - Etc.
- RMS assigns **higher priority** for tasks with higher rate/shorter period
  - Shorter period results higher priority
  - It always picks the task with the highest priority
  - It is preemptive



- We will implement RMS with an admission control policy as a kernel module
- The scheduler provides the following interface
  - **Registration**: save process info like pid, etc.
  - **Yield**: process notifies RMS that it has completed its period
  - **De-Registration**: process notifies RMS that it has completed all its tasks



- We only register a process if it passes admission control
- The module will answer this question every time:
  - Can the new set of processes still be scheduled on a single processor?
  - Yes if and only if:

$$\sum_{i \in T} \frac{C_i}{P_i} \leq 0.693$$

- Always assumes that

$$C_i < P_i$$

- $C_i$  is the runtime of a task
- $P_i$  is the period to deadline



Floating point operations are very expensive in  
the kernel.

You should NOT use them.

Instead use Fixed-Point arithmetic.



# MP2 User Process Behavior



```
void main (void)
{
    // Proc filesystem
    Register(PID, Period, ProcessTime);
    // Proc filesystem: Verify the process was admitted
    List = Read_Status();
    if (! process in the list) exit(1);

    Yield(PID); // Send yield to Proc filesystem
    while (exist jobs)
    {
        //wakeup_time = t0 - gettimeofday() and factorial
        computation
        do_job();
        Yield(PID); // Send yield to Proc filesystem
    }
    Unregister(PID); // Send yield to Proc filesystem
}
```

# MP2 User Process Behavior



```
void main (void)
{
    // Proc filesystem
    → Register(PID, Period, ProcessTime);
    // Proc filesystem: Verify the process was admitted
    List = Read_Status();
    if (! process in the list) exit(1);

    Yield(PID); // Send yield to Proc filesystem
    while (exist jobs)
    {
        //wakeup_time = t0 - gettimeofday() and factorial
        computation
        do_job();
        Yield(PID); // Send yield to Proc filesystem
    }
    Unregister(PID); // Send yield to Proc filesystem
}
```

# MP2 User Process Behavior



```
void main (void)
{
    // Proc filesystem
    Register(PID, Period, ProcessTime);
    // Proc filesystem: Verify the process was admitted
    → List = Read_Status();
    if (! process in the list) exit(1);

    Yield(PID); // Send yield to Proc filesystem
    while (exist jobs)
    {
        //wakeup_time = t0 - gettimeofday() and factorial
        computation
        do_job();
        Yield(PID); // Send yield to Proc filesystem
    }
    Unregister(PID); // Send yield to Proc filesystem
}
```

# MP2 User Process Behavior



```
void main (void)
{
    // Proc filesystem
    Register(PID, Period, ProcessTime);
    // Proc filesystem: Verify the process was admitted
    List = Read_Status();
    → if (! process in the list) exit(1);

    Yield(PID); // Send yield to Proc filesystem
    while (exist jobs)
    {
        //wakeup_time = t0 - gettimeofday() and factorial
        computation
        do_job();
        Yield(PID); // Send yield to Proc filesystem
    }
    Unregister(PID); // Send yield to Proc filesystem
}
```

# MP2 User Process Behavior



```
void main (void)
{
    // Proc filesystem
    Register(PID, Period, ProcessTime);
    // Proc filesystem: Verify the process was admitted
    List = Read_Status();
    if (! process in the list) exit(1);

    → Yield(PID); // Send yield to Proc filesystem
    while (exist jobs)
    {
        //wakeup_time = t0 - gettimeofday() and factorial
        computation
        do_job();
        Yield(PID); // Send yield to Proc filesystem
    }
    Unregister(PID); // Send yield to Proc filesystem
}
```

# MP2 User Process Behavior



```
void main (void)
{
    // Proc filesystem
    Register(PID, Period, ProcessTime);
    // Proc filesystem: Verify the process was admitted
    List = Read_Status();
    if (! process in the list) exit(1);

    Yield(PID); // Send yield to Proc filesystem
    → while (exist jobs)
    {
        //wakeup_time = t0 - gettimeofday() and factorial
        computation
        do_job();
        Yield(PID); // Send yield to Proc filesystem
    }
    Unregister(PID); // Send yield to Proc filesystem
}
```

# MP2 Process State



- A process in MP2 can be in one of three states
  - a. **READY**: a new job is ready to be scheduled
  - b. **RUNNING**: a job is currently running and using the CPU
  - c. **SLEEPING**: job has finished execution and process is waiting for the next period
- Those are states we should explicitly define in MP2 as they are specific to our scheduler.

# MP2 Extending the PCB



We should  
extend  
PCB to  
hold MP2  
specific  
information

```
struct mp2_task_struct {  
    struct task_struct *linux_task;  
    struct list_head task_node;  
    struct timer_list tasl_timer;  
    pid_t pid;  
    unsigned long period_ms;  
    unsigned long compute_time_ms;  
    unsigned long deadline_jiff;  
    int task_state;  
};
```



# MP2 Scheduling Logic



- What happens when userapp sends YIELD?  
(What does it actually mean when sending YIELD?)
  - Find the calling task
  - Change the state of the calling task to SLEEPING
  - Calculate the time when next period begins
  - Set the timer
    - What should happen if current deadline has passed, but no other tasks are preempting the currently running task?
  - Wake up dispatching thread
  - Put the calling task to sleep (in linux scheduler)

# MP2 Scheduling Logic



- What happens when a task is expired?
  - Change the task to READY
  - Wake up the dispatching thread

# MP2 Scheduling Logic



- What should dispatching thread do?  
Dispatching thread handles our main scheduling logic.
  - Trigger context switch
  - As soon as the context switch wakes up, find the **READY** task with **highest priority**
  - Preempt the currently running task
  - Set the state of new running task to **RUNNING**

# MP2 Scheduling Logic



- We are using a kernel thread to handle our main scheduling logic
- You will need to **explicitly put the kernel thread to sleep** when you're done with your work
- You also need to **explicitly check for signals**
  - Check if should stop working
  - **kthread\_should\_stop()**

# MP2 Scheduler API



- `schedule()` - trigger the kernel scheduler
- `wake_up_process (struct task_struct *)`
- `sched_setscheduler()`: set scheduling parameters
  - FIFO for real time scheduling,  
NORMAL for regular processes, etc.
- `set_current_state()`
- `set_task_state()`

# MP2 Scheduler API Example



- To sleep and trigger a context switch  
`set_current_state(TASK_INTERRUPTIBLE);`  
`schedule();`
- To wake up a process  
`struct task_struct * sleeping_task;`  
`.....`  
`wake_up_process(sleeping_task);`

# MP2 Final Notes



- Develop things incrementally, follow the mp2 description
- Test things one at a time
  - Try to test one feature after you are done with it
  - Use git commits to organize your developments. When things go wildly wrong, you can rollback to where it once worked.
- Use fixed point arithmetic. Don't use double or float
- Use global variables for persistent state
- Remember to cleanup everything
- If you get permission denied during login, you might have produced too many kernel logs. Post privately on piazza and I will help you (when I see it...)
- If your kernel freezes you might be asking too much from kmalloc (some other things could also happen)