

CS 423

Operating System Design:  
Processes and CPU Virtualization

1/24

Ram Alagappan

# Logistics

We have a second TA

Max Qian, [zq2@illinois.edu](mailto:zq2@illinois.edu)

## Office Hours

Ram Alagappan

W: 3:15-4pm, 1306 Everitt Laboratory

Max Qian

Tue: 1-2pm

Fri: 3:30-4:30pm

Xuhao Luo

Wed: 5-6pm

Fri: 2-3pm

# AGENDA / OUTCOMES

## Abstraction

What is a Process ? What is its lifecycle ?

## Mechanism

How does process interact with the OS ?

How does the OS switch between processes ?

ABSTRACTION: PROCESS

# PROGRAM VS PROCESS

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[]) {
    char *str = argv[1];

    while (1) {
        printf("%s\n", str);
        Spin(1);
    }
    return 0;
}
```

Program

Process

# WHAT IS A PROCESS?

Stream of executing instructions and their “context”


```
pushq    %rbp
movq     %rsp, %rbp
subq     $32, %rsp
movl     $0, -4(%rbp)
movl     %edi, -8(%rbp)
movq     %rsi, -16(%rbp)
cmpl     $2, -8(%rbp)
je       LBB0_2
```

Instruction  
Pointer

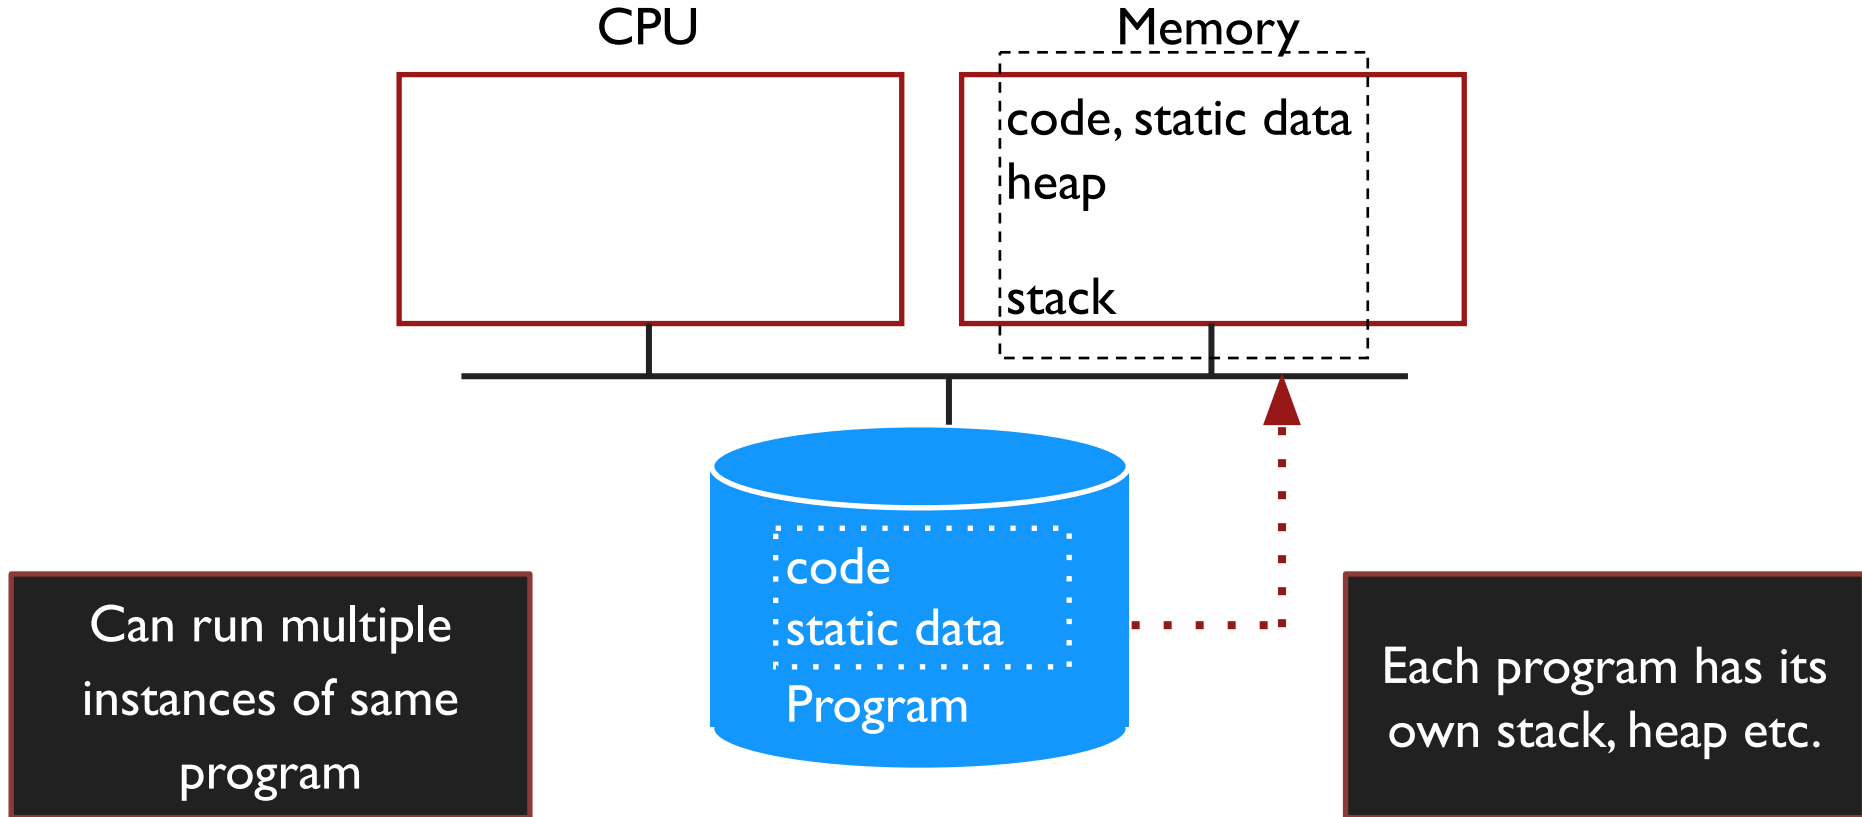
Registers  
Memory addrs

File descriptors

# PROCESS CREATION



# PROCESS CREATION





# PROCESS VS THREAD

Threads: “Lightweight process”


Execution streams that share an address space

Can directly read / write memory


Can have multiple threads within a single process

# SHARING THE CPU

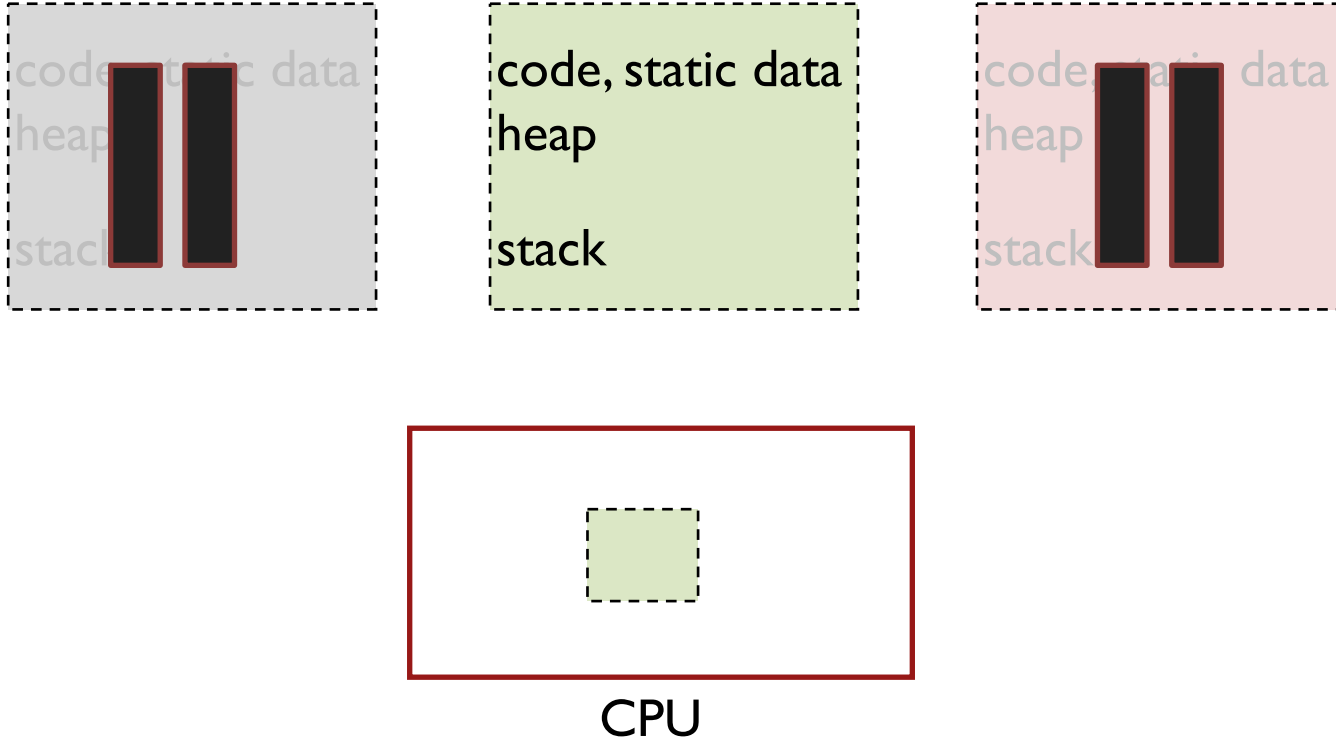
# SHARING CPU



# TIME SHARING



# TIME SHARING



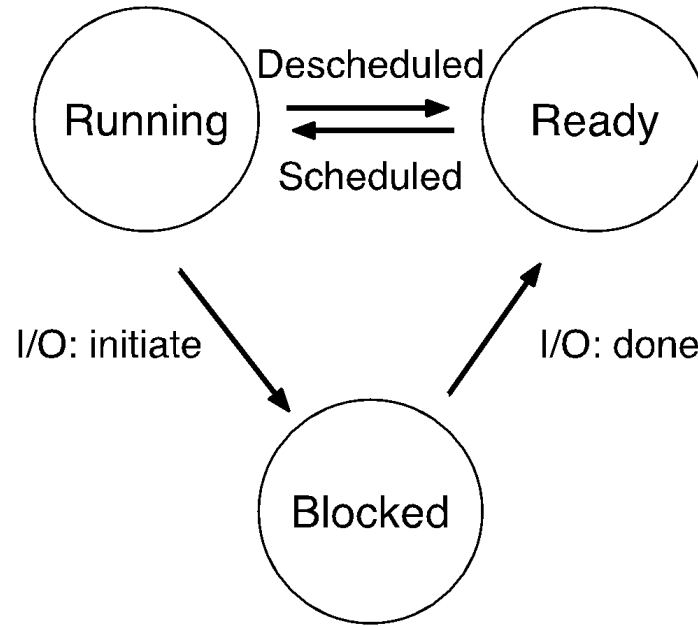
# WHAT TO DO WITH PROCESSES THAT ARE NOT RUNNING ?

OS Scheduler

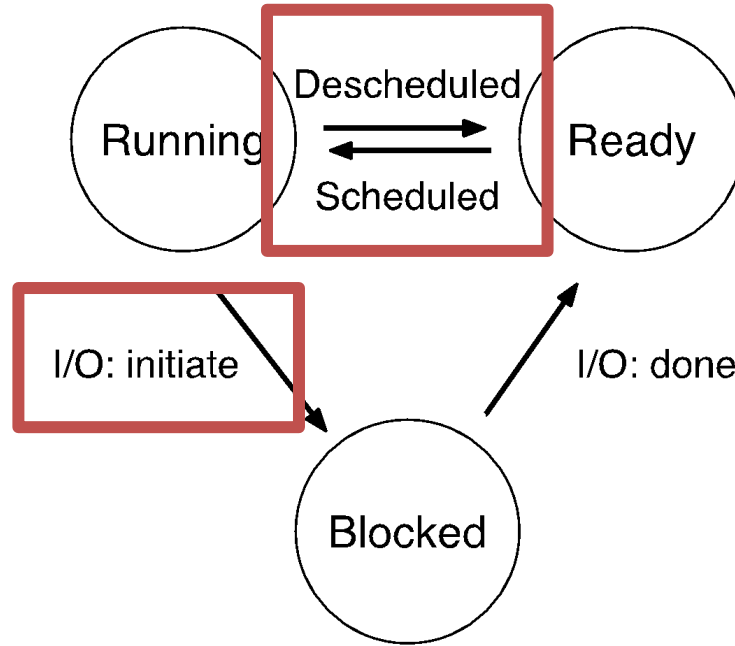
Save context when process is paused

Restore context on resumption

# STATE TRANSITIONS



# STATE TRANSITIONS





# Question

Process 0

io

io

cpu

Process 1

cpu

io

io

Each IO takes 5  
time units

Time	PID: 0	PID: 1
1	RUN:io	READY
2	WAITING	RUN:cpu
3	WAITING	RUN:io
4	WAITING	WAITING
5	WAITING	WAITING
6	RUN:io	WAITING
7	WAITING	WAITING

What happens at time 8?

# CPU SHARING

## Policy goals

- Virtualize CPU resource using processes

- Reschedule process for better CPU utilization? fairness?

## Mechanism goals

- Efficiency: Sharing should not add overhead

- Control: OS should be able to intervene when required

# EFFICIENT EXECUTION

Answer: Direct Execution

Allow user process to run directly on the CPU (no OS intervention)

Create process and transfer control to main()

# What's the Problem with DE?

# EFFICIENT EXECUTION

Problems with DE:

Restricted ops: What if the process wants to do something restricted like allocating more resources, access IO devices, etc?

Switching b/w processes: What if the process runs forever? Buggy? Malicious?

General solution: Limited Direct Execution (LDE)

# PROBLEM 1: RESTRICTED OPS

How can we ensure user process can't harm others?

Solution: privilege levels supported by hardware (bit of status)

User processes run in user mode (restricted mode)

OS runs in kernel mode (not restricted)

How can process access devices?

**System calls** (function call implemented by OS)

SYSTEM CALL

# Syscall

**Trap** instruction :

Jumps into the kernel, changes the processor mode (to kernel)

What is it in x86?

**Ret-from-trap** instruction:

Return from the kernel, change to user mode

What is it in x86?

Libraries usually hide these instructions and give a nicer interface like `read()/write()`



# Syscall

Must save callee registers and instruction pointer to resume after syscall

Where are these saved?

**Kernel stack:** every process has its own kernel stack

Operating System

Hardware

Program

Process A

Run main() ...  
Call system call  
trap into OS

save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

Handle the trap  
Do work of syscall  
return-from-trap

Restore regs (from kstack)  
move to user mode  
jump to PC after trap

# Syscall

How does the hardware know where to jump (i.e., trap handler location) !?

Solution: trap table and system call table

64 tells that this is syscall (other numbers for other exceptional events)

6 tells that this is a `sys_read`

During boot time, OS “configures” the hardware to say where the trap handlers are located...

On trap, hardware simply jumps to this location

OS then knows this is a syscall, uses the syscall number to decide which particular syscall to invoke

# SYSCALL SUMMARY

Separate user-mode from kernel mode for security

Syscall: call kernel mode functions

- Transfer from user-mode to kernel-mode (trap)

- Return from kernel-mode to user-mode (return-from-trap)

To call SYS\_read the instructions we used were

```
movl $6, %eax  
int $64
```

To call SYS\_exec what will be the instructions?

```
movl _____ %eax  
int    _____
```

// System call numbers

```
#define SYS_fork      1  
#define SYS_exit      2  
#define SYS_wait      3  
#define SYS_pipe      4  
#define SYS_write     5  
#define SYS_read      6  
#define SYS_close     7  
#define SYS_kill      8  
#define SYS_exec      9  
#define SYS_open     10
```

# PROBLEM2: HOW TO TAKE CPU AWAY

## Policy

To decide which process to schedule when

Decision-maker to optimize some workload performance metric

## Mechanism

To switch between processes

Low-level code that implements the decision

Separation of policy and mechanism: Recurring theme in OS

# DISPATCH MECHANISM

OS runs **dispatch loop**

```
while (1) {  
    run process A for some time-slice  
    stop process A and save its context  
    load context of another process B  
}
```


Question 1: How does dispatcher gain control?

Question 2: What must be saved and restored?

# HOW DOES DISPATCHER GET CONTROL?

Option 1: **Cooperative Multi-tasking**: Trust process to relinquish CPU through traps

- Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
- Provide special `yield()` system call





# PROBLEMS WITH COOPERATIVE ?

Disadvantages: Processes can misbehave

By avoiding all traps and performing no I/O, can take over entire machine

Only solution: Reboot!

Not performed in modern operating systems

# TIMER-BASED INTERRUPTS

Option 2: **Timer-based Multi-tasking**

Guarantee OS can obtain control periodically

Enter OS by enabling periodic alarm clock

Hardware generates timer interrupt (CPU or separate chip)

Example: Every 10ms

Can user code turn off timer?

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

Handle the trap  
Call switch() routine  
save kernel regs(A) to proc-struct(A)  
restore kernel regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

Note: difference b/w caller vs. callee-saved registers

Handle the trap  
Call switch() routine  
save kernel regs(A) to proc-struct(A)  
restore kernel regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler  
  
restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

Handle the trap

Call switch() routine

save kernel regs(A) to proc-struct(A)

restore kernel regs(B) from proc-struct(B)

switch to k-stack(B) —> this is the key point

return-from-trap (into B)

timer interrupt

save regs(A) to k-stack(A)

move to kernel mode

jump to trap handler

restore regs(B) from k-stack(B)

move to user mode

jump to B's IP

# xv6 Example

```
void
sched(void)
{
    int intena;
    struct proc *p = myproc();

    if(!holding(&ptable.lock))
        panic("sched ptable.lock");
    if(mycpu()->ncli != 1)
        panic("sched locks");
    if(p->state == RUNNING)
        panic("sched running");
    if(readeflags() & FL_IF)
        panic("sched interruptible");
    intena = mycpu()->intena;
    swtch(&p->context, mycpu()->scheduler);
    mycpu()->intena = intena;
}
```

- Scheduler switches to user process in “scheduler” function
- User process switches to scheduler thread in the “sched” function

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

## Swrch() function

- What is on the k-stack of A when a process A has just invoked the swrch?
- What does swrch do?
- What will swrch find on new kernel stack?
- Where does it return to?



# Swch() function

- What is on the k-stack when a process A has just invoked the swch?

Just the caller save registers of A, return address (eip) – where is this exactly?

- What does swch do?

Push remaining registers on old kernel stack (i.e., callee save registers or kernel registers of A)

Save pointer to this context into context structure pointer of old process (A)

Switch esp from old kernel stack (A k-stack) to new kernel stack (B k-stack)

ESP now points to saved context of new process (B k-stack)

Pop callee-save registers from new stack

Return (pops return address, caller save registers – hardware does this)

- What will swch find on new kernel stack?

Whatever was pushed when the new process gave up its CPU in the past

- Where does it return to?

We switched kernel stacks from old process to new process, CPU is now executing new process code, resuming where the process gave up its CPU by calling swch in the past

# Swch impl in xv6

```
# void swch(struct context *old, struct context *new);
#
# Save current register context in old
# and then load register context from new.
.globl swch
swch:
    # Save old registers
    movl 4(%esp), %eax # put old ptr into eax
    popl 0(%eax)      # save the old IP
    movl %esp, 4(%eax) # and stack
    movl %ebx, 8(%eax) # and other registers
    movl %ecx, 12(%eax)
    movl %edx, 16(%eax)
    movl %esi, 20(%eax)
    movl %edi, 24(%eax)
    movl %ebp, 28(%eax)

    # Load new registers
    movl 4(%esp), %eax # put new ptr into eax
    movl 28(%eax), %ebp # restore other registers
    movl 24(%eax), %edi
    movl 20(%eax), %esi
    movl 16(%eax), %edx
    movl 12(%eax), %ecx
    movl 8(%eax), %ebx
    movl 4(%eax), %esp # stack is switched here
    pushl 0(%eax)      # return addr put in place
    ret                # finally return into new ctxt
```

# SUMMARY

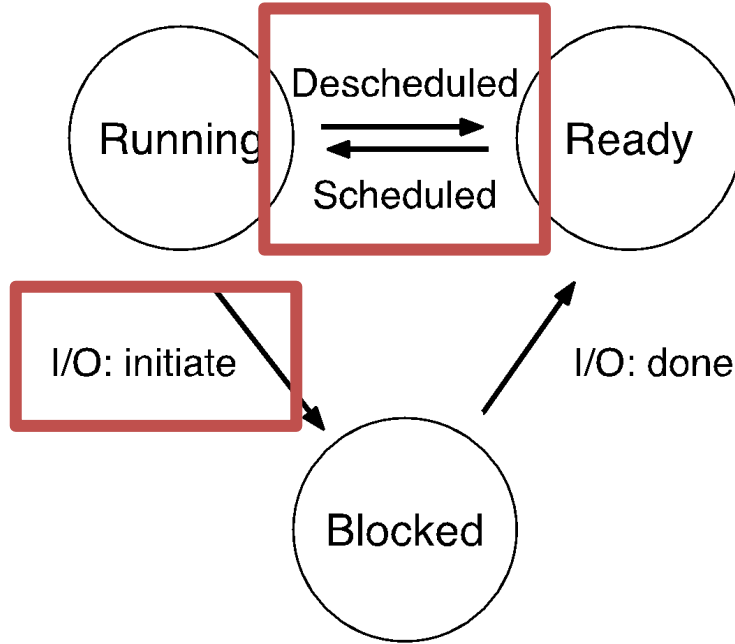
Process: Abstraction to virtualize CPU

Use time-sharing in OS to switch between processes

Key aspects

- Use system calls to run access devices etc. from user mode

- Context-switch using interrupts for multi-tasking



POLICY ?  
Next lecture!