



# CS 423

## Operating System Design: Disks and Disk Scheduling

Ram Alagappan

Acks: Prof. Tianyin Xu and Prof. Adam Bates for slides

# Updates and Logistics



Midterm grades released!

Max: 39.5

Min: 10 (except who didn't submit)

Avg: 32.67

Median: 35

April 25th and 30th lectures - possibly on Zoom  
(will be recorded)

# THREE EASY PIECES



## Three conceptual pieces

Make each application believe it has each resource to itself  
CPU and Memory

1. Virtualization

2. Concurrency

3. Persistence

Provide mutual exclusion, ordering

# MOTIVATION



What good is a computer without any I/O devices?  
keyboard, display, disks


We will focus on disks...

# Questions



- What's the difference between contents in RAM vs. Disk?
- What is the granularity of access in RAM vs. Disk?
- How does the access pattern affect performance in RAM vs. Disk?

# Disk



HDDs not SOTA  
by any means  
But still relevant!

# Hard Disk Internals




A: Track.

B: Sector.

C: Sector of Track.

D: File






# Disk Internals



Seek: move head to the target track

Rotate: wait for target sector to be under head

Transfer: access data





# HDD in Action



<https://www.youtube.com/watch?v=ojGvHDjHPb4>

# Disk Access Time Example



- Disk Parameters
  - Advertised average seek time is **12 ms**
  - Disk spins at **7200 RPM**
  - Transfer rate is **125 MB/sec**
- Assume idle disk (i.e., no queuing delay)

Disk Access Time=seek time +  
rotational delay +  
transfer time

# Disk Access Time Example



- Disk Parameters
  - Advertised average seek time is 12 ms
  - Disk spins at 7200 RPM
  - Transfer rate is 125 MB/sec
- Assume idle disk (i.e., no queuing delay)
- Q1: What is the total time to read 500 random sectors?
- Q2: What is the total time to read 500 sequential sectors (assume on same track)?

Discuss with your neighbors for 3 mins...

# Disk Access Time Example



- What is the total time to read 500 random sectors?

# Disk Access Time Example



- What is the total time to read 500 sequential sectors (assume on same track)?

# Bandwidth



Discuss for two mins...

Random bandwidth?

Sequential bandwidth?

# Disk Access Time Example



See the difference between random and sequential IO speeds on hard drives?

When have you noticed this difference?  
discuss for 1 min...

Always design for sequential IO on HDDs!

Random IO performance (slightly) better with SSDs

Why?




# Disk Scheduling



- Which disk request is serviced first?
  - FCFS
  - Shortest seek time first
  - SCAN (Elevator)
  - C-SCAN (Circular SCAN)

A: Track.  
B: Sector.  
C: Sector of Track.  
D: File




**Disk Scheduling Decision** — Given a series of access requests, on which track should the disk arm be placed next to maximize fairness, throughput, etc?

# FIFO (FCFS) Order




- Method
  - First come first serve
- Pros?
  - Fairness among requests
  - In the order applications expect
- Cons?
  - Arrival may be on random spots on the disk (long seeks)
  - When is it particularly bad?



# SSTF (Shortest Seek Time First)



- Method
  - Pick the one closest on disk (greedy approach)
- Pros?
  - Tries to minimize seek time
- Cons?
  - ???




98, 183, 37, 122, 14, 124, 65, 67  
(65, 67, 37, 14, 98, 122, 124, 183)

# SCAN (Elevator)



- Move outer to inner – service all requests along the way
- Move inner to outer – service all along the way
- Adv compared to SSTF:
  - Bounded time for each request




# C-SCAN (Circular SCAN)



Like SCAN

But, wrap around (i.e., only one direction)

- Adv over SCAN
  - By seeking to opposite side, moves head to where pending requests are likely to be denser
  - More fair
- Cons
  - Do nothing on the return (i.e., higher overhead)



# Scheduling Algorithms



| <i>Algorithm Name</i> | Description  |
|-----------------------|--|
| FCFS                  | First-come first-served  |
| SSTF                  | Shortest seek time first; process the request that reduces next seek time              |
| SCAN (aka Elevator)   | Move head from end to end (has a current direction)                                    |
| C-SCAN                | Only service requests in one direction (circular SCAN)                                 |
| LOOK                  | Similar to SCAN, but do not go all the way to the end of the disk.                     |
| C-LOOK                | Circular LOOK.<br>Similar to C-SCAN, but do not go all the way to the end of the disk. |


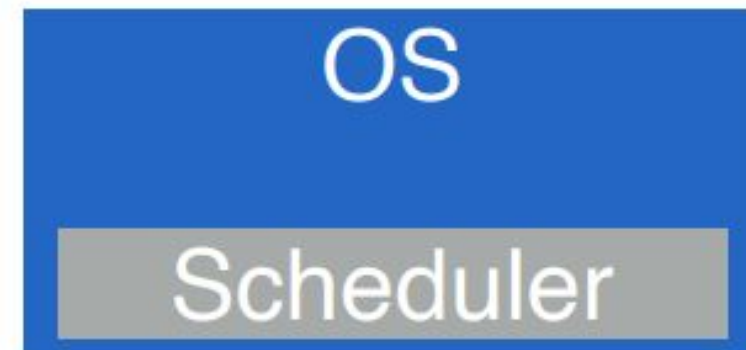
# Who does Scheduling?



The OS?

The disk itself?

Both?





# Linux I/O Schedulers



- What disk (I/O) schedulers are available in Linux?

```
$ cat /sys/block/sda/queue/scheduler  
noop deadline [cfq]
```

- As of Linux 2.6.10, it is possible to change the IO scheduler for a given block device on the fly!
- How to enable a specific scheduler?

```
$ echo SCHEDNAME > /sys/block/DEV/queue/scheduler
```

- SCHEDNAME = Desired I/O scheduler
- DEV = device name (e.g., sda)

# Linux NOOP Scheduler



- Insert all incoming I/O requests into a simple FIFO
- Merges duplicate requests
- When would this be useful?

# Linux NOOP Scheduler



- Insert all incoming I/O requests into a simple FIFO
- Merges duplicate requests (results can be cached)
- When would this be useful?
  - Solid State Drives! Avoids scheduling overhead
  - Scheduling is handled at a lower layer of the I/O stack (e.g., Disk firmware, RAID Controller, Network-Attached)
  - Host doesn't actually know details of sector positions

# Linux Deadline Scheduler



- Imposes a deadline on all I/O operations to prevent starvation of requests
- Maintains 4 queues:
  - 2 Sorted Queues (R, W), order by Sector
  - 2 Deadline Queues (R, W), order by Exp Time
- Scheduling Decision:
  - Check if 1st request in deadline queue has expired.
  - Otherwise, serve request(s) from Sorted Queue.
  - Prioritizes reads (DL=500ms) over writes (DL=5s)  
.Why?

# Linux CFQ Scheduler



- CFQ = Completely Fair Queueing!
- Maintain per-process queues.
- Allocate time slices for each queue to access the disk
- I/O Priority dictates time slice, # requests per queue
- Asynchronous requests handled separately —  
batched together in priority queues

# What Happens?



Assume 2 processes each calling read() with C-SCAN

```
void reader(int fd) {  
    char buf[1024]; int rv;  
    while((rv = read(fd, buf)) != 0) {  
        assert(rv);  
        // takes short time, e.g., 1ms  
        process(buf, rv);  
    }  
}
```

# What Happens?



Assume 2 processes each calling read() with C-SCAN

```
void reader(int fd) {
```

P1: read 100, 101

```
char buf[1024]; int rv;
```

P2: read 900, 901

```
while((rv = read(fd, buf)) != 0) {
```

```
assert(rv);
```

After 1 ms

```
// takes short time, e.g., 1ms
```

P1: read 102, 103

```
process(buf, rv);
```

P2: read 902, 903

```
}
```

```
}
```



# Work Conservation



Work conserving schedulers always try to do work if there's work to be done

Sometimes, it's **better to wait** instead if system anticipates another request will arrive

Possible improvements from I/O merging

# Linux Anticipatory Scheduler



- Deceptive Idleness: A process appears to be finished reading from disk, but is actually processing data. Another (nearby) request is coming soon!
- Bad for synchronous read workloads because seek time is increased.
- Anticipatory Scheduling: Idle for a few milliseconds after a read operation in *anticipation* of another close-by read request.

# Summary



Disks: specific geometry with platters, spindle, tracks, sector, head, etc

$\text{DAT} = \text{seek time} + \text{rotation delay} + \text{transfer time}$

Sequential bandwidth is much higher than random bandwidth

Scheduling approaches: FCFS, SSTF, SCAN, C-SCAN

Schedulers are at multiple layers of the stack

Need to think together (e.g., Linux NOOP)

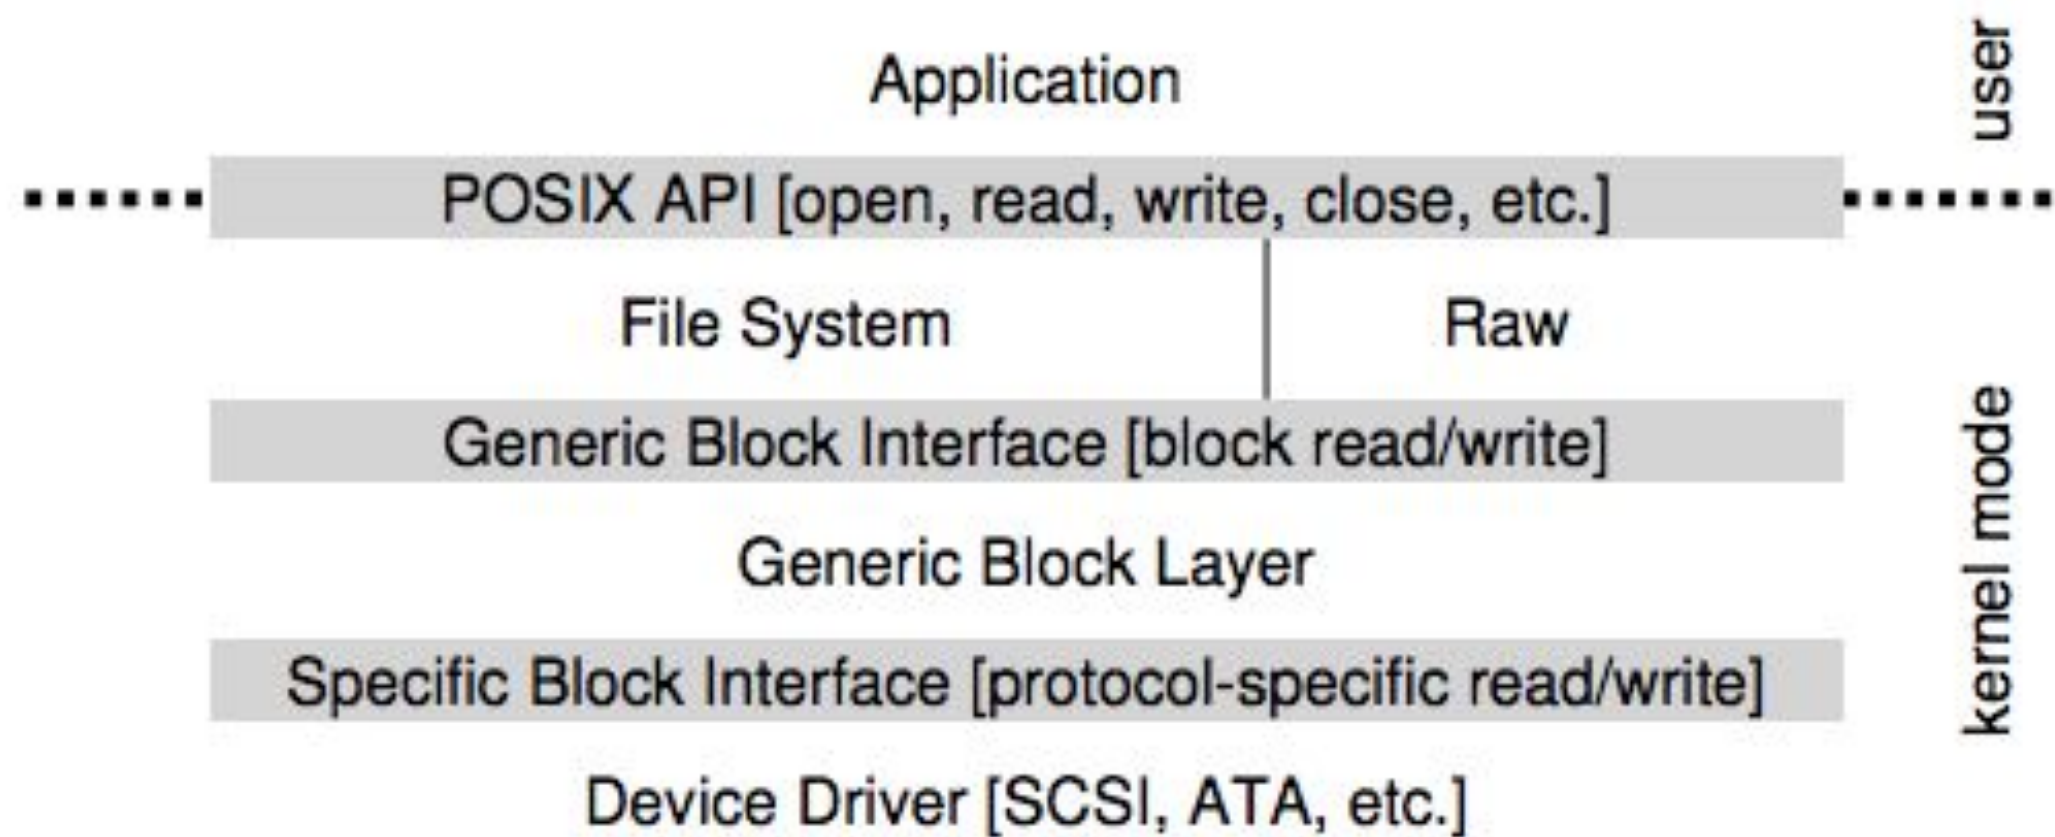
# What is above?



- Above the disk and IO scheduler? **The file system!**

Abstracts many of the underlying details to higher-level applications

1. Presents data as named files— neat, clean abstraction: need not work with sector #s
2. Can be byte-oriented instead of blocks/sectors
3. Offer protection and sharing among users
4. Ensures data reliability



# View of Disk



How does the FS view the disk?

A linear block array! – the block device

It often doesn't need to understand the disk details  
(e.g., positions of sectors)

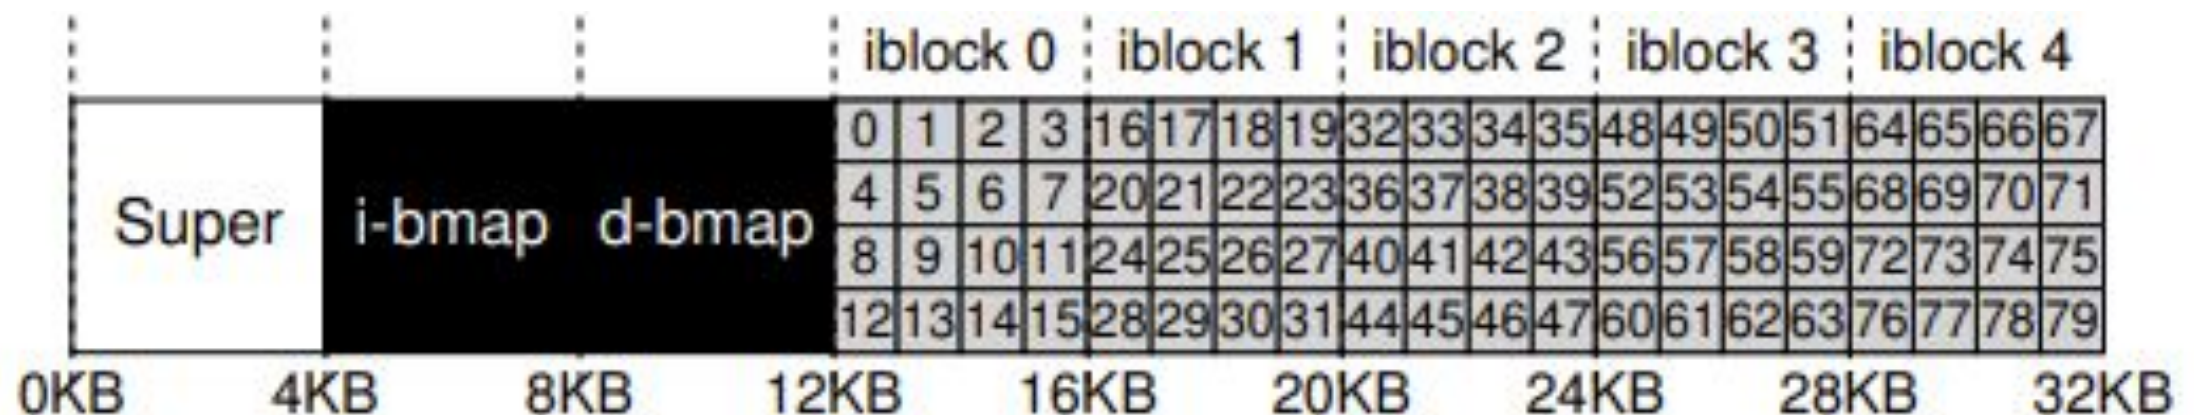
```
nvme0n1      259:0      0 953.9G  0 disk
├─nvme0n1p1  259:1      0  260M  0 part  /boot/efi
├─nvme0n1p2  259:2      0   16M  0 part
├─nvme0n1p3  259:3      0 494.2G  0 part
├─nvme0n1p4  259:4      0 1000M  0 part
└─nvme0n1p5  259:5      0 458.4G  0 part  /
```



# Disk Layout for a FS



Disk layout in a typical file system:



- Data Structures:
  - File data blocks: File contents (not shown)
  - Inodes: low-level file number
  - Directories: File names pointing to inodes
  - Bitmaps: track which disk blocks are free
    - Data bitmap (d-bmap)
    - Inode bitmap (i-bmap)