

CS 423

Operating System Design:
Persistence: FS Implementation

04/23

Ram Alagappan

RECAP

RAID LEVEL COMPARISONS

	Reliability	Capacity	Read latency	Write Latency	Seq Read	Seq Write	Rand Read	Rand Write
RAID-0	0	$C * N$	D	D	$N * S$	$N * S$	$N * R$	$N * R$
RAID-1	1	$C * N / 2$	D	D	$N / 2 * S$	$N / 2 * S$	$N * R$	$N / 2 * R$
RAID-4	1	$(N - 1) * C$	D	2D	$(N - 1) * S$	$(N - 1) * S$	$(N - 1) * R$	$R / 2$
RAID-5	1	$(N - 1) * C$	D	2D	$(N - 1) * S$	$(N - 1) * S$	$N * R$	$N / 4 * R$

What is a File?



Array of persistent bytes that can be read/written

File system consists of many files

Refers to collection of files

Also refers to part of OS that manages those files

Files need names to access correct one

Three types of names

- Unique id: inode numbers
- Path
- File descriptor

File API (attempt 3)



```
int fd = open(char *path, int flag, mode_t mode)
```

```
read(int fd, void *buf, size_t nbyte)
```

```
write(int fd, void *buf, size_t nbyte)
```

```
close(int fd)
```

advantages:

- string names
- hierarchical
- traverse once
- offsets precisely defined

FD Table (xv6)



```
struct file {
    int ref;
    char readable;
    char writable;
    struct inode *ip;
    uint off;
};

struct proc {
    ...
    struct file *ofile[NOFILE]; // Open files
    ...
};

struct {
    struct spinlock lock;
    struct file file[NFILE];
} ftable;
```

LSEEK and READ



```
off_t lseek(int filedesc, off_t offset, int whence)
```

If whence is **SEEK_SET**, the offset is set to offset bytes.

If whence is **SEEK_CUR**, the offset is set to its current location plus offset bytes.

If whence is **SEEK_END**, the offset is set to the size of the file plus offset bytes.

Assume head is on track 1

Suppose we do lseek to X and the sector for X is on track 4

Where is head immediately after lseek?

END RECAP

Shared Entries in OFT



Fork:

```
int main(int argc, char *argv[]) {
    int fd = open("file.txt", O_RDONLY);
    assert(fd >= 0);
    int rc = fork();
    if (rc == 0) {
        rc = lseek(fd, 10, SEEK_SET);
        printf("child: offset %d\n", rc);
    } else if (rc > 0) {
        (void) wait(NULL);
        printf("parent: offset %d\n",
               (int) lseek(fd, 0, SEEK_CUR));
    }
    return 0;
}
```


What is the parent trying to print?

What value will be printed?

Shared Entries in OFT



What's happening here?





```
int fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
int fd2 = open("file.txt"); // returns 4
int fd3 = dup(fd2);         // returns 5
```

DUP



```
int fd1 = open("file.txt"); // returns 12
int fd2 = open("file.txt"); // returns 13
read(fd1, buf, 16);
int fd3 = dup(fd2);          // returns 14
read(fd2, buf, 16);
lseek(fd1, 100, SEEK_SET);
```

How many entries in the OFT (assume no other process)?

Offset for fd1?

Offset for fd2?

Offset for fd3



File system keeps newly written data in memory for a while

Write buffering improves performance (why?)

But what if system crashes before buffers are flushed?

fsync(int fd) forces buffers to flush to disk, tells disk to flush its write cache

Makes data durable



rename(char *old, char *new):

Do we need to copy/move data?

How does the FS implement this?

Does it matter whether the old and new names are in the same directory or different directories?

Rename



rename(char *old, char *new):

- deletes an old link to a file
- creates a new link to a file

Just changes name of file, does not move data

Even when renaming to new directory

What can go wrong if system crashes at wrong time?



(Hard) Link

Inode has a field called “nlinks”

When is it incremented?

When is it decremented?

Deleting Files



What is the system call for deleting files?

Inode (and associated file) is **garbage collected** when there are no references

Paths are deleted when: `unlink()` is called

FDs are deleted when: `close()` or process quits

Symbolic or soft links



A different type of link

Hard links don't work with directory and cannot be cross-FS

```
touch foo; echo hello > foo;
```

Hardlink: `ln foo foo2`

`Stat foo`; what will be the size and inode?

`Stat foo2`; what will be the size and inode?

Softlink: `ln -s foo bar`

`Stat bar`; what will be the size and inode?

Using vs. Implementing



So far, focus on interface of FS

how apps view FS

Today, more about how to implement the
FS itself

Then, crash consistency



FILE SYSTEM IMPLEMENTATION




Very Simple File System

Two aspects:

Data structures – how are files, directories, etc stored on disk

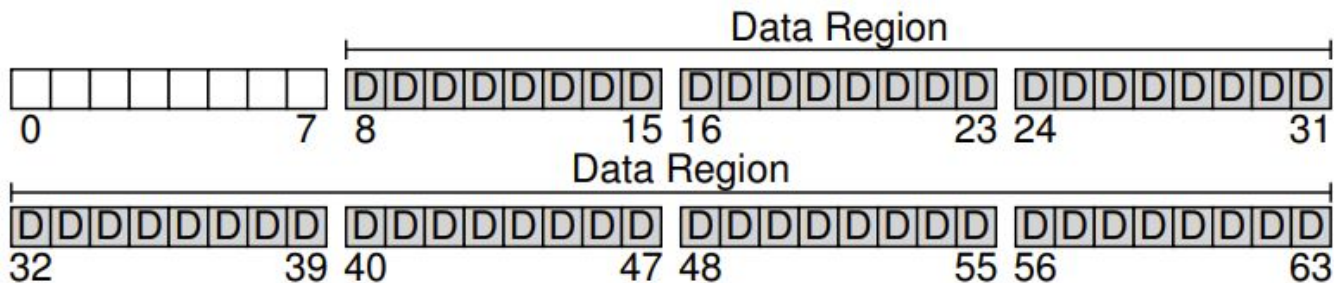
Access methods – how are high-level operations like open, read, write mapped to these DS operations

Assume a small disk partition with 64 blocks

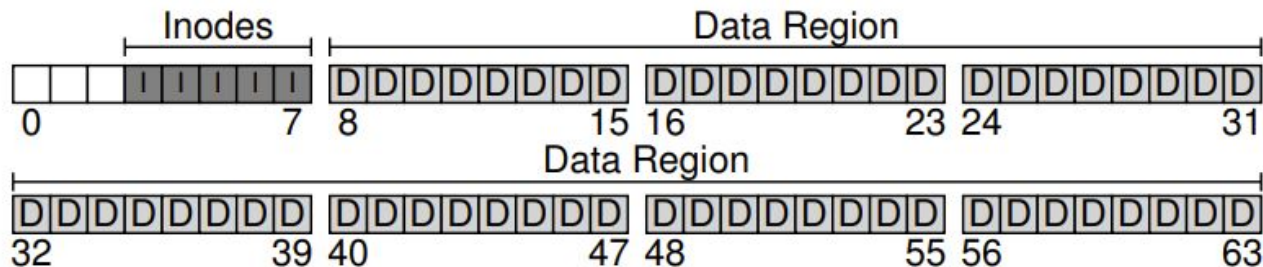


Data and metadata – most space must go for data blocks

VSFS – Data blocks



VSFS – Inodes (metadata)



Called the inode table

With 256-byte inodes, we can store 16 inodes in a block, so totally 80 files can be stored in VSFS

But we can simply scale VSFS to a larger disk

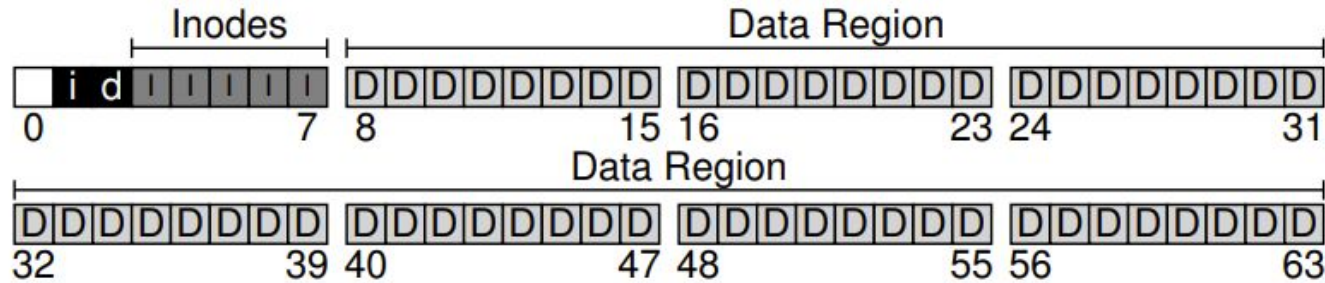
VSFS – Bitmaps (metadata)



Need allocation structures

Free lists – linked list is an option

Most commonly used: bitmap (ib, db)

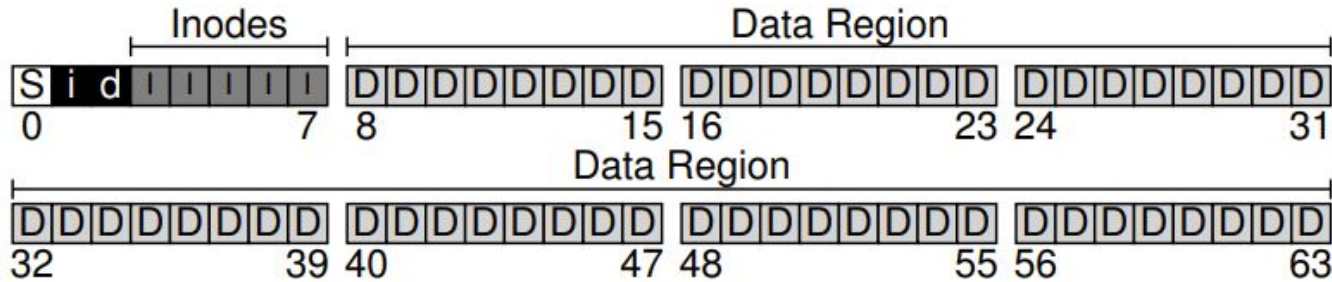


VSFS – (metadata)



What's stored in the first block?

VSFS – Superblock (metadata)




How many data blocks, how many inode blocks

Inode table starting block #

INODE



The Inode Table (Closeup)



Implicitly know the block/sector number

INODE



Size	Name	What is this inode field for?
2	mode	can this file be read/written/executed?
2	uid	who owns this file?
4	size	how many bytes are in this file?
4	time	what time was this file last accessed?
4	ctime	what time was this file created?
4	mtime	what time was this file last modified?
4	dtime	what time was this inode deleted?
2	gid	which group does this file belong to?
2	links_count	how many hard links are there to this file?
4	blocks	how many blocks have been allocated to this file?
4	flags	how should ext2 use this inode?
4	osd1	an OS-dependent field
60	block	a set of disk pointers (15 total)
4	generation	file version (used by NFS)
4	file_acl	a new permissions model beyond mode bits
4	dir_acl	called access control lists

What is the max file size?

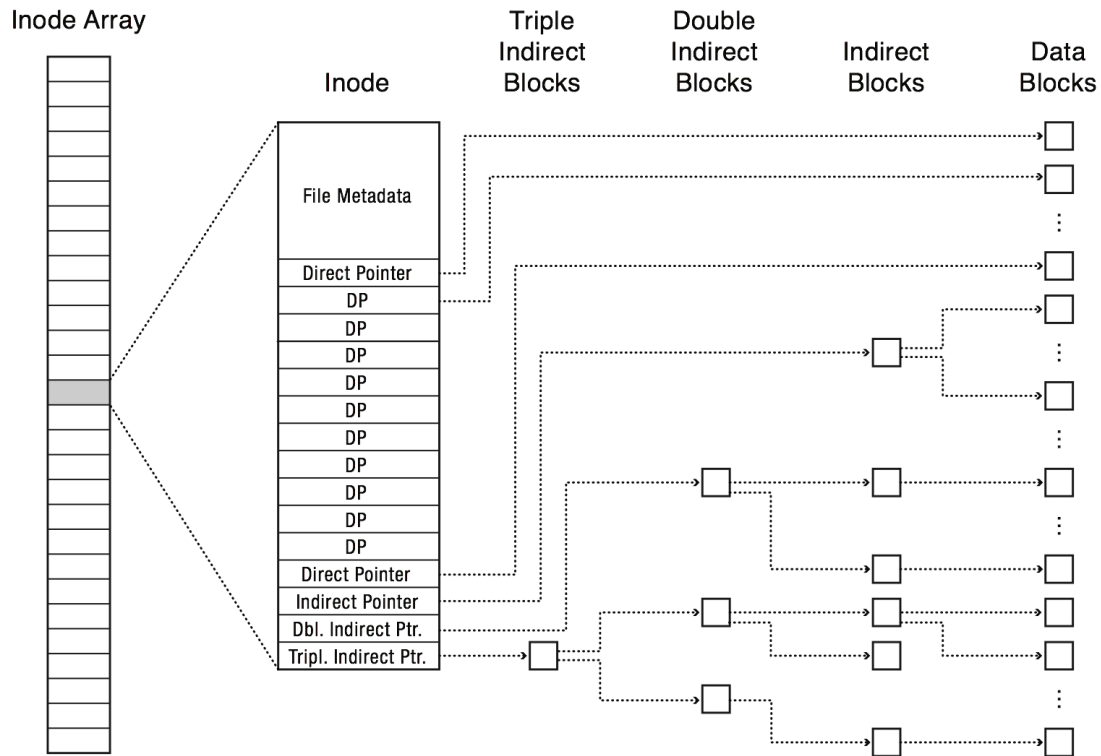


We have 15 block pointers

What is the max file size?

How can we support larger files?

Direct and Indirect Pointers



File Size



File size with one indirect pointer + 12 direct:

$$1024 * 4K + 12 * 4K \text{ -- roughly 4MB}$$

File size with 1 double ID + 1 ID + 12 direct:

$$1024 * 1024 * 4K + 1024 * 4K + 12 * 4K \text{ -- roughly 4GB}$$

Extent based approach



No pointer for every block

<Starting block, num blocks>

Adv compared to pointer approach?

Cons?

Small files: Inlined



- Really small files
- No need to have a separate data block
- Inline them into the inode – can access with fewer disk accesses

Directory Organization



inum	reclen	strlen	name
5	12	2	.
2	12	3	..
12	12	4	foo
13	12	4	bar
24	36	28	foobar_is_a_pretty_longname

What is the inode of this directory?

Where is the directory's content stored?

Creating and Writing File



	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data [0]	bar data [1]
create (/foo/bar)		read write	read	read		read	read		
					read write	write			
write()	read write			read			write		
				write read					
write()	read write				write				write

Why read foo data?

What is written in foo data?

What is written in foo inode?

why is bar inode written upon data write?

Page Cache



Disk access is expensive

Can cache blocks in memory – all FS do this

Integrated with virtual memory

can balance fs cache vs. vm

Also helps write buffering (need to fsync for persistence)

Flushing daemon

Crash Consistency



Basic problem:

Must update many data structure on disk as a unit


What if failure happens in the middle

Types of failure:

- kernel panic

- power failures


Append a Block Example



How many blocks do we need to write to accomplish the append?

Which ones?

Problems



What if only Db is written?

Only i[V2] is written to disk? (2 problems)

Data bitmap is alone written to disk?

Bitmap and data are written:

Data and inode are written:

Bitmap and inode are written:

What's special about the last case?

Metadata vs. Data



FS Metadata consistency vs. Data consistency

FS metadata consistency: internal structures agree with each other

Data consistency: additionally, the data must “make sense” to applications and users



Let inconsistencies happen and take care during reboot

```
UNEXPECTED SOFT UPDATE INCONSISTENCY
** Last Mounted on /
** Root file system
** Phase 1 - Check Blocks and Sizes
** Phase 2 - Check Pathnames
** Phase 3 - Check Connectivity
** Phase 4 - Check Reference Counts
UNREF FILE I=9470237 OWNER=mysql MODE=100600
SIZE=0 MTIME=Feb  9 06:52 2016

CLEAR? no

** Phase 5 - Check Cyl groups
FREE BLK COUNT(S) WRONG IN SUPERBLK
SALVAGE? no

SUMMARY INFORMATION BAD
SALVAGE? no

BLK(S) MISSING IN BIT MAPS
SALVAGE? no

722171 files, 11174066 used, 8118876 free (156260 frags, 995327 blocks, 0.8% fra
gmentation)
\[\033[01;34m\]root@\[\033[00m\]:\[\033[01;34m\]/\[\033[00m\]#
```



Do superblocks match?

Is the list of free blocks correct?

Do number of dir entries equal inode link counts?


Do different inodes ever point to same block?

Are there any bad block pointers?


Do directories contain “.” and “..”?

...


Free Blocks Example



Link Count Example




Duplicate Pointers



FSCK PROBLEMS



Not always obvious how to fix file system image - don't know “correct” state, just consistent one
Simply too slow!



Checking a 600GB disk
takes ~70 minutes

ffsck: The Fast File System Checker

Ao Ma, Chris Dragga, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau

Journaling or WAL




Main idea: write a “note” to a well-known location before actually writing the blocks

If crash, know what to fix and how to do so from the note (instead of scanning the entire disk)

Journaling in Linux ext3



Append a block to an existing file example



Journaling or WAL



First write the txn to journal


Once that is safe, write the actual blocks (this is called checkpointing)

What if crash happens during journal write?

Journal Writes



Can issue one write at a time but is too slow
Must maximize how many writes can be concurrently sent
But sending all 5 blocks together is problematic



How to solve this?

One solution



Incurs a wait or flush between TxB + Data and TxE...
How to do without waiting?

Solution without Wait



What is the problem with DJ?



Next Lecture



Continue CC (more journaling + LFS)

Then:

Advanced storage-1: RAID, NFS

Advanced storage-2: AFS, GFS