

CS 423
Operating System Design:
Persistence: LFS
04/30

Xuhao Luo/Ram Alagappan

LFS PERFORMANCE GOAL



Motivation:

- Growing gap between sequential and random I/O performance
- RAID-5 especially bad with small random writes

Idea: use **disk purely sequentially**

Design for writes to use disk sequentially – how?




File system buffers writes in main memory until “enough” data

- How much is enough?
- Enough to get good sequential bandwidth from disk (MB)

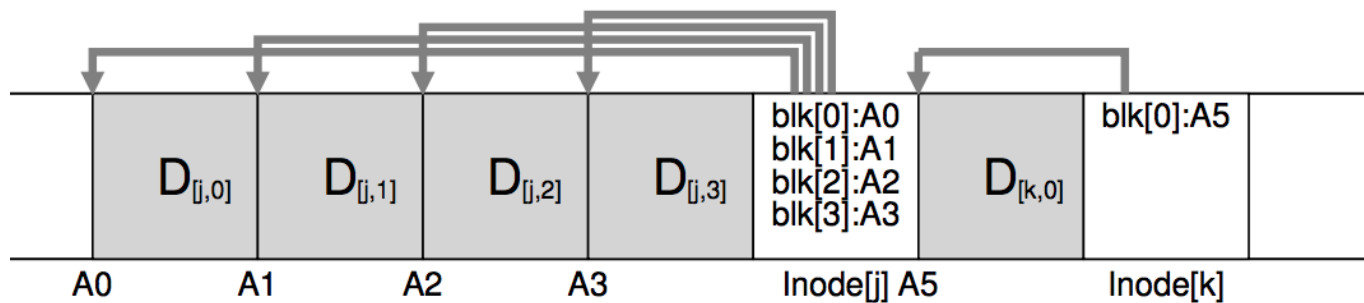
Write buffered data sequentially to new **segment** on disk

Never overwrite old info: old copies left behind

WHERE DO INODES GO?



BUFFERED WRITES



What else is diff from VSFS?



What data structures has LFS removed?

allocation structs: data + inode bitmaps

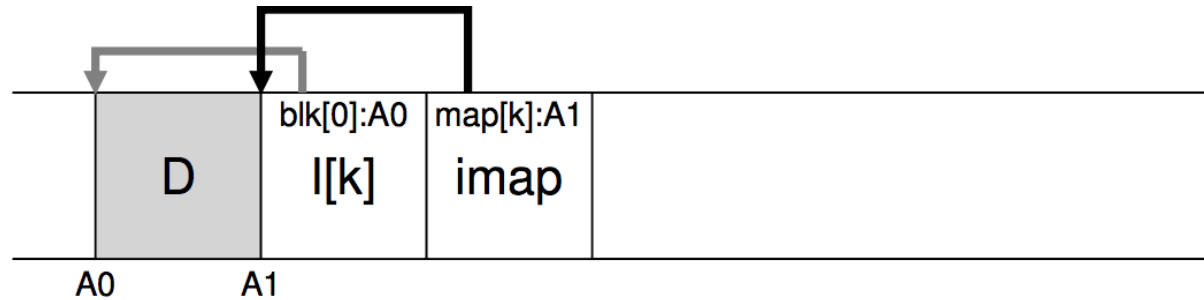

How to do reads?

Inodes are no longer at fixed offset


Use **imap** structure to map:

inode number => inode location on disk

IMAP EXPLAINED




READING IN LFS



1. Read the Checkpoint region
2. Read all imap parts, cache in mem
3. To read a file:
 1. Lookup inode location in imap
 2. Read inode
 3. Read the file block

GARBAGE COLLECTION



WHAT TO DO WITH OLD DATA?



Old versions of files ☐ garbage

Approach 1: garbage is a feature!

- Keep old versions in case user wants to revert files later
- Versioning file systems
- Example: Dropbox

Approach 2: garbage collection




Need to reclaim space:

1. When no more references (any file system)
2. After newer copy is created (COW file system)


LFS reclaims **segments** (not individual inodes and data blocks)

- Want future overwrites to be to sequential areas
- Tricky, since segments are usually partly valid

GARBAGE COLLECTION



GARBAGE COLLECTION



When moving data blocks, copy new inode to point to it

When move inode, update imap to point to it



General operation:

Pick M segments, compact into N (where $N < M$).

Mechanism:

How does LFS know whether data in segments is valid?

Policy:

Which segments to compact?



Is an inode the latest version?

- Check imap to see if this inode is pointed to
- Fast!


Is a data block the latest version?

- Scan ALL inodes to see if any point to this data
- Very slow!

How to track information more efficiently?

- **Segment summary** lists inode and data offset corresponding to each data block in segment (reverse pointers)

SEGMENT SUMMARY



```
(N, T) = SegmentSummary[A];  
  
inode = Read(imap[N]);  
  
if (inode[T] == A)  
    // block D is alive  
else  
    // block D is garbage
```




General operation:

Pick M segments, compact into N (where $N < M$).

Mechanism:

Use segment summary, imap to determine liveness

Policy:

Which segments to compact?

- clean most empty first
- clean coldest (ones undergoing least change)
- more complex heuristics...



What data needs to be recovered after a crash?

- Need imap (lost in volatile memory)


Better approach?

- Occasionally save to **checkpoint region** the pointers to imap pieces

How often to checkpoint?

- Checkpoint often: random I/O
- Checkpoint rarely: lose more data, recovery takes longer
- Example: checkpoint every 30 secs

CRASH RECOVERY



CHECKPOINT SUMMARY



Checkpoint occasionally (e.g., every 30s)

Upon recovery:

- read checkpoint to find most imap pointers and segment tail
- find rest of imap pointers by reading past tail

What if crash during checkpoint?

CHECKPOINT STRATEGY



Have two checkpoint regions

Only overwrite one checkpoint at a time

Use checksum/timestamps to identify newest checkpoint





Journaling:

Put final location of data wherever file system chooses
(usually in a place optimized for future reads)

LFS:

Puts data where it's fastest to write, assume future reads
cached in memory

Other COW file systems: WAFL, ZFS, btrfs