

CS 423

Operating System Design:

Persistence: NFS/AFS

05/02

Ram Alagappan

# NFS Distributed File System

NFS: more of a protocol than a particular file system

Many companies have implemented NFS: Oracle/Sun, NetApp, EMC, IBM

# Recap



So far, local file systems:

VSFS – in-place update file system

LFS – always append to the end of the log

Everything including the inode, imap, data blocks live on the log

The log is the file system...



# Questions (2 min)

What is the biggest downside of LFS over VSFS (or in general in-place update file systems?)

What is one workload that will perform almost the same in both?

# Questions (2 mins)

Can you think of one workload where LFS will be much better than VSFS?

Can you think of one workload where VSFS will be much better than LFS?

What is a useful feature that LFS can provide that VSFS cannot (at least efficiently)?

# End Recap

Today:


Switching from local FS to distributed file systems

Will look into two popular ones:

- NFS (network FS)

- AFS (Andrew FS)

# NFS Arch





# Benefits?

Sharing across machines

Central admin

# Goal: Simple Server Crash Recovery

Why do servers crash?

# Goal: Simple Crash Recovery


Strategy-1: server returns fd upon open, client passes fd on each call

```
int fd = open("foo", O_RDONLY);  
read(fd, buf, MAX);  
read(fd, buf, MAX);  
...  
read(fd, buf, MAX);
```

# Goal: Simple Crash Recovery

Strategy-1: server returns fd upon open, client passes fd on each call

```
int fd = open("foo", O_RDONLY);  
read(fd, buf, MAX);  
read(fd, buf, MAX);  
...  
read(fd, buf, MAX);
```



Server crash!

# Problems

Complicates crash recovery.

Server crash – what happens? What must client do?

Client crashes – what happens?

# General idea: Statelessness

Server keeps no state: no fd □ file map, no file position pointer

Server doesn't keep any state about a client

Client passes all info needed in each call to server

Advantage:

- no special crash recovery - the server just starts running again

- client might have to retry a request

# Pass all info – option-1

*Stateless* protocol: server maintains no state about clients

Need API change. One possibility:

`read(char *path, buf, size, offset)`

Specify path and offset each time

Pros? Cons?

# Pass all info – option-2

*Stateless* protocol: server maintains no state about clients

Use file handles

```
fh = open(char *path);  
pread(fh, buf, size, offset);  
pwrite(fh, buf, size, offset);
```

File Handle = <volume ID, inode #>

Think of volumeID as the file system ID, for simplicity assume only one FS



# Pass all info – option-2

What is the problem with using the inode number as the file handle?

What can go wrong?

# Pass all info – option-2

*Stateless* protocol: server maintains no state about clients

Use file handles

```
fh = open(char *path);  
pread(fh, buf, size, offset);  
pwrite(fh, buf, size, offset);
```

File Handle = <volume ID, inode #, **generation#**>

Opaque to client, purpose of generation#? when incremented?

# Some NFS calls

Lookup – notice no open (open == series of lookups)

GetAttr

Read

Write

Who keeps the fd to fh mapping?

# Virtual File System

Allows different FSES to be plugged into the OS

Acts a common thin layer above all file systems

Can also hide distributed file systems!

NFS engineers designed VFS to make using NFS easier but today VFS is used for all FS in Linux

# Reading a File on NFS

---

## Client

```
fd = open("/foo", ...);  
Send LOOKUP (rootdir FH, "foo")
```

Receive LOOKUP reply  
allocate file desc in open file table  
store foo's FH in table  
store current file position (0)  
return file descriptor to application

---

```
read(fd, buffer, MAX);  
Index into open file table with fd  
get NFS file handle (FH)  
use current file position as offset  
Send READ (FH, offset=0, count=MAX)
```

Receive READ reply  
update file position (+bytes read)  
set current file position = MAX  
return data/error code to app

---

---

## Server

Receive LOOKUP request  
look for "foo" in root dir  
return foo's FH + attributes

Receive READ request  
use FH to get volume/inode num  
read inode from disk (or cache)  
compute block location (using offset)  
read data from disk (or cache)  
return data to client

# Close() a file?

What happens?

No server communication

# Failures

What do clients do when they don't get a response?

Request lost

Server down

Reply lost

# Simplifying Recovery with Idempotency

All cases are handled uniformly

read

write

mkdir

creat



# Client-side caching

Cache data for performance

For both reads and writes

What are the problems?

# PI: update visibility

Scenario: edit a file and move on to a different workstation

Solution: flush-on-close

Drawbacks?

# P2: stale cache

Cached content could be old

Solution: getattr

What problems will this introduce?

# Write buffering on server

Can server buffer writes?

BB-Ram to make writes fast with failure-resilience

NFS - a huge success! most possibly due to “open market” approach

NetApp, EMC, IBM all sell NFS servers

# AFS

Main goal: scale

Better cache consistency semantics

Successful system from CMU

- many novel ideas

- many institutions use AFS

# AFSv1

Whole-file caching instead of block caching

Keep file on disk (not memory like NFS)

Open a file “/path/to/file” – fetch it from the server

Writes and reads – no interaction with server... purely local

Close a file – send it to the server

Next open – use TestAuth message to check if file is up-to-date, if yes  
use cached copy else fetch again

# AFSv1 Problems

Path traversal costs – too much CPU to convert paths to inodes

Too many TestAuth message – most unnecessary...why?

Load imbalance across servers

Server process/threads structure

# AFSv2 - Callbacks

How to reduce TestAuth messages

Callback – promise from server to client that it will inform when the file changes

Implications – ?



# AFSv2 - FIDs

How to reduce CPU cost for path resolution

Use an FID

FID = <volume\_ID, inode, uniqifier> - similar to NFS

# AFSv2

fd = open(/foo/bar.txt, ...)

fetch(/ FID, "foo")

Look for "foo" in /

create callback for foo for client

return foo's content and FID

write foo to disk

record callback of foo

fetch(foo FID, "bar.txt")

Look for "bar.txt" in foo

create callback for bar.txt for client

return bar's content and FID

write bar.txt to disk

record callback of bar.txt

local open cached bar.txt

return fd to app

# AFSv2

read(fd...) – local

write(fd...) – local

close(fd) – flush to server using Store message – send whole file

Next open...

fd = open("/foo/bar.txt", ...);

if (callback(foo) == VALID)

    use local copy foo

else Fetch

if (callback(bar.txt) == VALID)

    open local cached copy

    return file descriptor

else Fetch then open and return fd

# Cache consistency

Update visibility: when?

Processes on different machines

upon close, flush to server

visible to other clients now

Processes on the same machine

immediately visible

Last closer wins – one consistent version uploaded – compare to NFS?

# Cache consistency

Invalidate stale cache: when?

after server gets latest, breaks callback

Is it possible for a client to use an older copy after the file has been changed at the server?

Is it possible for a client to open an older copy?

# Crash Recovery

Client crashes – how to handle?

Treat cache as suspect – should throw cached files?

Server crashes

Keeps callback state in memory – can lose, how to handle?

use heartbeats ... if not heard from server, suspect cache

# Performance Comparison

N = blocks, Lnet, Lmem, Ldisk – latency of net, mem, and disk

NFS

AFS

Small file, sequential read

Small file, sequential re-read

Large file, sequential read

Large file, sequential re-read

Large file, single read

Small file, sequential write (new file)

Large file, sequential write (new file)

Large file, sequential overwrite

Large file, single overwrite