

CS 423

# Operating System Design: Segmentation & Paging

Feb 05

Ram Kesavan

# Logistics

Lecture slides are initially uploaded as “tentative”; finalized soon after

MP0: Due on 02/09 11:59pm CT

MP1:

Will get posted by tonight

Walkthrough by Gabriella on 2/10

4Cr: Linux APIs paper; due by 2/10 2pm CT

Unsubmitted reviews will hurt your participation grade

Google Apps @Illinois problems

# AGENDA / LEARNING OUTCOMES

Memory virtualization

- Segmentation: technique used by older systems


- Pros and cons?

- Paging & modern systems

# RECAP

# ABSTRACTION: ADDRESS SPACE


Virtual Memory



Virtual Address Space:  
Each process has its  
own address range

OS provides that  
illusion by mapping to  
physical memory

Physical Memory




# HOW TO VIRTUALIZE MEMORY

Problem: Addresses are “hardcoded” into process binaries  
How to avoid collisions?

End of Recap

# Mechanisms for Virtualization

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation



Limited practicality, has many problems

More practical, still has some problems

x86 and Linux: Paging, TLB, Efficient Page Tables

# SEGMENTATION


Divide VM address space into logical segments

- Each segment corresponds to logical entity in address space  
**(code, stack, heap)**

Each segment with its own base + bounds register

Each segment can independently:

1. Be placed in physical memory
2. Grow and shrink (not code & static data)
3. Be protected (read/write/exec)



# SEGMENTED ADDRESSING

Process should specify segment and offset within segment

How does process designate a particular segment?

- Use portion of logical/VM address
  - MSBs of logical address => segment
  - LSBs of logical address => offset within segment

# Segmentation Implementation

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 16KB VM address space = 14 bit logical address with 4 segments

How many bits  
for segment?

Segment	Base	Bounds	R	W
00 (code)	32KB	2KB	1	0
01 (heap)	34KB	3KB	1	1
11 (stack)	28KB	2KB	1	1
10 (unused)	0x0000	0x000	0	0

How many bits  
for offset?

# Example Translation

Segment	Base	Bounds	R W
0	32KB	2KB	1 0
1	34KB	3KB	1 1
3	28KB	2KB	1 1
2	0x0000	0x000	0 0

Translate logical (in hex) to physical  
0x1108:

Bottom 12 bits =

Offset =

Is Offset < Bounds?

Top 2 bits =

Segment number:

Physical addr:

base + offset =

# Example Translations

Segment	Base	Bounds	R W
0	32KB	2KB	1 0
1	34KB	3KB	1 1
3	28KB	2KB	1 1
2	0x0000	0x000	0 0

Translate logical (in hex) to physical  
0x1108:

Bottom 12 bits = 0x108  
Offset = 264


Is  $264 < 3\text{KB}$  ( $0x108 < 0xc00$ )? Yes  
Top 2 bits = 01

Segment number: 1  
Physical addr:

$$34\text{KB} + 264 = 35080$$

$$0x8800 + 0x108 = 0x8908$$

# VIRTUAL->PHYSICAL TRANSLATION



Heap Segment

Virtual address range : 4 KB to 7 KB

Physical address range: 34 KB to 37 KB

Virtual address 6 KB (hex: 0x1800)

Which segment?


What is the offset?

How does the MMU translate that?

Physical address = Base + Offset

=

# VIRTUAL->PHYSICAL TRANSLATION



## Heap Segment

Virtual address range : 4 KB to 7 KB

Physical address range: 34 KB to 37 KB

Virtual address 6 KB (hex: 0x1800)

2 msb (0x1): heap segment ref

12 lsb (0x800) = 2 KB offset


How does the MMU translate that?

Bounds check: 2KB < 3KB? Yes

Physical address = Base + Offset

$$= 34 \text{ KB} + 2 \text{ KB} = 36 \text{ KB}$$

# STACK ADDRESS TRANSLATION



Stack Segment

Virtual address range : 16 KB to 14 KB

Physical address range: 28 KB to 26 KB

Virtual address 15 KB (hex: 0x3C00)


Which segment?

What's the offset?

How does the MMU translate that?

Physical address =?

# STACK ADDRESS TRANSLATION



Stack Segment

Virtual address range : 16 KB to 14 KB

Physical address range: 28 KB to 26 KB

Virtual address 15 KB (hex: 0x3C00)

top 2 bits (0x3): stack segment ref,  
offset is 0xC00 = 3 KB

How does the MMU translate that?

**Negative offset** = offset - max segment size

$$= 3 \text{ KB} - 4\text{KB} = -1 \text{ KB}$$

**Add to base** =  $28 \text{ KB} - 1 \text{ KB} = 27 \text{ KB}$

Bounds check on  $\text{abs}(\text{offset})$ , so is  $1\text{KB} < 2\text{KB}$ ? Yes

HW needs to track one more bit of info for every segment  
(grows positively/negatively)

# SEGMENTATION IMPLEMENTATION

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 16KB VM address space = 14 bit logical address with 4 segments

Segment	Base	Bounds	R W	Grows?
00 (code)	32KB	2KB	1 0	1
01 (heap)	34KB	3KB	1 1	1
11 (stack)	28KB	2KB	1 1	0
10 (unused)	0x0000	0x000	0 0	0

# SEGMENTATION: ADVANTAGES

Enables sparse allocation of address space

Stack and heap can grow dynamically

- Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc lib calls sbrk())
- Stack: OS recognizes reference outside legal segment, extends stack implicitly

Different protection for different segments

- Enables sharing of selected segments (2 processes share code)
- With no write permission for code segment


Supports dynamic relocation of each segment

# SEGMENTATION: DISADVANTAGES



Each segment must be allocated contiguously

Physical memory gets fragmented

May not have sufficient physical memory for large segments?



# COMPACT & REARRANGE SEGMENTS



# PAGING

(most modern systems including Linux use this)

# PAGING


Goal: Eliminate requirement that segment is contiguous

Eliminate external fragmentation

Idea:

Divide address spaces and physical memory into fixed-sized pages/frames


Size:  $2^n$ , Example: 4KB



# TRANSLATION of 4KB PAGE ADDRESS

How to translate logical address to physical address?

- MSBs of virtual address => page number
- LSBs bits of virtual address => offset within page




# ADDRESS FORMAT

Page Size	Low Bits (offset)	N-Bit Address	High Bits (vpn)	Virt Pages
16 bytes		10		
1 KB		20		
1 MB		32		
512 bytes		16		
4 KB		32		

# VIRTUAL -> PHYSICAL PAGE MAPPING

#bits in virtual address need not be equal to #bits in physical address



How should OS translate VPN to PPN/PFN?

# LINEAR PAGE TABLE


# What is an obvious data structure?

## Linear page table (array)

## Indexed by VPN

# Yields PFN (Physical Frame Number)

## A Single PTE (Page Table Entry):



# ADVANTAGES OF PAGING

No external fragmentation

- A page can be placed in any frame in physical memory

Fast to allocate and free

- Alloc: No searching for suitable free space
- Free: No need to coalesce/compact free space

Simple to swap-out portions of memory to disk (later lecture)

- Page size matches disk block size!
- Can run process when some pages are on disk
- Add “present” bit to PTE

# HOW BIG IS A PAGE TABLE?

Assume 32-bit address

Assume 4KB pages

Assume 4 byte page table entries (PTE)

How large is PT for each process?

Implications?

# HOW BIG IS A PAGE TABLE?

Assume 32-bit address

Assume 4KB pages => 12 bits for offset  
20 bits for VPN

Assume 4 byte page table entries (PTE)

How large is PT for each process?  $2^{20} * 4 = 4 \text{ MB}$

Implications? For 1K processes, we need 4GB of PTs!

# IMPLICATIONS

Page tables may be substantial

- Simple page table: requires PTE for all pages in address space  
Entry needed even if page not allocated ?

Additional memory reference to page table

- Very inefficient, so...
- ...page table must be stored in memory
- MMU stores only base address of page table

Could choose larger page size

- Leads to (internal) fragmentation

**Next Lecture: Paging and TLBs**